Lecture Notes for

Neural Networks and Machine Learning

Fully Convolutional Learning

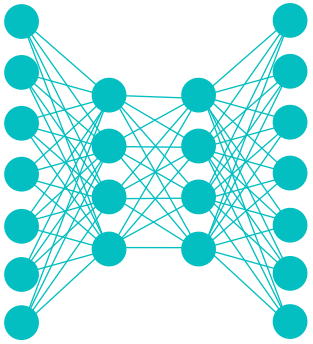

Logistics and Agenda

- Logistics
  - Lab one due soon!
- Agenda
  - Town Hall
  - Paper Presentation
  - Segmentation
    - Semantic (this time)
    - Object (partially this time, maybe)
    - Instance (next time)

Lab One Town Hall

Student Paper Presentation

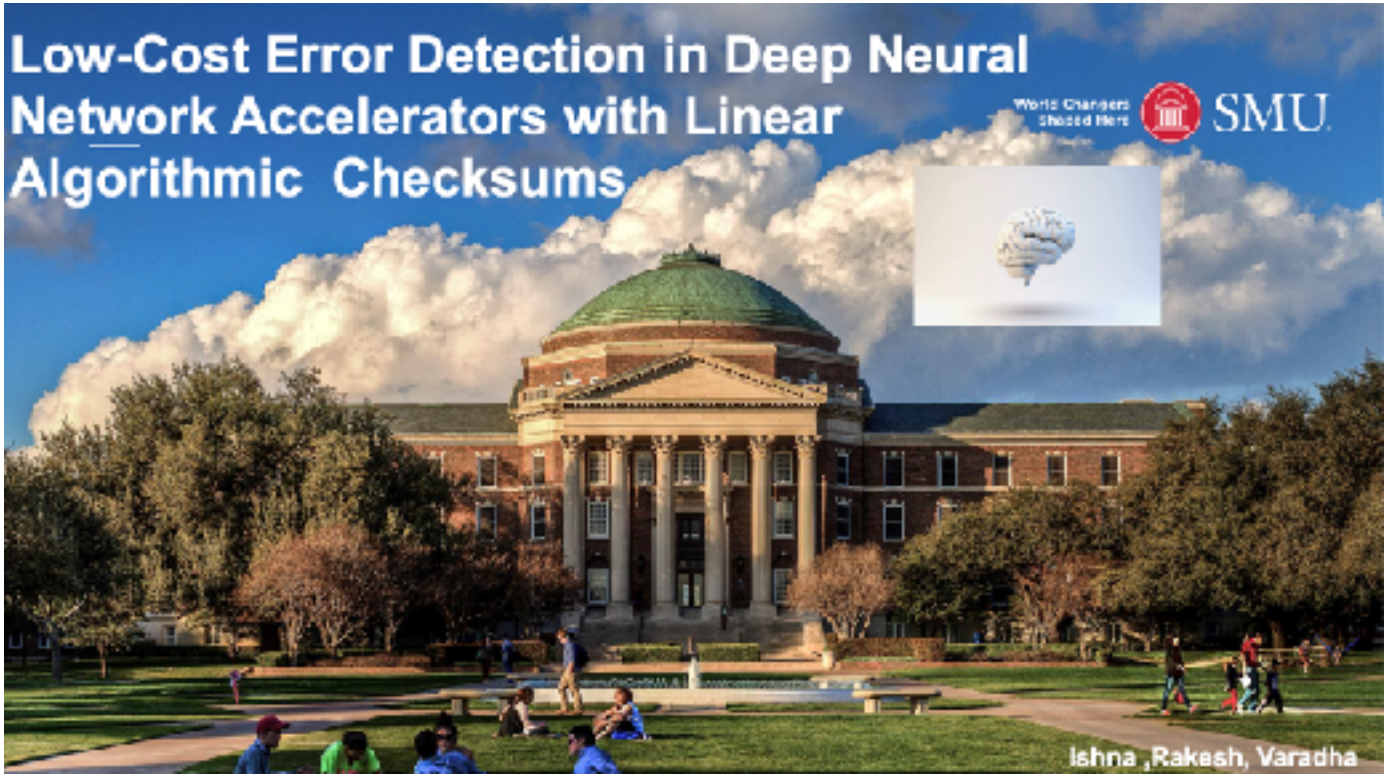

Types of Fully Convolutional Problems

- Semantic Segmentation
- Object Detection
- Instance Segmentation

Semantic Segmentation

Karandeep Singh @kdpsinghlab · 10h

Statistician: Do you ever use statistics?

ML researcher: Nope. Never.

Statistician: What about when reading a paper?

ML: Nope. Never.

Statistician: Ok. So if you're reading an ML paper comparing lots of models, how do you know which one is the best?

ML: Bold font.

Semantic Segmentation

- Given a set of pixels, classify each pixel according to what instance it belongs

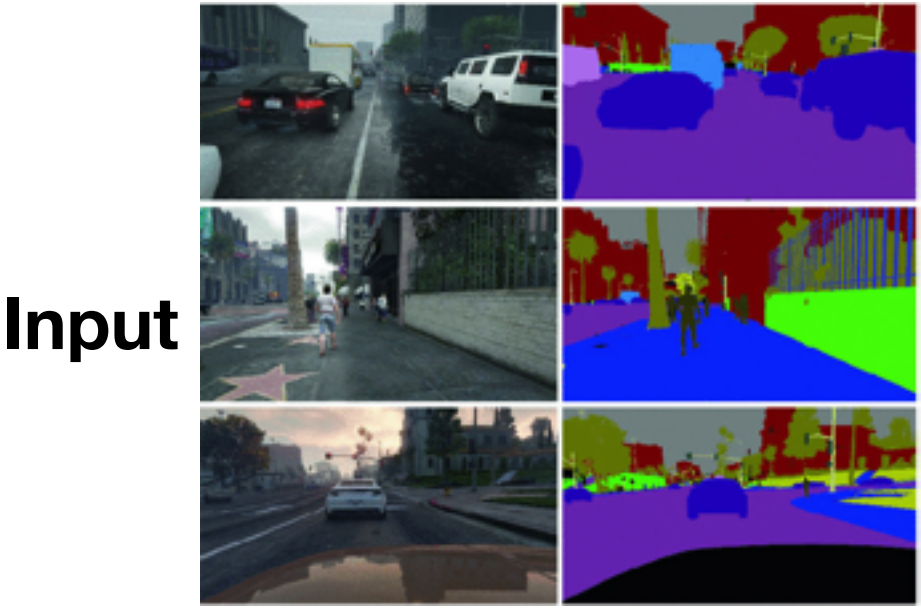

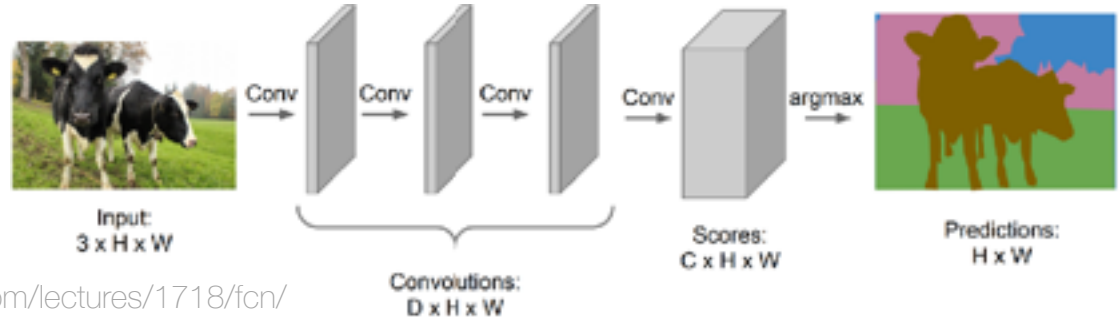

https://tjmachinelearning.com/lectures/1718/fcn/

Popular Semantic Segmentation Datasets

COCO http://cocodataset.org/

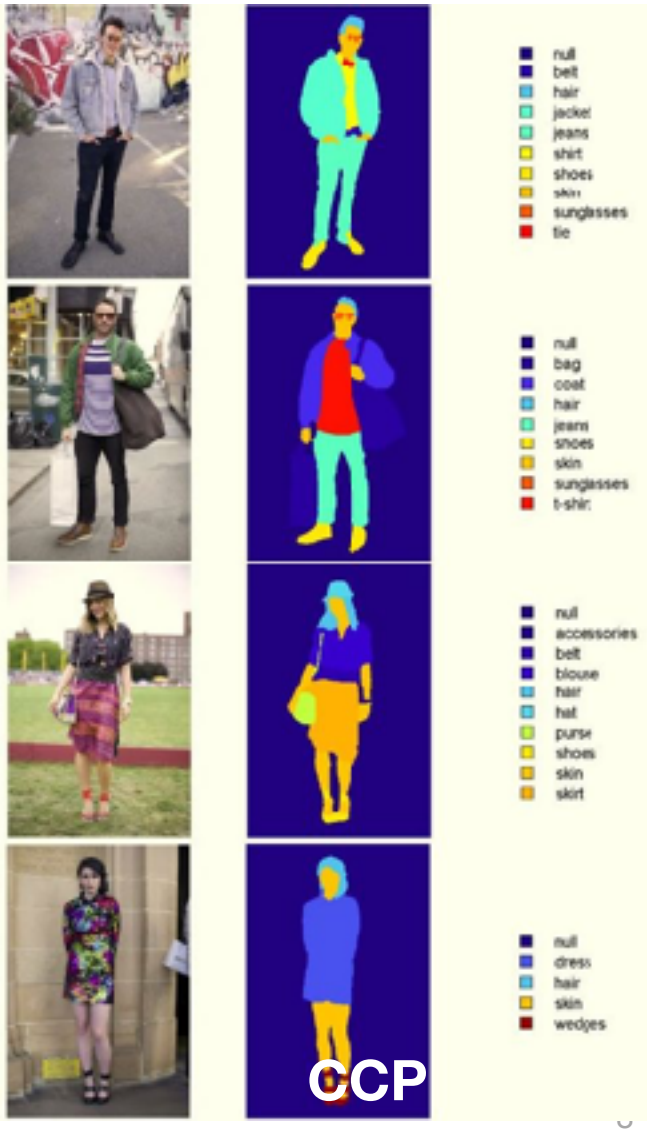

Early Training Methods

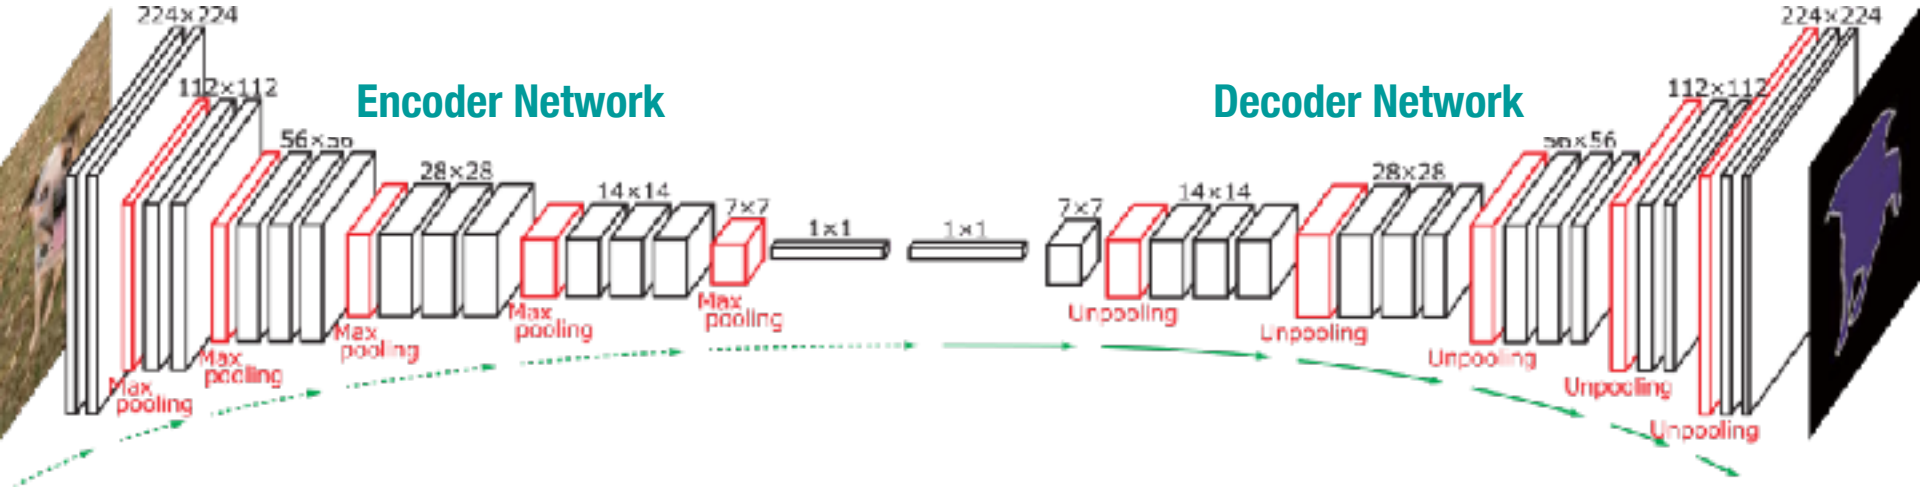

- Init Encoder with traditional CNN (like VGG or DarkNet)
- Freeze encoder and train decoder with segmented image maps
- Unfreeze encoder and fine tune
  - Repeat tuning as needed

https://tjmachinelearning.com/lectures/1718/fcn/

Aside: Integer Upsampling via Interpolation

Nearest Neighbor

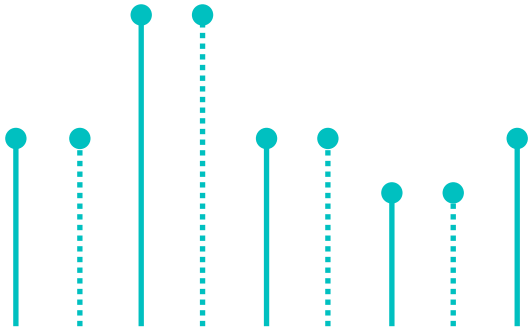

All are equivalent to inserting zeros and applying convolutional filter

Linear

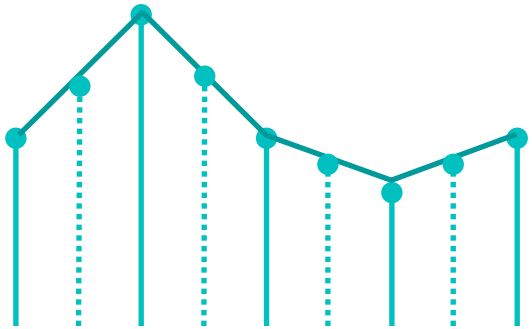

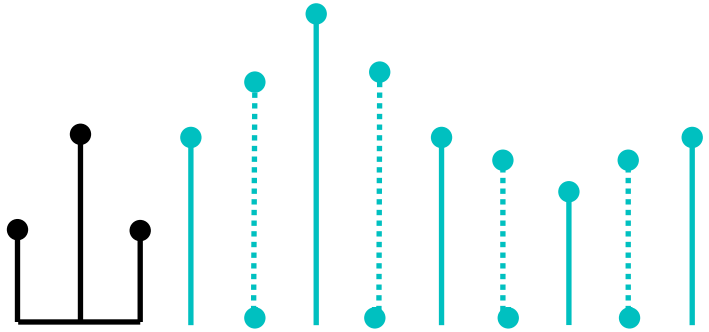

Cubic

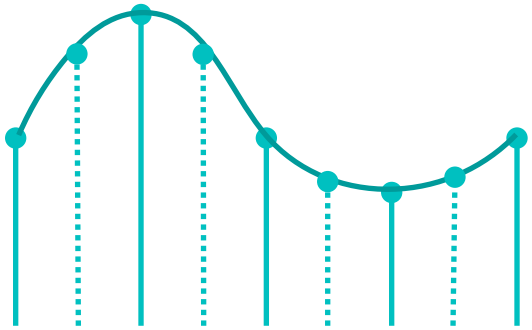

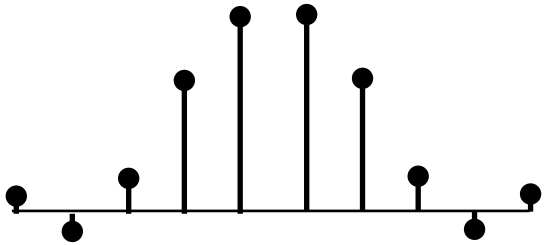

Aside: Image Upsampling, Integer Factor

- Insert Zeros
- Convolve

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

1		2		3		4	
5		6		7		8	
9		10		11		12	
13		14		15		16	

0.25	0.5	0.25
0.5	1	0.5
0.25	0.5	0.25

Bilinear Filtering

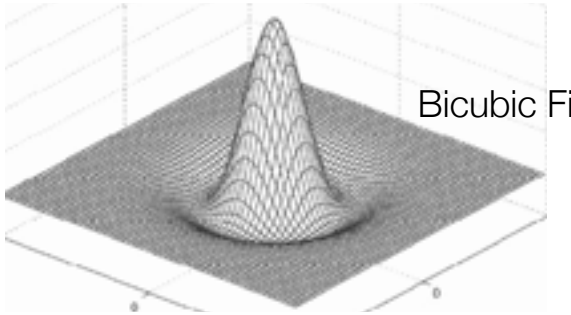

Bicubic Filter

Aside: Image Upsampling, Integer Factor

Nearest Neighbor

Bilinear

Bicubic

https://www.cs.toronto.edu/~guerzhoy/320/lec/upsampling.pdf

Semantic Segmentation

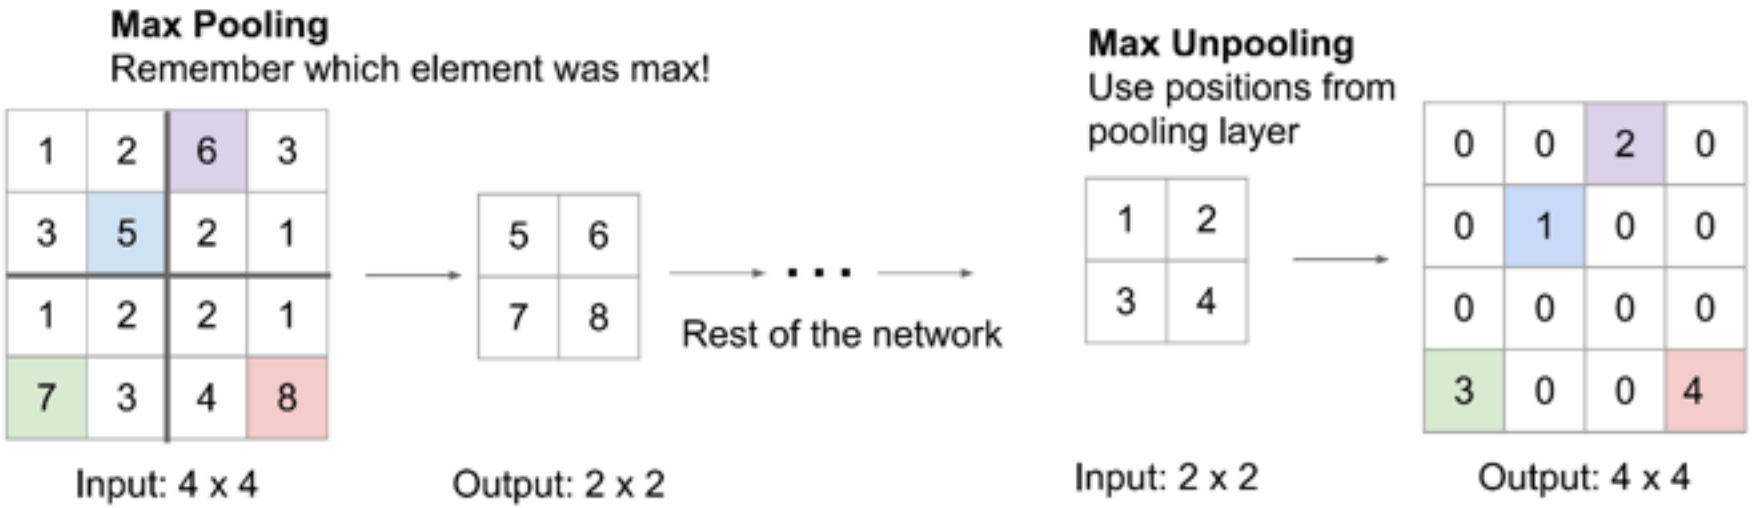

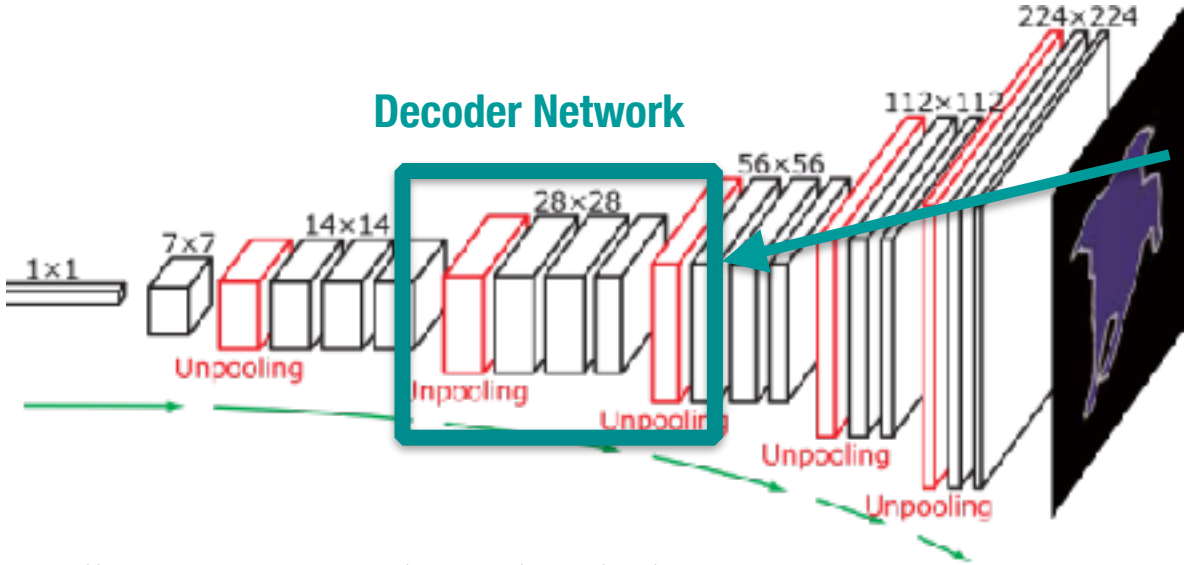

Some knucklehead started calling this **deconvolution**. If you use that term in this class, **you fail**.

This is unpooling and then convolution, but **now the interpolation filters are learned**!!

https://tjmachinelearning.com/lectures/1718/fcn/

What about transpose convolution?

Convolution as Matrix Multiplication

y	x	0	0	0		0	ax
z	y	x	0	0		a	$ay+bx$
0	z	y	x	0	x	b	$az+by+cx$
0	0	z	y	x		c	$bz+cy$
0	0	0	z	y		0	cz

Transpose

y	z	0	0	0		0	az
x	y	z	0	0		a	$ay+bz$
0	x	y	z	0	x	b	$ax+by+cz$
0	0	x	y	z		c	$bx+cy$
0	0	0	x	y		0	cx

like convolving with “reversed coefficients”

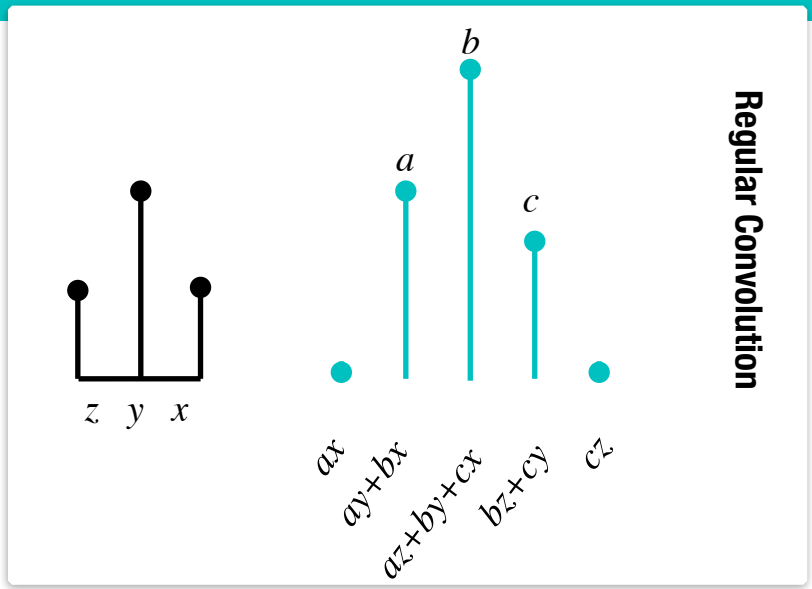

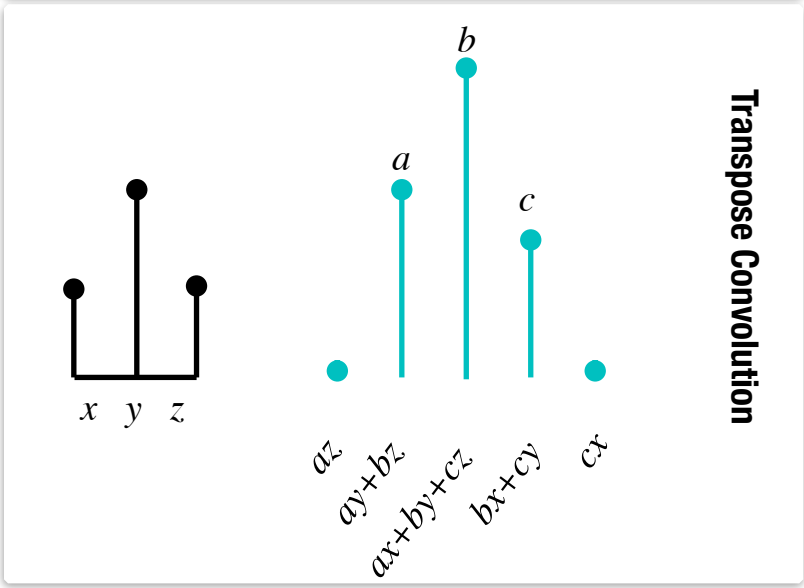

Transpose Convolution: Strides

Strided Convolution as Matrix Multiplication

$$\begin{bmatrix} y & x & 0 & 0 & 0 \\ 0 & z & y & x & 0 \\ 0 & 0 & 0 & z & y \end{bmatrix} \times \begin{bmatrix} 0 \\ a \\ b \\ c \\ 0 \end{bmatrix} = \begin{bmatrix} ax \\ az+by+cx \\ cz \end{bmatrix}$$

Transpose

$$\begin{bmatrix} y & 0 & 0 \\ x & z & 0 \\ 0 & y & 0 \\ 0 & x & z \\ 0 & 0 & y \end{bmatrix} \times \begin{bmatrix} a \\ b \\ c \end{bmatrix} = \begin{bmatrix} ay \\ ax+bz \\ by \\ bx+cz \\ cy \end{bmatrix}$$

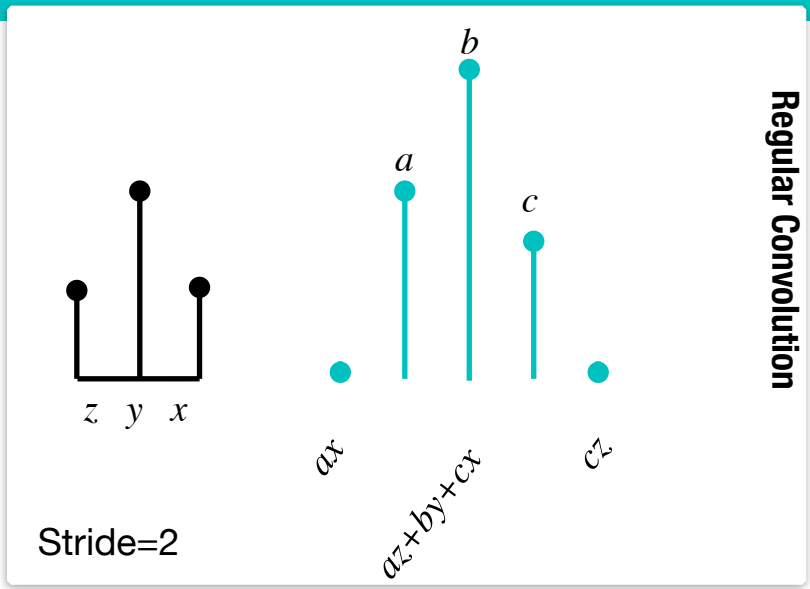

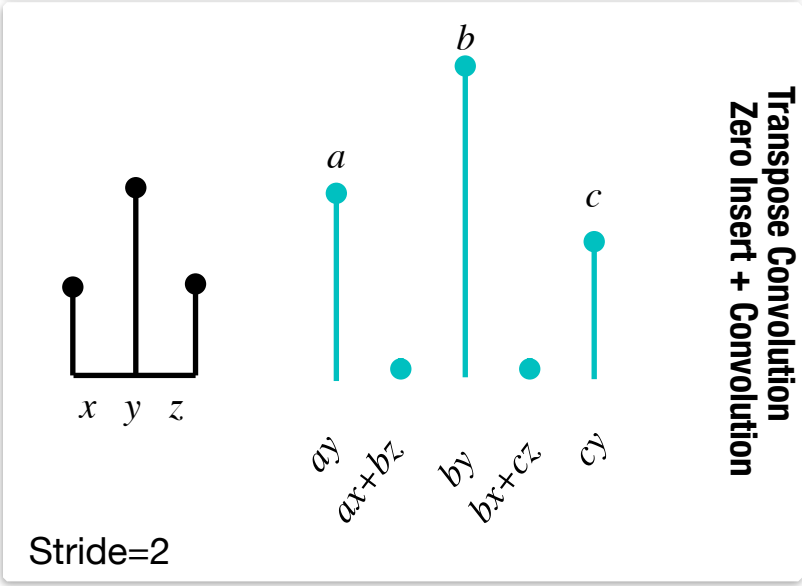

Aside: Convolution after zero insertion

- Kernel size should be a symmetric multiple of the stride

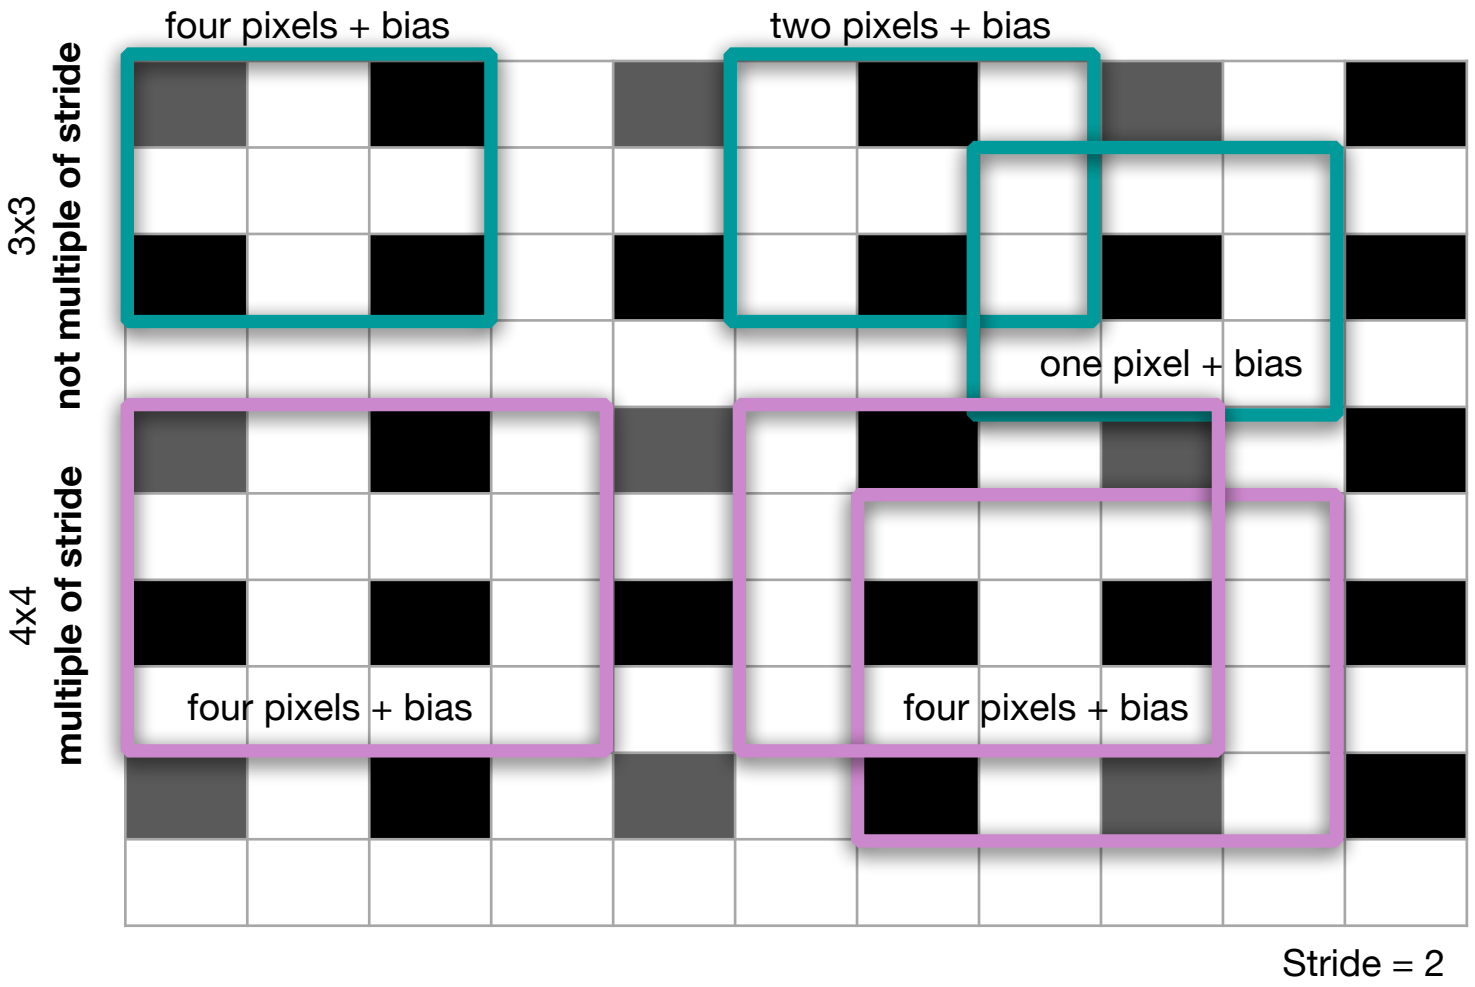

Bias needs to account for both when different numbers of pixels overlap with the kernel

Multiple of stride ensures that same number of active pixels overlap the kernel.