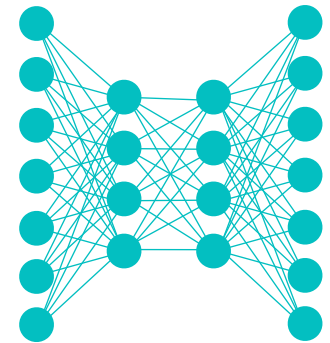


Lecture Notes for **Neural Networks and Machine Learning**



Adaptive, Self-supervised,
Multi-modal, & Multi-task
Learning



Logistics and Agenda

- Logistics
 - Lab three uses multi-task and multi-modal learning
- Agenda
 - Adaptive Learning
 - Self-Supervised Learning
 - Paper Presentation
 - Multi-modal/task Learning
 - ◆ Techniques
 - ◆ Applications and domains
- Next Time:
 - Paper Presentation: Speaker Verification with X-Vectors and SincNet



Paper Presentation: The Lottery Hypothesis

Published as a conference paper at ICLR 2019

THE LOTTERY TICKET HYPOTHESIS: FINDING SPARSE, TRAINABLE NEURAL NETWORKS

Jonathan Frankle
MIT CSAIL
jfrankle@csail.mit.edu

Michael Carbin
MIT CSAIL
mcarbin@csail.mit.edu

ABSTRACT

Neural network pruning techniques can reduce the parameter counts of trained networks by over 90%, decreasing storage requirements and improving computational performance of inference without compromising accuracy. However, contemporary experience is that the sparse architectures produced by pruning are difficult to train from the start, which would similarly improve training performance.

We find that a standard pruning technique naturally uncovers subnetworks whose initializations made them capable of training effectively. Based on these results, we articulate the *lottery ticket hypothesis*: dense, randomly-initialized, feed-forward networks contain subnetworks (*winning tickets*) that—when trained in isolation—reach test accuracy comparable to the original network in a similar number of iterations. The winning tickets we find have won the initialization lottery: their connections have initial weights that make training particularly effective.

We present an algorithm to identify winning tickets and a series of experiments that support the lottery ticket hypothesis and the importance of these fortuitous initializations. We consistently find winning tickets that are less than 10-20% of the size of several fully-connected and convolutional feed-forward architectures for MNIST and CIFAR10. Above this size, the winning tickets that we find learn faster than the original network and reach higher test accuracy.



Last Time

$$X = x_1, x_2, \dots, x_N \in \mathcal{X}$$

$$Y = y_1, y_2, \dots, y_N \in \mathcal{Y}$$

$$\mathcal{D} = \{\mathcal{X}, p(X)\}$$

Domain Feature Space Probability Observation

- Domain defines the features used
- Marginal Distribution of observing instances in the feature space
 - Typically intractable to calculate (generative)

$$\mathcal{T} = \{\mathcal{Y}, p(Y|X)\}$$

Task Label Space Learned Probability

- Task is within a domain
- Label space is typically one specific classification or regression task
- Probability of observing label given the feature space:
 - Not intractable (discriminative)

	Training		Testing	
Transfer Learning	Task 1		Task 2	
Multi-task Learning	Task 1	... Task N	Task 1	... Task N
Lifelong Learning	Task 1	... Task N	Task N+1	

Humans can learn to ride a bike and use that to understand better about driving a car. Machine Learning in its current form is far from this capability. How can we move our stoned version of artificial intelligence closer to the process of human based learning? How can we accumulate knowledge from model to model?

Does biology of human learning hold any clues to success? How does a human learn to crawl? To talk? To ride a bike? What is a human's motivation to learn?

Feature Extraction Transfer

- Most well known: use learned parameters from one task in another task in same domain
- Most useful when labels for target domain are sparse



Ian Goodfellow's Definition:

"Transfer learning refers to any situation where what has been learned in one setting is exploited to improve generalization in another setting."



Active Transfer Learning


Theory:



Practice:



Machine Learning :

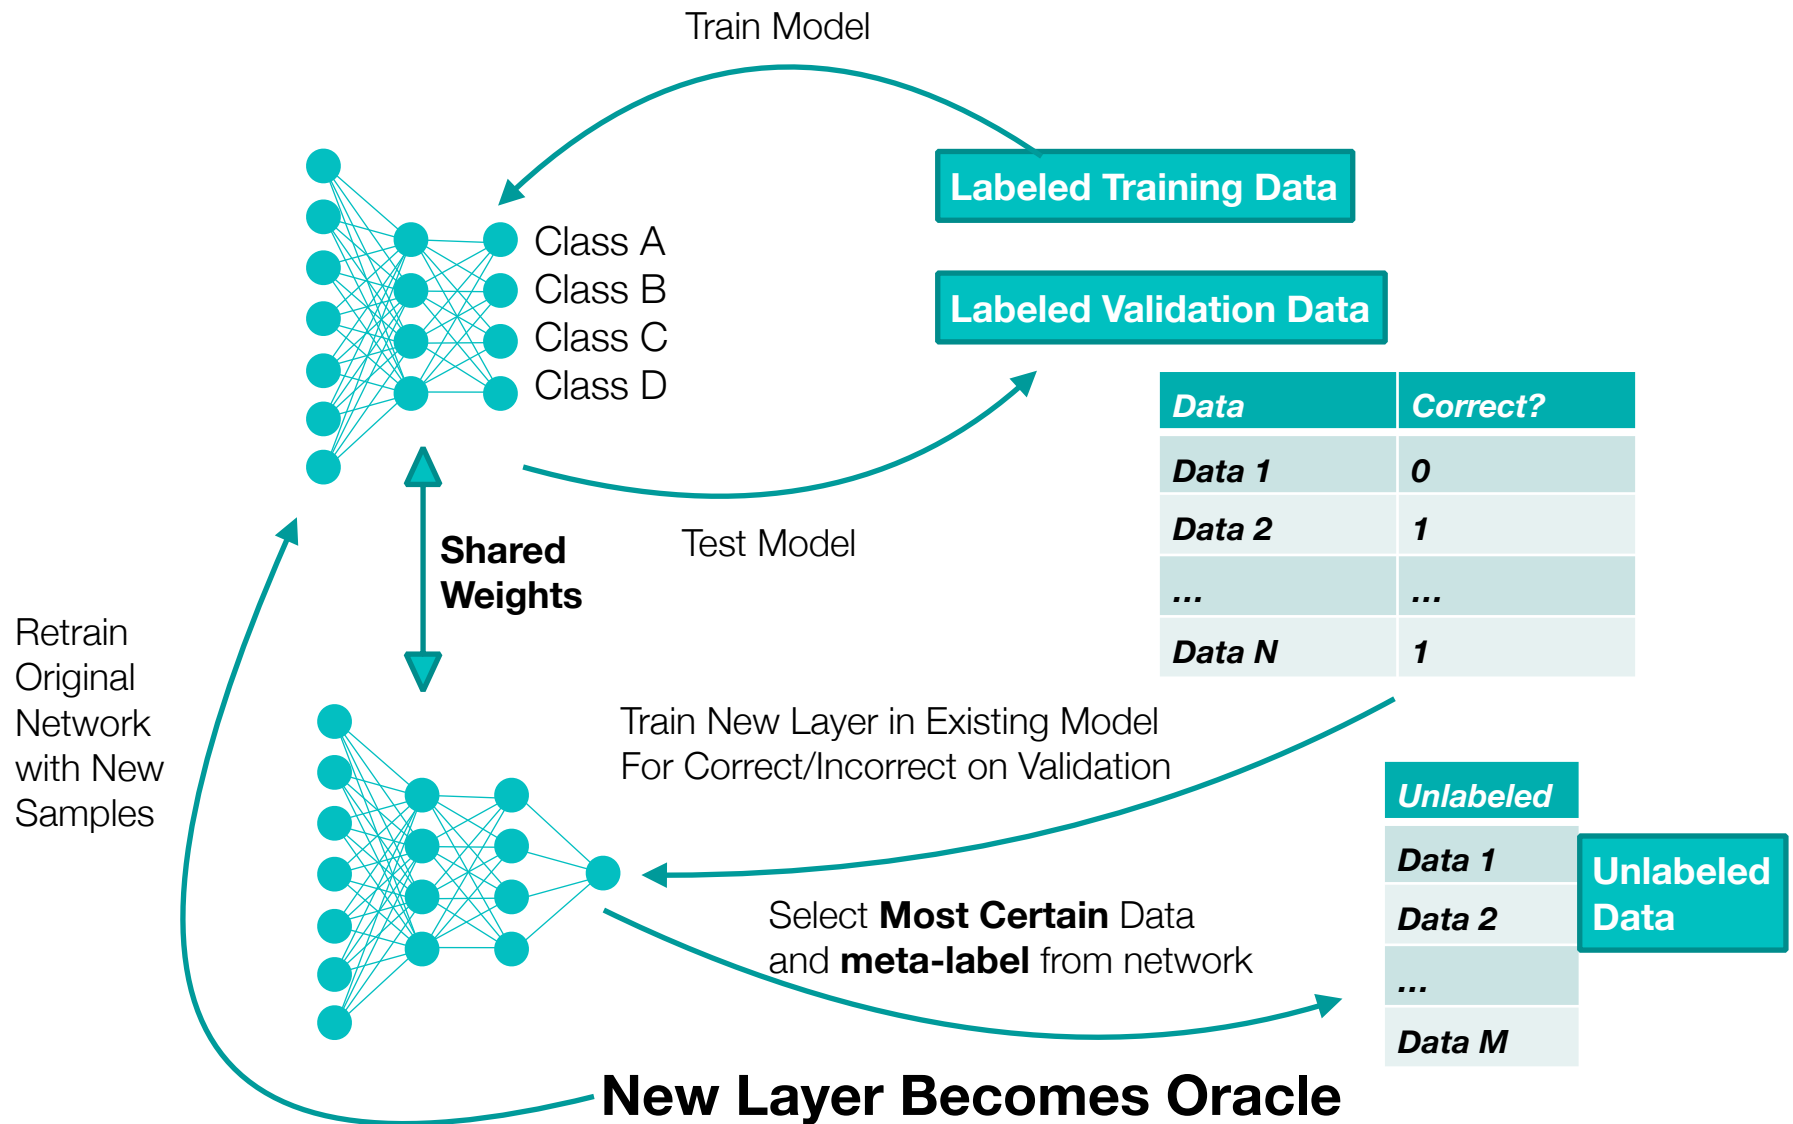


Active Learning Overview

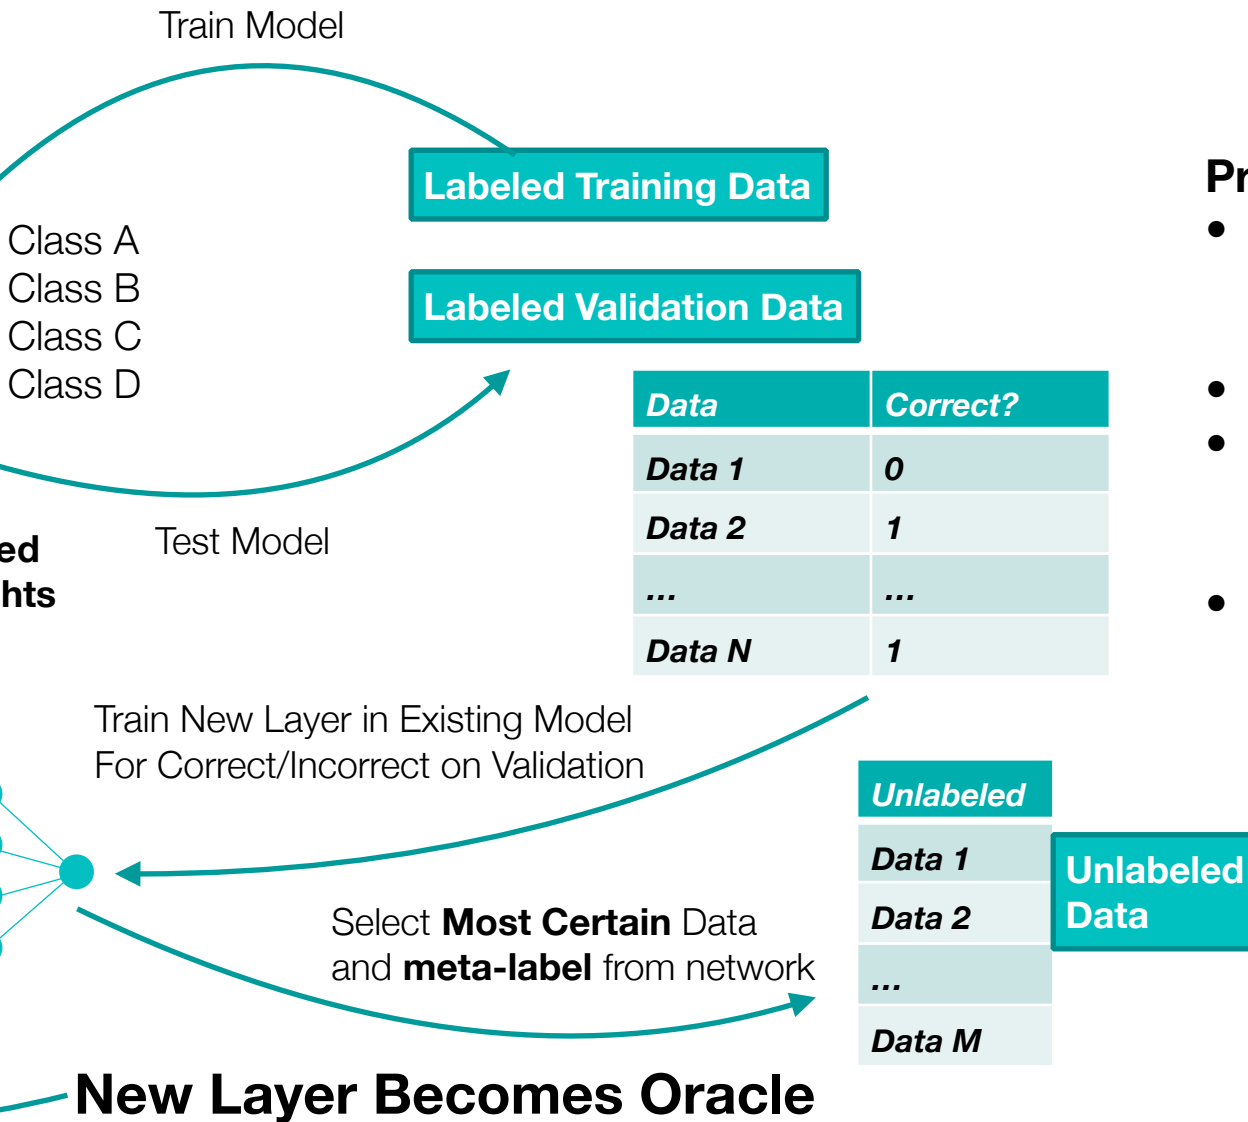
- **Basic Idea:** Use a trained model to sample from an oracle that can magically give you a new label
 - We are asking:
What labels should we ask the oracle about?
- Uncertainty Sampling
 - Choose instances where the model is most uncertain or most certain
 - Various ways to measure certainty
- Diversity Sampling
 - Choose instances that are similar or different from training distribution



Uncertainty Sampling with a Neural Network



Uncertainty Sampling with a Neural Network

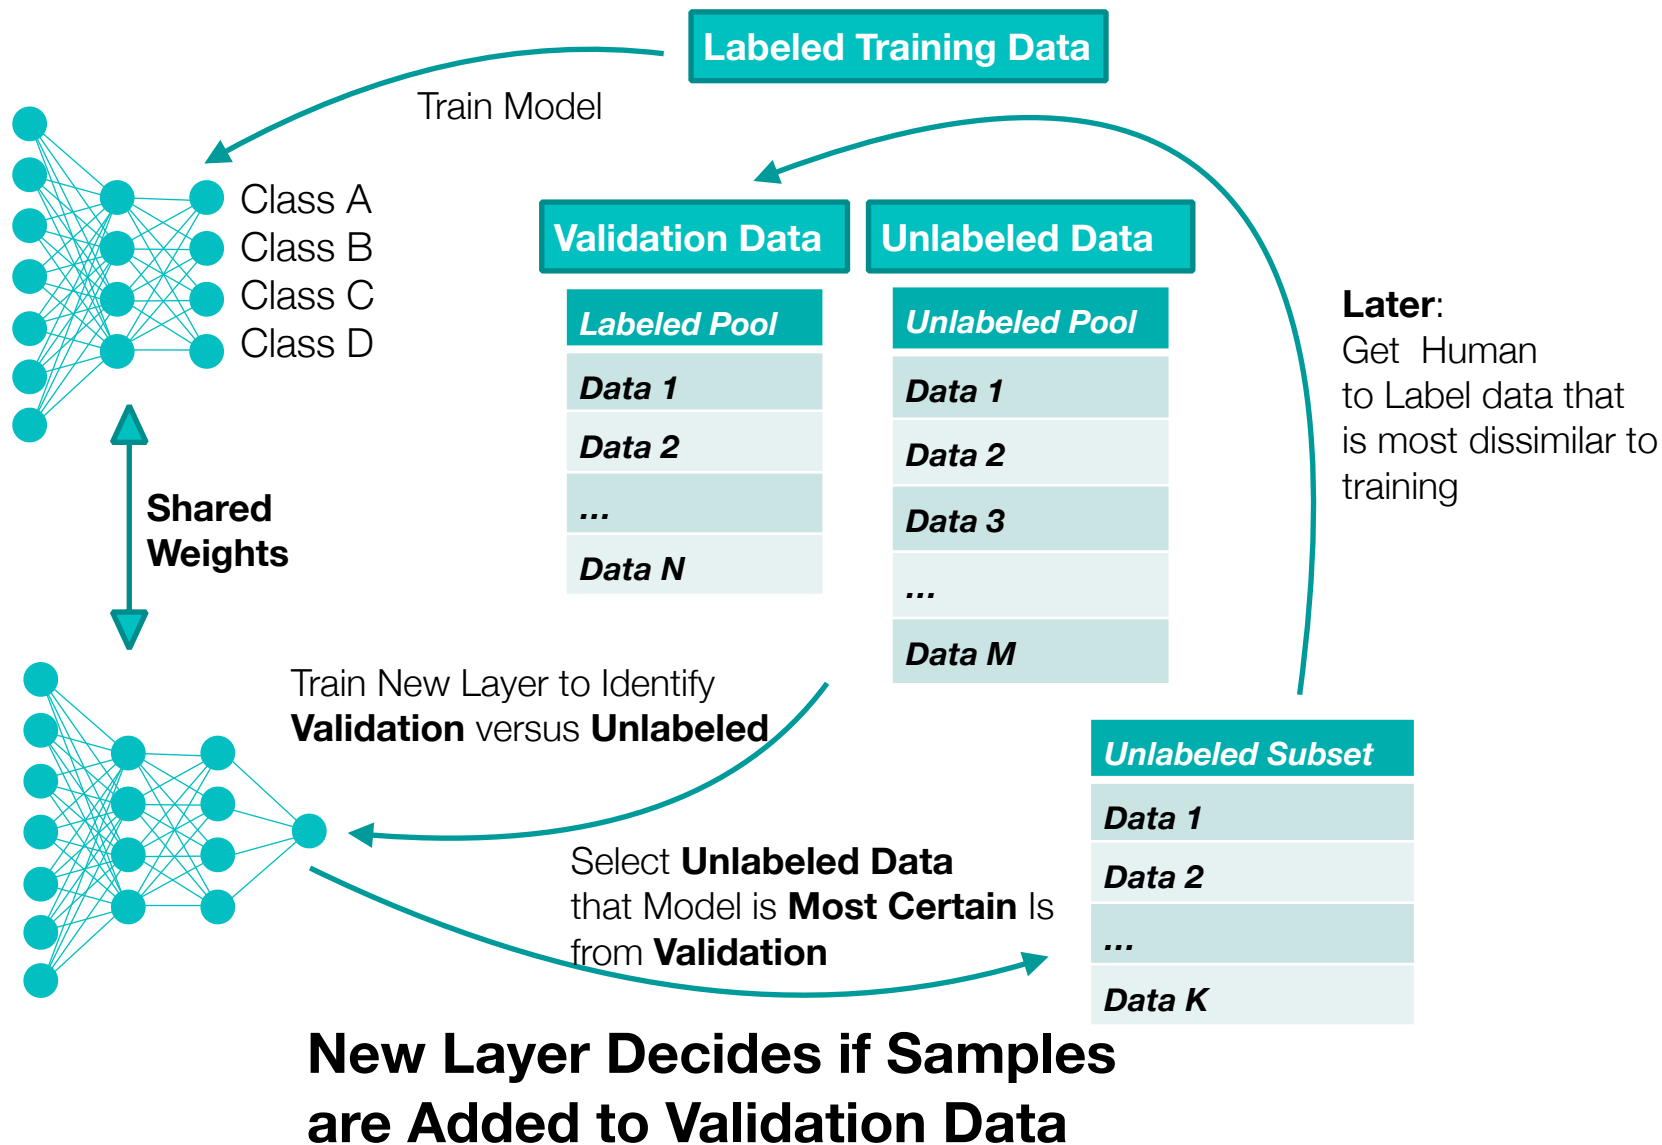


Problems:

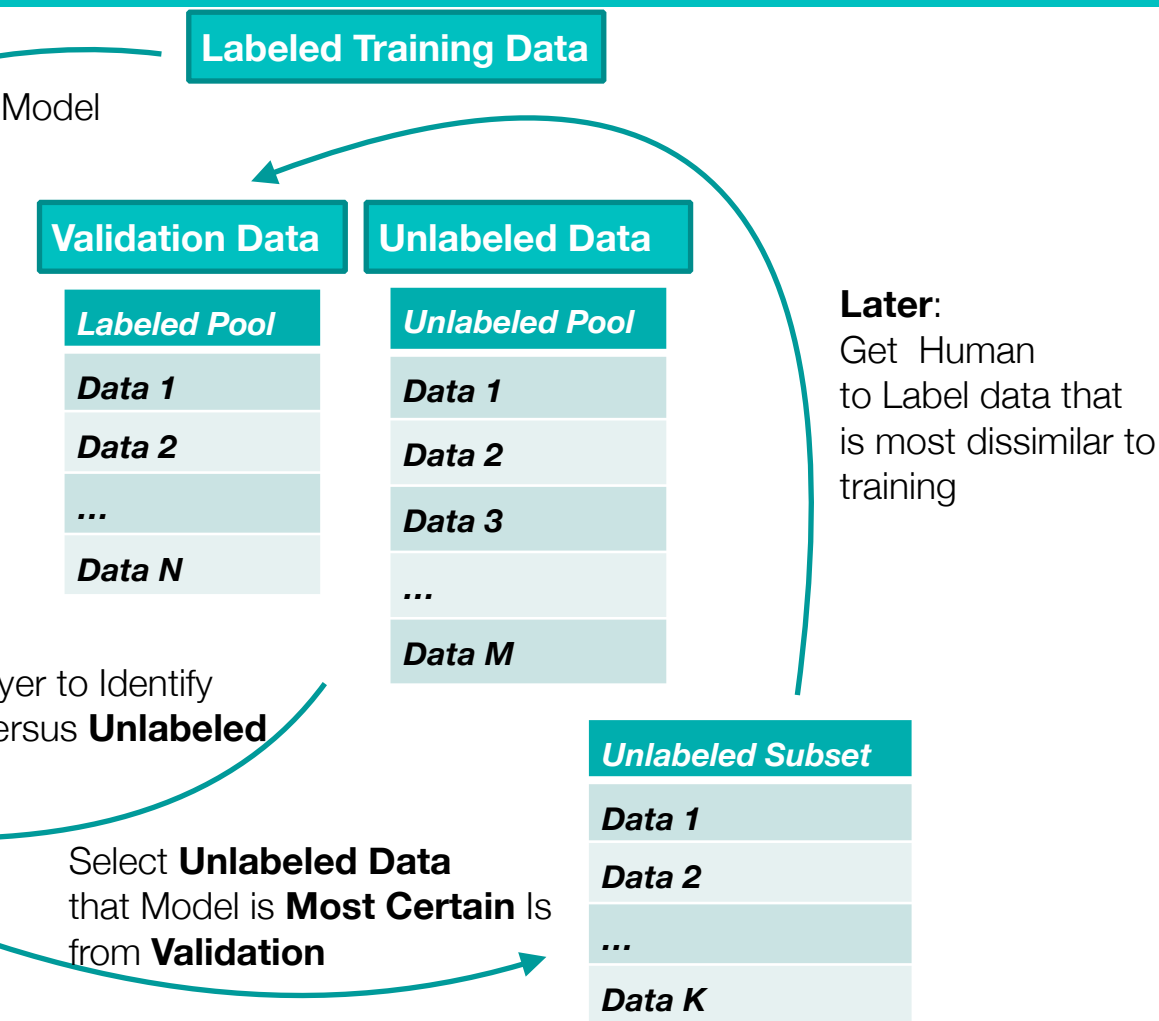
- Training pool is represented by classes the model already does well predicting
- Limited diversity of Samples
- Training pool can become contaminated easily from a few wrong predictions
- For Oracle: we might be asking to get labels that the model is already good at classifying



Diversity Sampling with a Neural Network



Diversity Sampling with a Neural Network



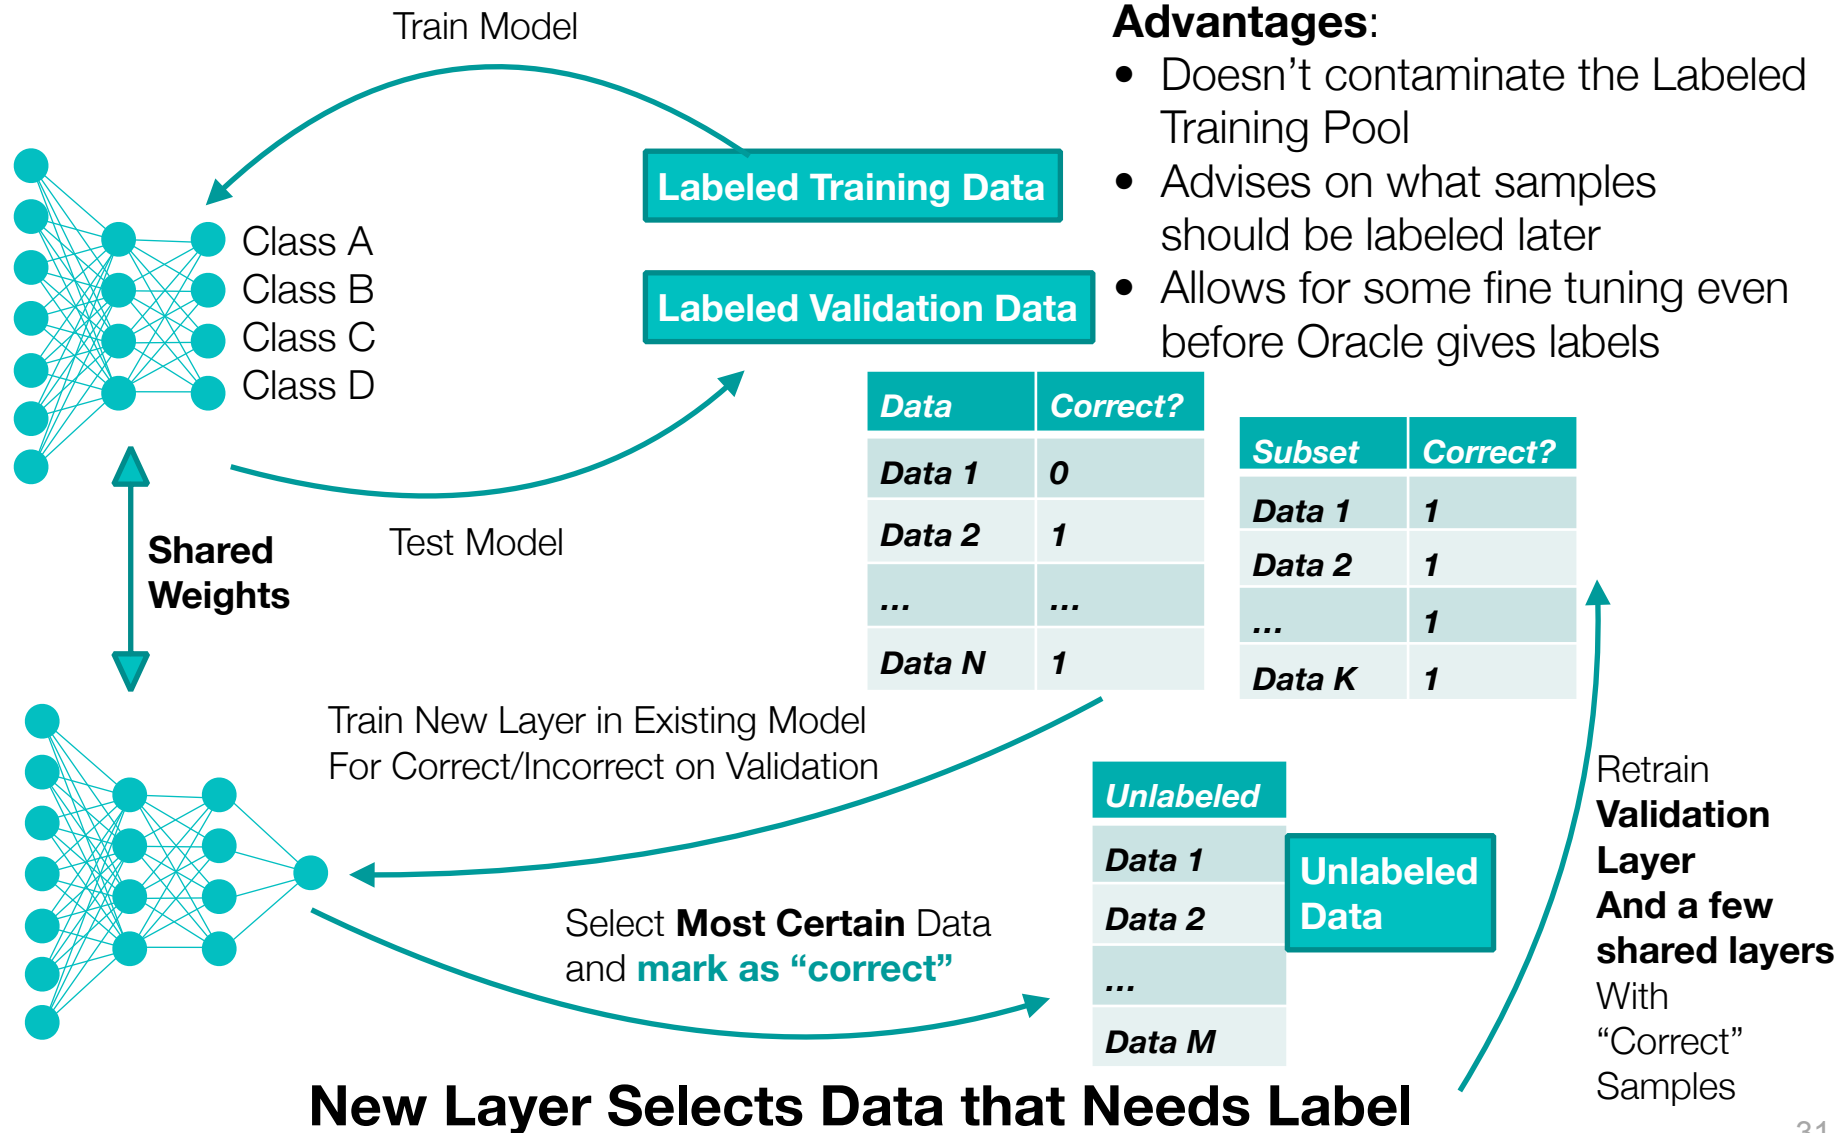
Discussion:

- Training pool is not contaminated
- Expands validation data in well mannered way, not adding too “far away” samples
- Validation versus Unlabeled might not be the best comparison, because it ignores confusions in the training data
- For Oracle: we can get labels to inputs that the model is likely to be unsure about
- But... this only helps us when we have an Oracle to give us labels

ayer Decides if Samples
led to Validation Data



ATLAS: Active Transfer Learning for Adaptive Sampling



Time Period	Protocol	Expected Feedback
First Week	<p>Homeowner provides 8-20 examples over the first week:</p> <ul style="list-style-type: none"> • 1-2 Shower usages • 1 run of the dishwasher • 1 run of the laundry machine • 2 examples of each toilet • 1 example of hot and cold water use for each dual handle faucet • 1 example of hot, cold, and mixed water use for each single handle faucet (2 examples if in kitchen) 	<p>HydroSense relies on the rule based classifier for the first week.</p> <p>Pressure waves are saved in order to create a sparse codebook of features.</p> <p>Results are displayed at the fixture category for dishwashers, showers, and washing machines.</p>
Start of Second Week	Homeowner provides 2-4 labels every other day when the system messages them on their mobile device	<p>Results are displayed at the full fixture category level from the CoDBN-VE algorithm. Expected accuracy:</p> <ul style="list-style-type: none"> • 85% at fixture category level
End of Second Week	Homeowner has supplied 9-12 examples that were flagged by active learning.	<p>HydroSense now displays results at the Lumped Fixture level.</p> <p>Expected accuracy:</p> <ul style="list-style-type: none"> • 82% at fixture level • 87% at fixture category level
End of Third Week	Homeowner continues to supply sparsely selected examples every other day. About 9-12 additional examples provided.	<p>Valve level accuracy now provided.</p> <p>Expected accuracy:</p> <ul style="list-style-type: none"> • 80% at valve level • 87% at fixture level • 92% at fixture category level
Fourth Week	Homeowner can optionally continue to provide examples to the system for increased accuracy.	<p>Expected accuracy:</p> <ul style="list-style-type: none"> • 81% at valve level • 89% at fixture level • 93% at fixture category level

Table 8-2. Expected feedback and calibration protocol for semi-supervised HydroSense system

Self-Supervised Learning

From
Yoshua Bengio

Three challenges for Deep Learning

- ▶ Deep Supervised Learning works well for perception
 - ▶ When labeled data is abundant.
- ▶ Deep Reinforcement Learning works well for action generation
 - ▶ When trials are cheap, e.g. in simulation.

Three problems the community is working on:

1. Learning with fewer labeled samples and/or fewer trials
 - ▶ Self-supervised learning / unsup learning / learning to fill in the blanks
 - ▶ learning to represent the world before learning tasks
2. Learning to reason, beyond "system 1" feed forward computation
 - ▶ Making reasoning compatible with gradient-based learning.
3. Learning to plan complex action sequences
 - ▶ Learning hierarchical representations of action plans

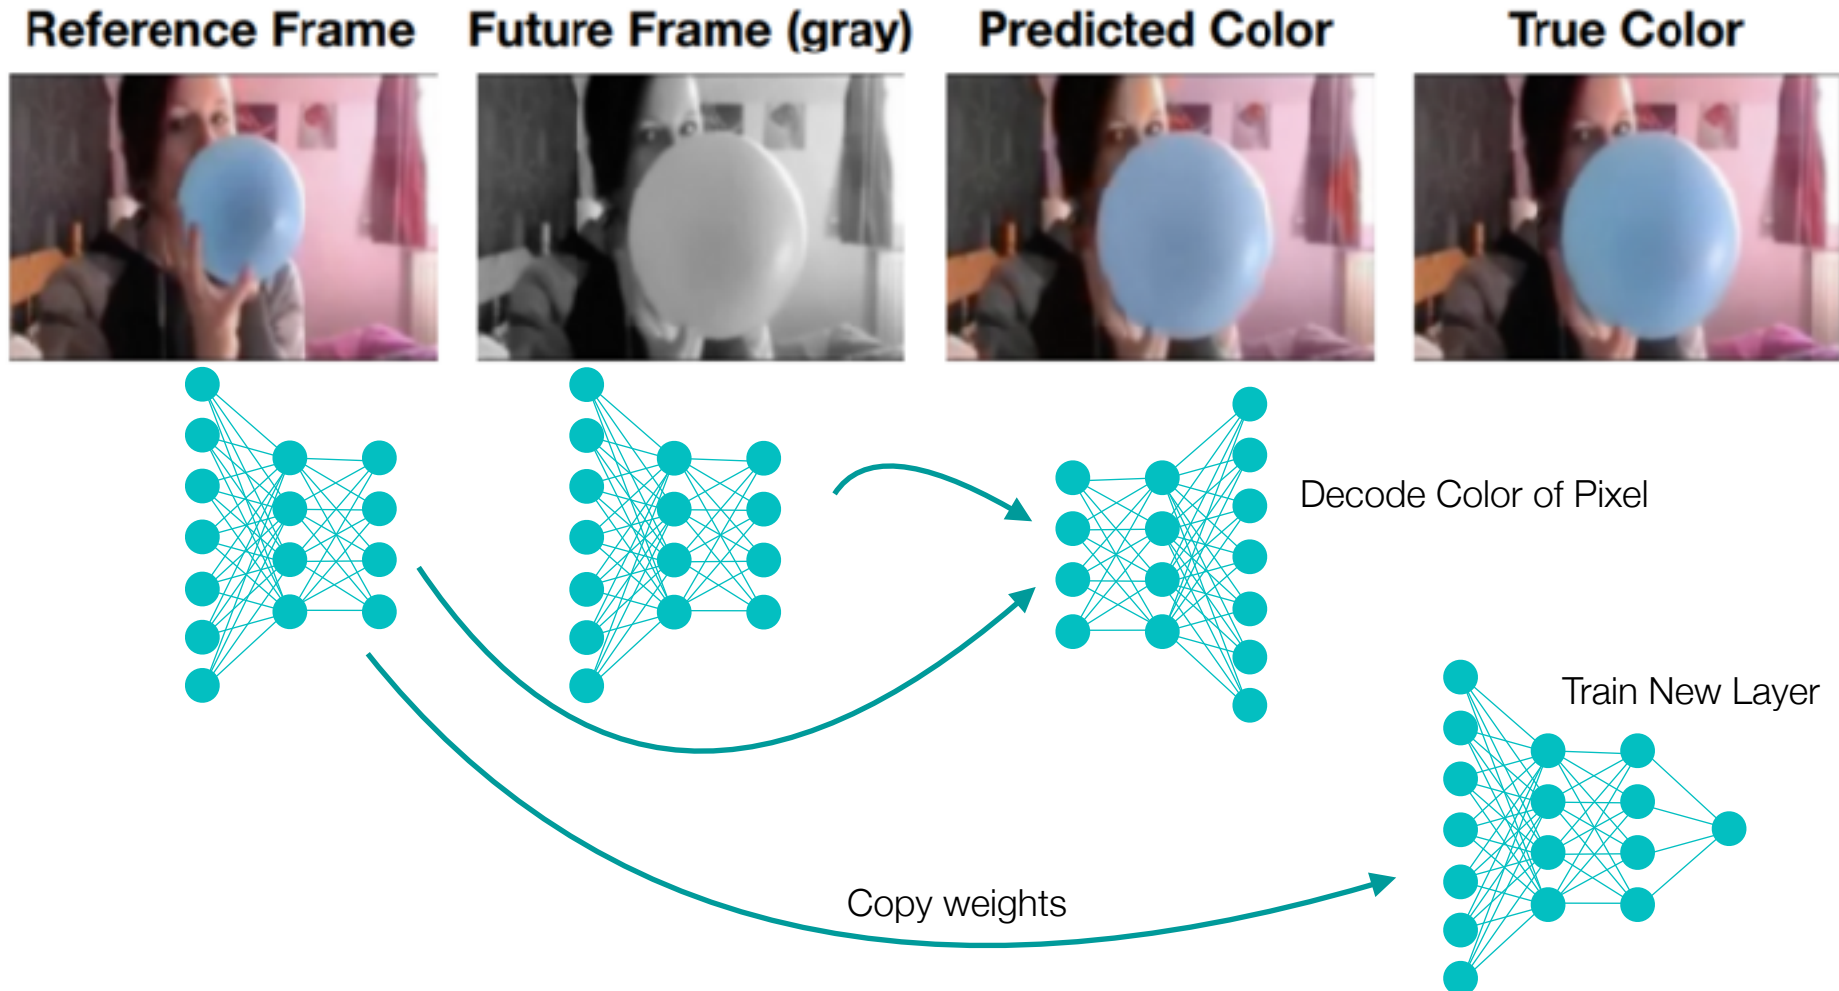


Self-supervised Learning

- **Problem:** deep learning is not sample efficient
- **Idea:** learn about the world before learning the task
- **New Problem:** how do we learn about the world?
- **Solution:** transfer learning on toy problem
 - 1. train on auxiliary task that is easy to label
 - 2. throw away anything specific to auxiliary task
 - 3. train new network with task of interest, transferring knowledge (downstream task)
 - 4. profit



Examples of Self Supervised Learning

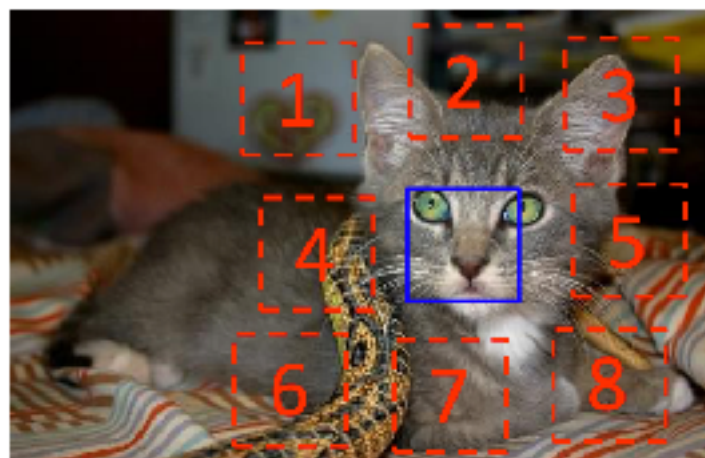


https://www.fast.ai/2020/01/13/self_supervised/

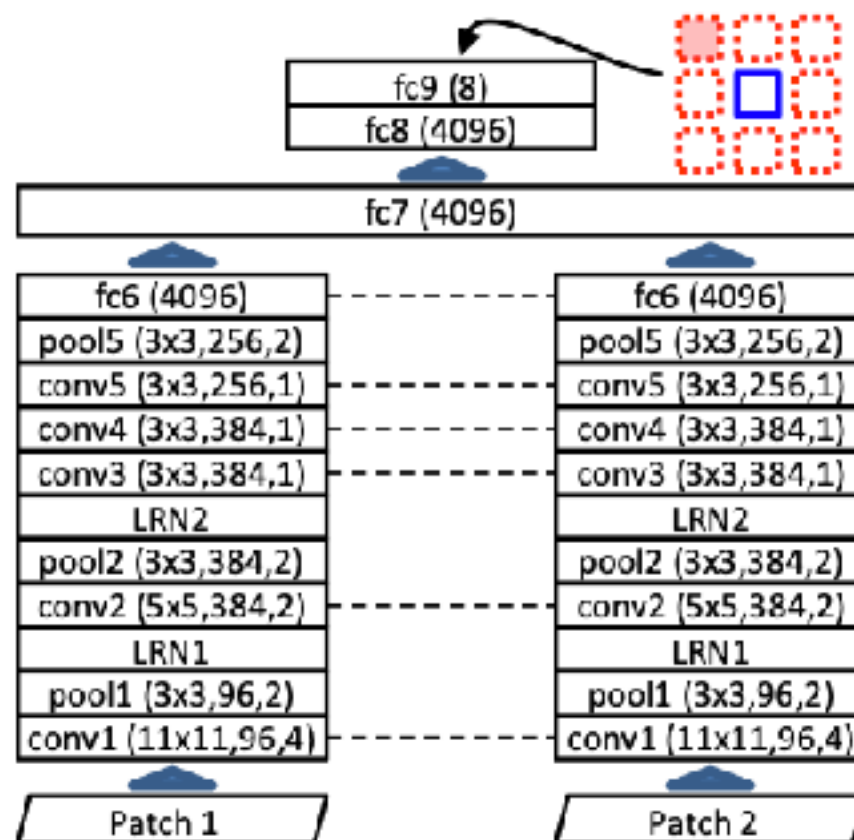
35



Examples of Self Supervised Learning



$$X = \left(\begin{array}{c} \text{cat face patch} \\ \text{cat ear patch} \end{array} \right); Y = 3$$



Unsupervised Visual Representation Learning by Context Prediction

Carl Doersch^{1,2}

Abhinav Gupta¹

Alexei A. Efros²

¹ School of Computer Science
Carnegie Mellon University

² Dept. of Electrical Engineering and Computer Science
University of California, Berkeley



Examples of SSL

Shuffle and Learn: Unsupervised Learning using Temporal Order Verification

Ishan Misra¹ C. Lawrence Zitnick² Martial Hebert¹


¹ The Robotics Institute, Carnegie Mellon University

² Facebook AI Research



Table 2: Mean classification accuracies over the 3 splits of UCF101 and HMDB51 datasets. We compare different initializations and finetune them for action recognition.


Dataset	Initialization	Mean Accuracy
UCF101	Random	38.6
	(Ours) Tuple verification	50.2
HMDB51	Random	13.3
	UCF Supervised	15.2
	(Ours) Tuple verification	18.1



Examples of Self Supervised Learning



Pretraining Method	Supervision	Pretraining time	Classification	Detection	Segmentation
ImageNet [26]	1000 class labels	3 days	78.2%	56.8%	48.0%
Random Gaussian	initialization	< 1 minute	53.3%	43.4%	19.8%
Autoencoder	-	14 hours	53.8%	41.9%	25.2%
Agrawal <i>et al.</i> [1]	egomotion	10 hours	52.9%	41.8%	-
Wang <i>et al.</i> [39]	motion	1 week	58.7%	47.4%	-
Doersch <i>et al.</i> [7]	relative context	4 weeks	55.3%	46.6%	-
Ours	context	14 hours	56.5%	44.5%	30.0%



Context Encoders: Feature Learning by Inpainting

Deepak Pathak

Philipp Krähenbühl

Jeff Donahue

Trevor Darrell

Alexei A. Efros

https://www.fast.ai/2020/01/13/self_supervised/

University of California, Berkeley

38



Unsupervised Consistency Loss

$$\min_{\mathbf{w}} \underbrace{\mathbf{E}_{\mathbf{x}, y \in L} [-\log p_{\mathbf{w}}(y | \mathbf{x})]}_{\text{cross entropy}} + \lambda \underbrace{\mathbf{E}_{\mathbf{x} \in U} \mathbf{E}_{\hat{\mathbf{x}} \leftarrow q(\hat{\mathbf{x}} | \mathbf{x})} \left[\mathcal{D}_{KL} (p_{\mathbf{w}}(y | \mathbf{x}) || p_{\mathbf{w}}(y | \hat{\mathbf{x}})) \right]}_{\text{consistency in augmentation}}$$


no back prop yes back prop

Neural Network approximates $p(y|\mathbf{x})$ by \mathbf{w}
Use labeled data to minimize network

Sample new \mathbf{x} from unlabeled pool with function q
function q is augmentation procedure
Minimize cross entropy of two models

**Get accustomed
to this notation**

**Update Model with
Back-propagation along these paths**



Unsupervised Consistency Loss

$$\min_{\mathbf{w}} \underbrace{\mathbf{E}_{\mathbf{x}, y \in L} [-\log p_{\mathbf{w}}(y | \mathbf{x})]}_{\text{cross entropy}} + \lambda \underbrace{\mathbf{E}_{\mathbf{x} \in U} \mathbf{E}_{\hat{\mathbf{x}} \leftarrow q(\hat{\mathbf{x}} | \mathbf{x})} \left[\mathcal{D}_{KL} (p_{\mathbf{w}}(y | \mathbf{x}) || p_{\mathbf{w}}(y | \hat{\mathbf{x}})) \right]}_{\text{consistency in augmentation}}$$




Figure 2: Augmented examples using back-translation and RandAugment.




Unsupervised Consistency Loss

Augmentation (# Sup examples)	Sup (50k)	Semi-Sup (4k)
Crop & flip	5.36	16.17
Cutout	4.42	6.42
RandAugment	4.23	5.29


Table 1: Error rates on CIFAR-10.

Augmentation (# Sup examples)	Sup (650k)	Semi-sup (2.5k)
\times	38.36	50.80
Switchout	37.24	43.38
Back-translation	36.71	41.35

Table 2: Error rate on Yelp-5.



(a) CIFAR-10



(b) SVHN



Unsupervised Consistency Loss

Method	Model	# Param	CIFAR-10 (4k)	SVHN (1k)
II-Model (Laine & Aila, 2016)	Conv-Large	3.1M	12.36 ± 0.31	4.82 ± 0.17
Mean Teacher (Tarvainen & Valpola, 2017)	Conv-Large	3.1M	12.31 ± 0.28	3.95 ± 0.19
VAT + EntMin (Miyato et al., 2018)	Conv-Large	3.1M	10.55 ± 0.05	3.86 ± 0.11
SNTG (Luo et al., 2018)	Conv-Large	3.1M	10.93 ± 0.14	3.86 ± 0.27
VAdD (Park et al., 2018)	Conv-Large	3.1M	11.32 ± 0.11	4.16 ± 0.08
Fast-SWA (Athiwaratkun et al., 2018)	Conv-Large	3.1M	9.05	-
ICT (Verma et al., 2019)	Conv-Large	3.1M	7.29 ± 0.02	3.89 ± 0.04
Pseudo-Label (Lee, 2013)	WRN-28-2	1.5M	16.21 ± 0.11	7.62 ± 0.29
LGA + VAT (Jackson & Schulman, 2019)	WRN-28-2	1.5M	12.06 ± 0.19	6.58 ± 0.36
mixmixup (Hataya & Nakayama, 2019)	WRN-28-2	1.5M	10	-
ICT (Verma et al., 2019)	WRN-28-2	1.5M	7.66 ± 0.17	3.53 ± 0.07
MixMatch (Berthelot et al., 2019)	WRN-28-2	1.5M	6.24 ± 0.06	2.89 ± 0.06

Methods	SSL	10%	100%
ResNet-50	✗	55.09 / 77.26	77.28 / 93.73
w. RandAugment		58.84 / 80.56	78.43 / 94.37
UDA (RandAugment)	✓	68.78 / 88.80	79.05 / 94.49

Table 5: Top-1 / top-5 accuracy on ImageNet with 10% and 100% of the labeled set. We use image size 224 and 331 for the 10% and 100% experiments respectively.



Lecture Notes for **Neural Networks and Machine Learning**

Ada, SSL,

Next Time:
M-Modal/task
Reading: Papers

