

Lecture Notes for **Neural Networks and Machine Learning**



Neural Style Transfer
Model Optimization



Logistics and Agenda

- Logistics
 - Next Assignment: Style Transfer
- Agenda
 - *A History of Style Transfer (last time)*
 - *Image Optimization Algorithms (last time)*
 - Student Paper Presentation
 - Model Optimization Algorithms (today)
 - One Shot Algorithms (partially today)
 - Evaluating Style Transfer Performance
 - Extensions in Other Domains



Paper Presentation

Perceptual Losses for Real-Time Style Transfer and Super-Resolution

Justin Johnson, Alexandre Alahi, Li Fei-Fei
{jcjohns, alahi, feifeili}@cs.stanford.edu

Department of Computer Science, Stanford University

Abstract. We consider image transformation problems, where an input image is transformed into an output image. Recent methods for such



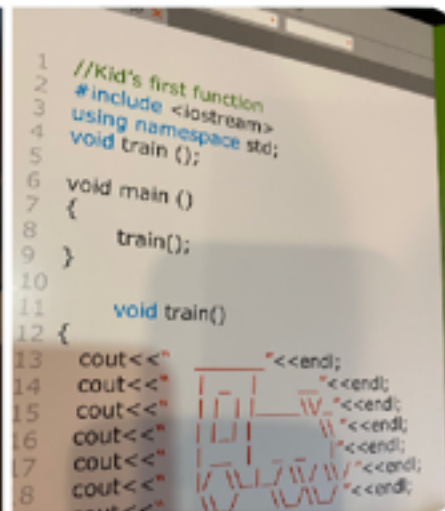

Model Optimization Based Style Transfer



zach lieberman
@zachlieberman



C++ for kids



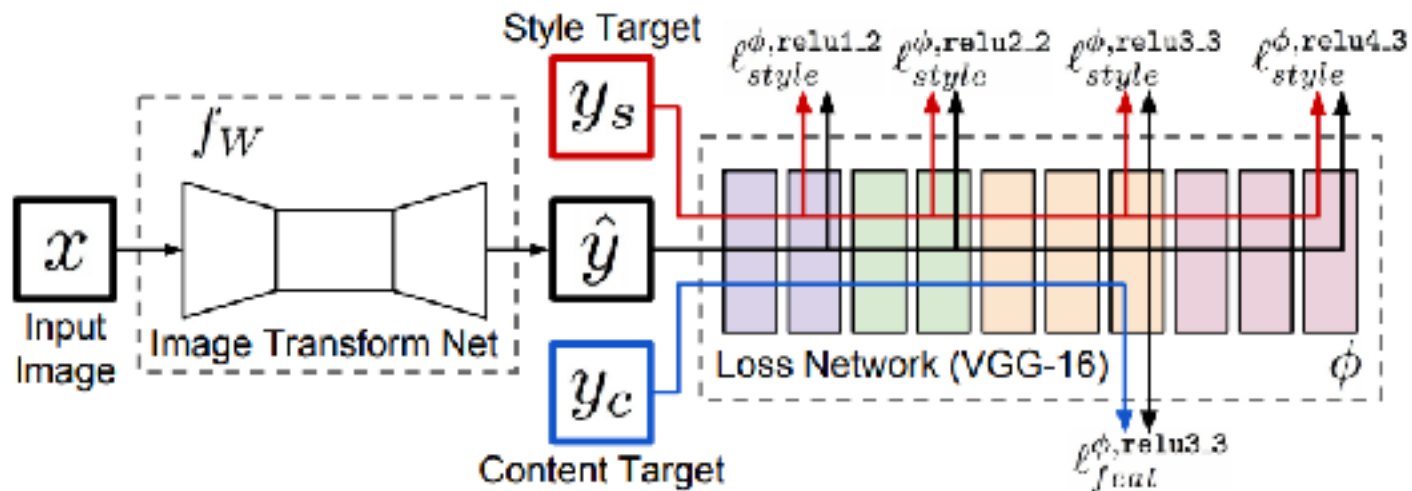
One other thought!

- I hate the term **perceptual loss**!
- It's marketing! Joy!
- ...but not good science verbiage
- ...like the term global warming
 - it can introduce subjective opinion, misunderstanding through misleading labeling
- There is nothing perceptual here, its a neural network
- A better description of the loss:
 - Convolutional Gram Loss?
 - Information Distillation Covariance?
 - Weighted Grammian Norm (WGN Loss)?
- Just don't say "perceptual" in this class (so you don't fail)




Johnson Paper Recap (if needed)

- **Basic Idea:** replace image optimization with single pass, fully convolutional network to perform the transformation
- Loss functions stay identical to Gatys
 - “Content” through activations difference
 - “Style” through Gramian differences
- Instances are normalized along the channel input
- No bias in filters



Specifics of the Architecture

Layer	Activation size
Input	$3 \times 256 \times 256$
$32 \times 9 \times 9$ conv, stride 1	$32 \times 256 \times 256$
$64 \times 3 \times 3$ conv, stride 2	$64 \times 128 \times 128$
$128 \times 3 \times 3$ conv, stride 2	$128 \times 64 \times 64$
Residual block, 128 filters	$128 \times 64 \times 64$
Residual block, 128 filters	$128 \times 64 \times 64$
Residual block, 128 filters	$128 \times 64 \times 64$
Residual block, 128 filters	$128 \times 64 \times 64$
$64 \times 3 \times 3$ conv, stride 1/2	$64 \times 128 \times 128$
$32 \times 3 \times 3$ conv, stride 1/2	$32 \times 256 \times 256$
$3 \times 9 \times 9$ conv, stride 1	$3 \times 256 \times 256$



Methods	Time(s)			Styles/Model
	256×256	512×512	1024×1024	
Gatys et al. [10]	14.32	51.19	200.3	∞
Johnson et al. [47]	0.014	0.045	0.166	1



Johnson Paper Extensions

- Multi-style transfer:

Set of Style Images

For each Style Image

$$\mathcal{L}_s(I_{s0}, \dots, I_{sN}, I_{new}) = \sum_{i=0}^N \beta_i \sum_{l \in L_s} \|G_{si}^{(l)} - G_{new}^{(l)}\|^2$$



(a) A Starry Night



(b) 100% / 0%



(c) 75% / 25%



(d) 50% / 50%



(e) 25% / 75%



(f) 0% / 100%



(g) The Scream



(h) The Great Wave



(i) 100% / 0%



(j) 75% / 25%




(k) 50% / 50%



(l) 25% / 75%



(m) 0% / 100%



(n) Rain Princess



Fast Style Paper Extensions


- Color Preservation:
 - Apply style to luminance only
 - Reapply color channels from original image (blurred)
 - $HS\{V_{\text{transform}}\}$






Model Based Optimization Style Transfer

Johnson, *et. al* Fast Style Transfer



Jake Carlson



Justin Ledford



Luke Wood

Follow Along: `LectureNotesMaster/05 LectureStyleTransfer.ipynb`
Training: https://github.com/8000net/StyleTransfer/blob/master/Style_Transfer_Training.ipynb

Optional Demo

33



One Shot Transfer



The Problem of Training

- One network is only capable of one style transformation
 - Great for trained filters in an app
 - Does not work for generic transfer of one style to another
- How to define a transformation on a content and style image that transfers style into the content?
 - ...without any “per style” training
 - ...but still using the “grammian style losses”
 - ...and allows for infinite customization?
- We need to first cover **whitening** and **coloring** transformations



Aside: Whitening and Coloring

- Signal processing terms applied typically to spectra of signals
- **Whitening** is the process of removing correlation in a signal
 - For a matrix, the whitening transform:
 - ◆ makes covariance nearly diagonal (identity)
 - ◆ using only linear operations
- **Coloring** is the process of adding the observed correlation back into a signal
 - For a matrix, the coloring transform
 - ◆ makes the covariance of signal A as close as possible to signal B (the target)
 - ◆ with only linear operations to A



Aside: Whitening with SVD

- Recall that Singular Value Decomposition (SVD) decomposes a matrix into three elements
 - $A = U\Sigma V^T$
 - U is eig-vec of AA^T and V is eig-vec of A^TA
 - where U and V are orthogonal matrices such that
$$UU^T=I \quad \text{and} \quad VV^T=I$$
 - Σ is a diagonal matrix of the singular values
- the matrix $UV^T=A_w$ is a whitened version of the signal A such that
$$A_w A_w^T=I$$
- since the Gram of a layer activation is $A A^T=G$, the whitened signal has $A_w A_w^T = I = G$
 - which would have “no style” according to Gatys



Coloring with SVD

- Suppose we have two activations
 - $A_c = U_c \Sigma_c V_c^T$ and $G_c = A_c A_c^T$
 - $A_s = U_s \Sigma_s V_s^T$ and $G_s = A_s A_s^T$
- We can transfer the Gram matrix of A_s to A_c using coloring
 - To color the matrix:
$$A_{new} = U_s \Sigma_s V_c^T$$
$$A_{new} A_{new}^T = G_s$$
 - But A_{new} is more similar to A_c than A_s
- That is exactly what we want for style transfer!
- *However, there are some computational problems with this method of coloring, so in practice, we use the SVD of the Grammian to perform coloring.*





Whitening and Coloring

Computational Problems with coloring
and the Eigen Decomposition

Demo will introduce concept of coloring with PCA

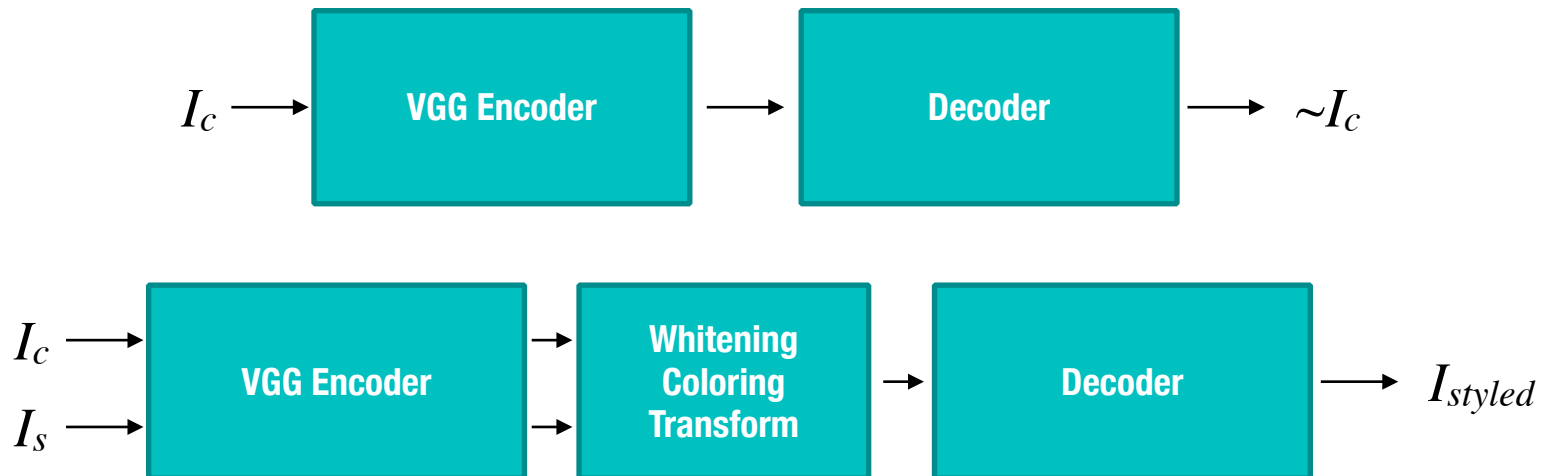
Follow Along:

`05b WhiteningColoringExample.ipynb`



How to use this for style transfer?

- If we can learn to **reconstruct an image from VGG...**
 - ...we can **whiten content activations**
 - ...then **color with desired style Grammmian**
- The resulting reconstruction should have largely the same content, but style from the colored activations
- ...One transformation network for any style!



Multi-Staged WCT

Auto
Encoder



I_c


VGG Encoder

Decoder

$\sim I_c$



Single
Style
Transfer



I_c

I_s

VGG Encoder


WCT

Decoder

I_{style}



Multi-
Staged
Style
Transfer



I_c


I_s

VGG Encoder
(Layer L)

Whitening
Coloring
Transform

Decoder
(From Layer L)

I_{s1}



I_{s1}

I_s

VGG Encoder
(Layer L-1)

Whitening
Coloring
Transform

Decoder
(From Layer L-1)

I_{s2}



I_{s2}

I_s

VGG Encoder
(Layer L-2)

Whitening
Coloring
Transform

Decoder
(From Layer L-2)

I_{styled}



Lecture Notes for **Neural Networks and Machine Learning**

Style Transfer: Model Opt.



Next Time:
Photo Realistic WCT
Reading: None

