

# Lecture Notes for **Neural Networks and Machine Learning**



Cross Entropy and  
Value Iteration



# Logistics and Agenda

- Logistics
  - Grading Update
  - Class schedule
- Agenda
  - Finish: Cross Entropy Method
  - Value Iteration (and demo)
  - Tabular Q-Learning
  - Deep Q-Learning (next time?)



# Last Time

```
import gym
```

```
if __name__ == "__main__":
    env = gym.make("CartPole-v0")
```

```
total_reward = 0.0
total_steps = 0
obs = env.reset()
```

```
while True:
```

```
    action = env.action_space.sample()
```



```
    obs, reward, done, _ = env.step(action)
```

```
    total_reward += reward
```

```
    total_steps += 1
```

```
    if done:
```

```
        break
```




**Action Space:** One input, [0, 1] pull left or pull right

**Obs Space:** Dynamic state variables (continuous and four dimensional)

**End:** When more than 15 degrees off or too far from center

**Reward:** +1 for each time step



Edward Thorndike



B.F. Skinner



Ivan Pavlov



Bernard Widrow



Marvin Minsky



Ted Hoff



Claude Shannon



# Cross Entropy Method





# Direct Policy Exploration and Optimization

- Instead of defining what is optimal, just setup a comparison of different actions we might take (**policy**)
- A **policy** is defined as  $\pi(a, s) = P(a_t = a \mid s_t = s)$ 
  - Given the current state, we have a certain probability of selecting each action
  - Action selection is **probabilistic**, but easy to discover **deterministic** actions (*set one action to 1.0, all others to 0.0*)
- Try different policies, select one with best average reward
- First try: Cross Entropy Method



# Cross Entropy Method

- Create a random neural network, with output  $p(a|s)$
- Let it interact with the environment (randomly)
  - For some set of episodes (e.g., 20)
    - ◆ Use network output to sample from possible actions
    - ◆ Run episode to completion
    - ◆ Repeat
- Calculate reward for each episode
- Keep best episodes (some percentile, e.g., best five)
- For the given best episodes, develop loss function incentivizing the actions taken based upon the input observations



**Repeat until desired performance!**



# Cross Entropy Method

- Model based or Model Free?
  - Model Free (no assumptions of problem)
- Value or Policy Based?
  - Policy Based (randomly sample actions based on policy)
- On-policy or Off-Policy?
  - On-Policy (need to interact with environment to get better)



# Mathematical Motivation

- If we have the optimal policy  $p(x)$  and a reward function  $H(x)$ , then maximize

$$\mathbf{E}_{x \leftarrow p(x)}[H(x)] = \mathbf{E}_{x \leftarrow q(x)}\left[\frac{p(x)}{q(x)}H(x)\right]$$

- We can approximate the distribution by:  $\frac{1}{N} \sum_i \frac{p(x_i)}{q(x_i)} H(x_i)$
- Proven that this is optimized when  $\text{KL}(q(x) \parallel p(x)H(x))$  is minimized. But its intractable, so we can only optimize upper bound ... minimizing (neg) cross entropy of samples

$$\pi_{k+1}(a | s) = \arg \max_{\pi_k} \mathbf{E}_{z \leftarrow \pi_k}^{\text{Performance Measure}} [\mathbf{1}_{R(z) > \psi} \log \pi_k(a | s)]$$

min CrossEntropy( *neural\_net\_actions*, *best\_actions* )







# Review: Basics of Cartpole

```
import gym

if __name__ == "__main__":
    env = gym.make("CartPole-v0")

    total_reward = 0.0
    total_steps = 0
    obs = env.reset()

    while True:
        action = env.action_space.sample()
        obs, reward, done, _ = env.step(action)
        total_reward += reward
        total_steps += 1
        if done:
            break
```



**Action Space:** One input,  $[0, 1]$  pull left or pull right


**Obs Space:** Dynamic state variables (continuous and four dimensional)

**End:** When more than 15 degrees off or too far from center

**Reward:** +1 for each time step




# Another Example: Frozen Lake




- **State:** Every square in grid
- **Action:** Move to make (l,r,u,d), *with probability*
- **Reward:** Goal, Death
- **Policy:** Given state, where should we move?
- **Optimal Policy:**


$$\pi^* = \arg \max_{\pi} \mathbf{E} \left[ \sum_k \gamma^k R_{t+k+1} \mid \pi \right]$$



Random Policy



Another Policy



Another Policy





# Cross Entropy Reinforcement Learning

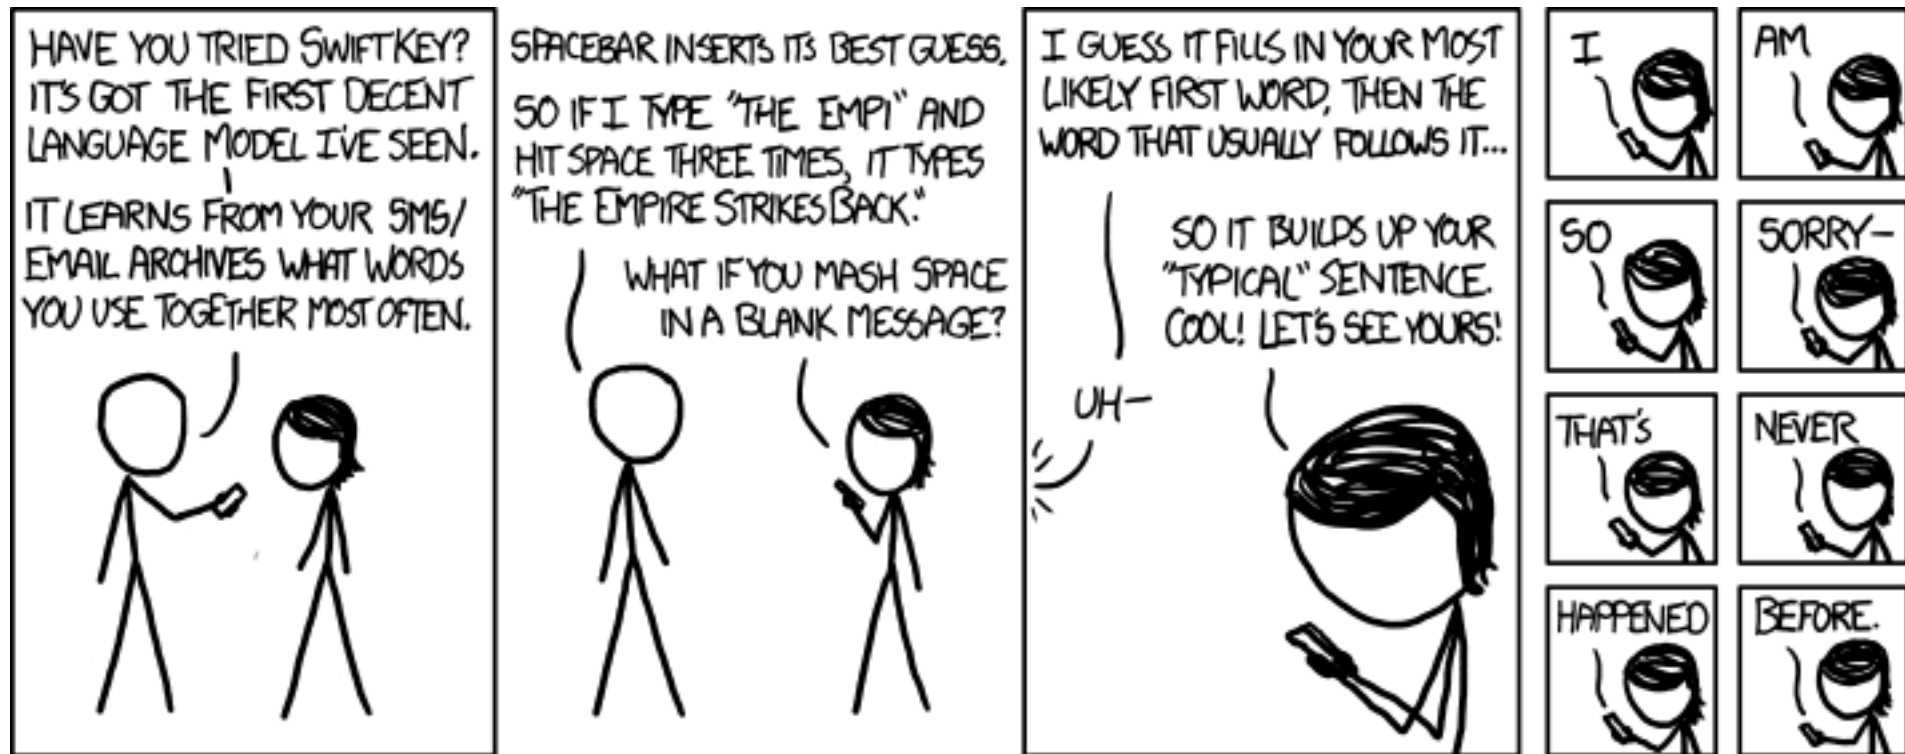
M. Lapan Implementation for CartPole  
and Frozen Lake

Follow Along:

`08a_Basics_Of_Reinforcement_Learning.ipynb`



# Markov Building Blocks



# Markov Processes (MP)


- **Definition:** Any process that can be explained (or simplified) through a sequential set of states that depend only on the previous state
- **Practical Meaning:** For  $N$  states, there will be the probability of transition to any other state, encoded through an  $N \times N$  transition matrix of discrete probabilities
- State sequences are not deterministic, they are sampled from these distributions
- Despite **simplicity**, MP can model a number of real processes with **good enough** precision

|                      | Next State, $s_{t+1}$ |     |     |     |     |
|----------------------|-----------------------|-----|-----|-----|-----|
|                      | 0.1                   | 0.2 | 0.1 | 0.6 | 0.0 |
| Current State, $s_t$ | 0.9                   | 0.0 | 0.1 | 0.0 | 0.0 |
|                      | 0.0                   | 0.4 | 0.0 | 0.4 | 0.2 |
|                      | 0.0                   | 0.4 | 0.2 | 0.0 | 0.4 |
|                      | 0.0                   | 0.0 | 0.6 | 0.0 | 0.4 |



# MP Example from Maxim Lapan

|       | Sunny' | Rainy' |
|-------|--------|--------|
| Sunny | 0.8    | 0.2    |
| Rainy | 0.6    | 0.4    |



|              |  |  |  | ... |  |
|--------------|--|--|--|-----|--|
| Sun+Summer   |  |  |  |     |  |
| Rainy+Summer |  |  |  |     |  |
| Sun+Fall     |  |  |  |     |  |
| Rainy+Fall   |  |  |  |     |  |
| Sun+Else     |  |  |  |     |  |
| Rainy+Else   |  |  |  |     |  |

**Adding One Variable Can Have  
Drastic Effect on State Space Size**

