

Lecture Notes for **Neural Networks and Machine Learning**



Deep-Q Demo



Logistics and Agenda

- Logistics
 - Final paper, OK?
- Agenda
 - Student Presentations
 - ◆ SAC
 - ◆ AlphaFold
 - Deep Q-Learning Demonstration



Paper Presentation



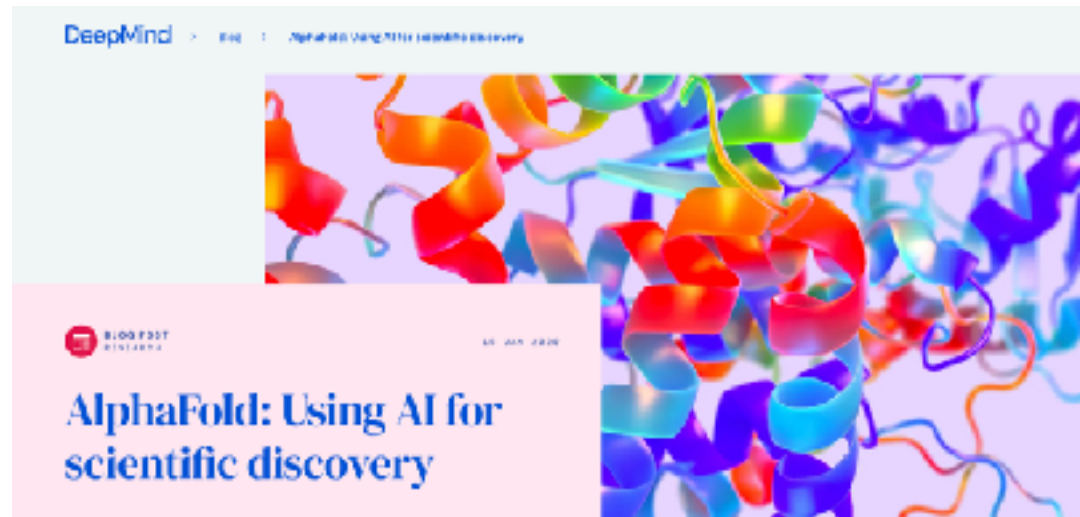
Jürgen Schmidhuber
@SchmidhuberAI

...

Big news [#AlphaFold](#) uses [#DeepLearning](#) for [#ProteinFolding](#) prediction. This approach was pioneered by Sepp Hochreiter et al. in 2007 when compute was 1000 times more expensive than today. Their LSTM was orders of magnitude faster than the competitors.



Fast model-based protein homology detection without alignment - Pub...
pubmed.ncbi.nlm.nih.gov





Deep Q-Learning

Reinforcement Learning

M. Lapan Implementation for
Frozen Lake and Atari!

$$Q(s, a) \leftarrow (1 - \alpha) \cdot Q(s, a) + \alpha \cdot [r_{s,a} + \gamma \max_{a' \in A} Q(s', a')]$$

$$\mathcal{L} = \left[\underset{\substack{\text{from current network} \\ \text{params}}}{Q(s, a)} - \underset{\substack{\text{from older network params} \\ \text{(better stability)}}}{[r_{s,a} + \gamma \max_{a' \in A} Q^*(s', a')]} \right]^2$$

$$\mathcal{L} = [Q(s, a) - [r_{s,a}]]^2$$

if no next state (env is done)

Follow Along:

`08a_Basics_Of_Reinforcement_Learning.ipynb`



World Models



The Problem

World Models

Can agents learn inside of their own dreams?

DAVID HA	JÜRGEN SCHMIDHUBER
Google Brain	NNAISENSE
Tokyo, Japan	Swiss AI Lab, IDSIA (USI & SUPSI)

March 27
2018

NIPS 2018
Paper

YouTube
Talk

Download
PDF



<https://worldmodels.github.io>



A Motivation

Agents can dream! What a time to be alive!

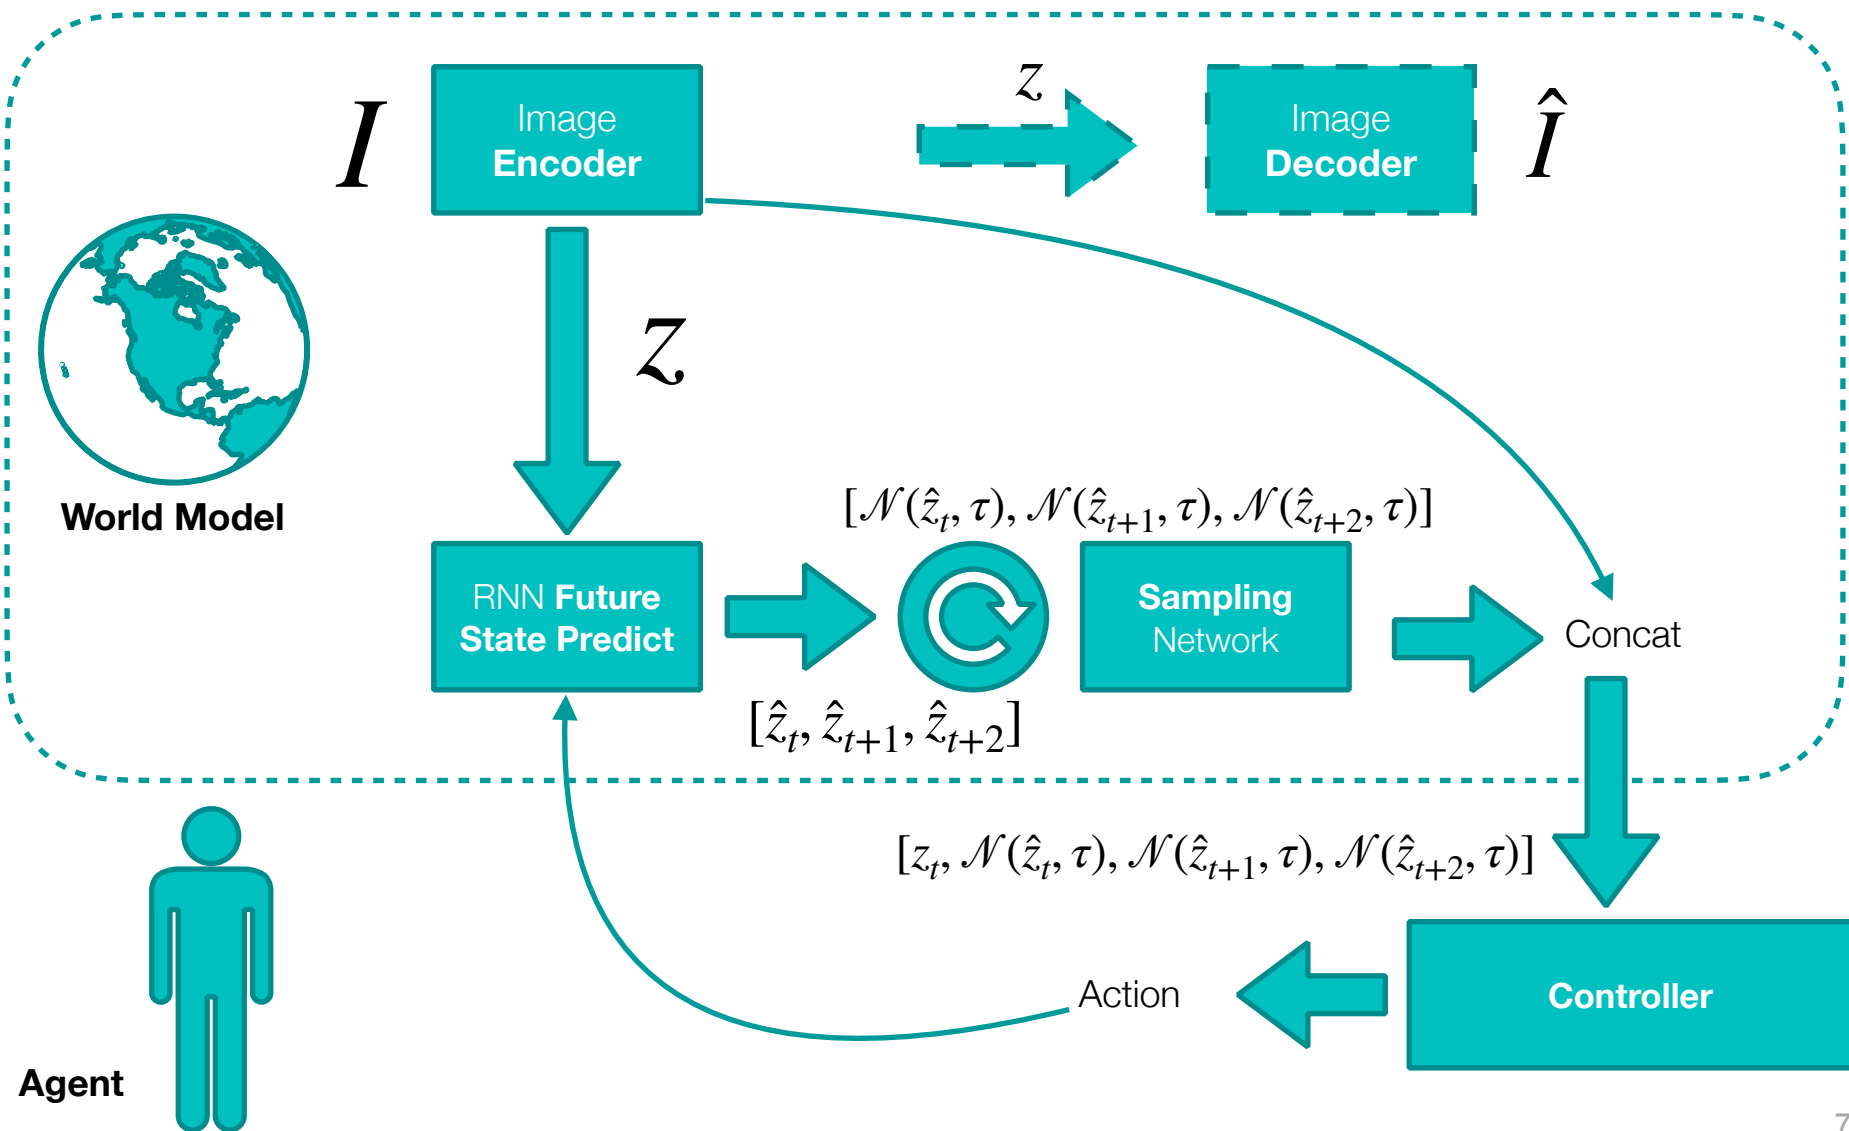
And academia can dream about driving the hype train!



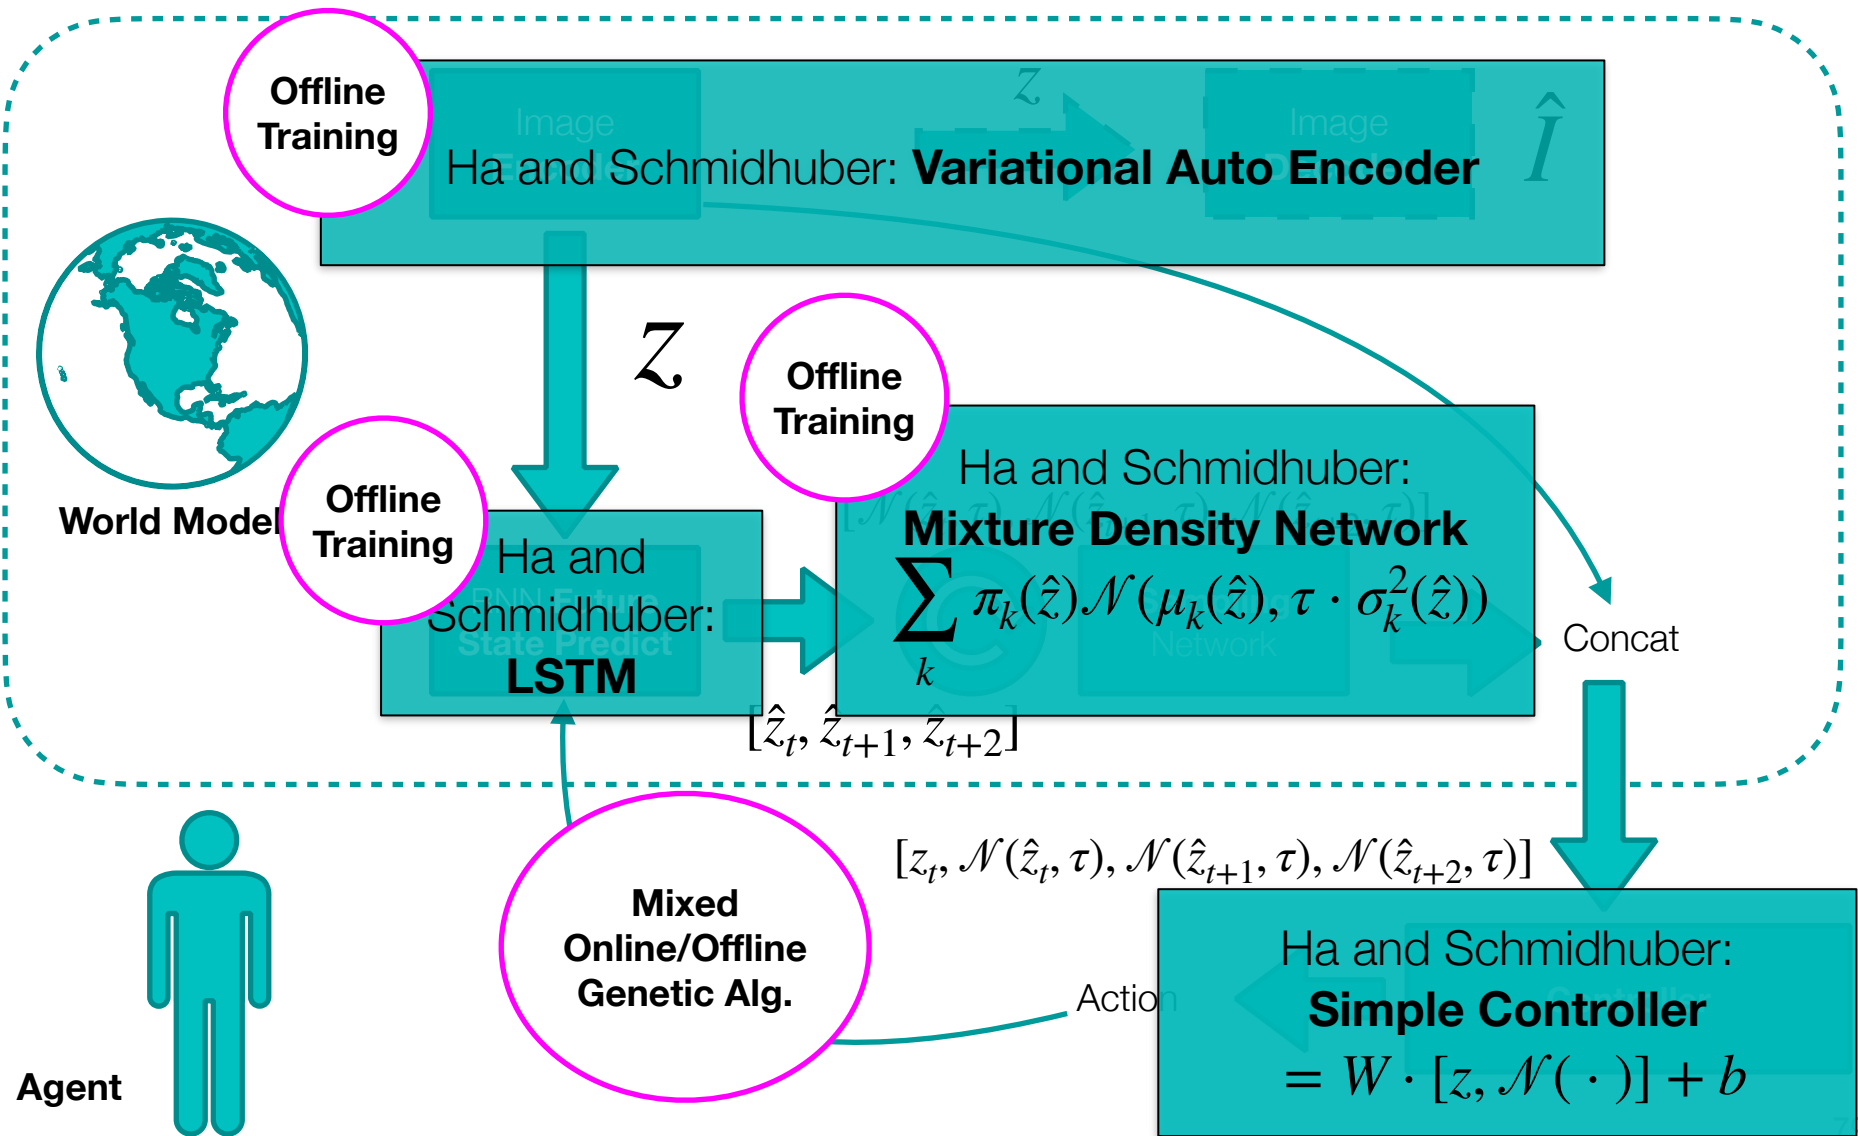
Maybe we should be more careful about the way we describe what an agent does... because they don't dream. That's fluff.



The Main Idea





Implementation



An Example, Racing

- Schmidhuber and Ha Methods:
 - Collect 10,000 rollouts from a random policy.
 - Train VAE (Λ) to encode each frame
 - Train
 - Evolve cumulative

Method	Average Score over 100 Random Tracks	Model	Parameter Count
DQN [53]	343 \pm 18	VAE	4,348,647
A3C (continuous) [52]	591 \pm 45		422,368
A3C (discrete) [51]	652 \pm 10		867
ceobillionaire's algorithm (unpublished) [47]	838 \pm 11		
V model only, z input	632 \pm 251		
V model only, z input with a hidden layer	788 \pm 141		
Full World Model, z and h	906 \pm 21		



Only use VAE Encoding

<https://worldmodels.github.io>

Full World Model

80



Lecture Notes for **Neural Networks and Machine Learning**

Deep Q Learning



Next Time:
World Models
Reading: None

