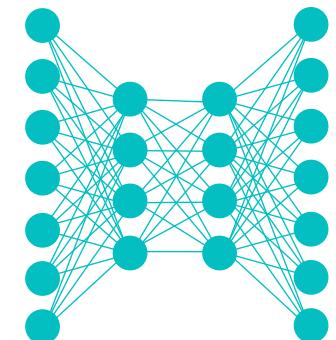


Lecture Notes for **Deep Learning II**



Course Introduction
Lecture: Deep Learning Review



Logistics and Agenda

- Logistics
 - This class evolves across semesters (drastically!)
 - ◆ First offered in 2019
 - Using Canvas
 - Office Hours: after class, Mondays
 - GitHub: Basically one repository
 - Job Posting: Audio Annotation
- Agenda
 - Syllabus and Introductions
 - Presentation Selection
 - Topic Review



Syllabus

- Course Schedule 😅
- Reading / Distance
- GitHub
- Grading
 - Labs x4 (60%)
 - Final Presentation and Project x1 (30%)
 - ◆ Could be replaced with comprehensive exam
 - Participation (Pass/Fail)
 - Class Presentation Summary (10%)
- LLM Usage



People



3



8000net

This organization houses a number of repositories for Dr. Larson's 8000 Level Neural Networks Course, Offered at SMU



Presenting

- First Presentation is coming up soon!
- During Semester: 12 Presentations Total (as a team)
- First Presentation → **Who wants to go first?**
 - ~10-15 Minutes (~5 mins per group member)
 - Present the topic of the paper, summarize important aspects
 - Make about 5 Visuals
 - ◆ typically: Slides
 - ◆ AND/OR Handouts
 - ◆ AND/OR Notebooks



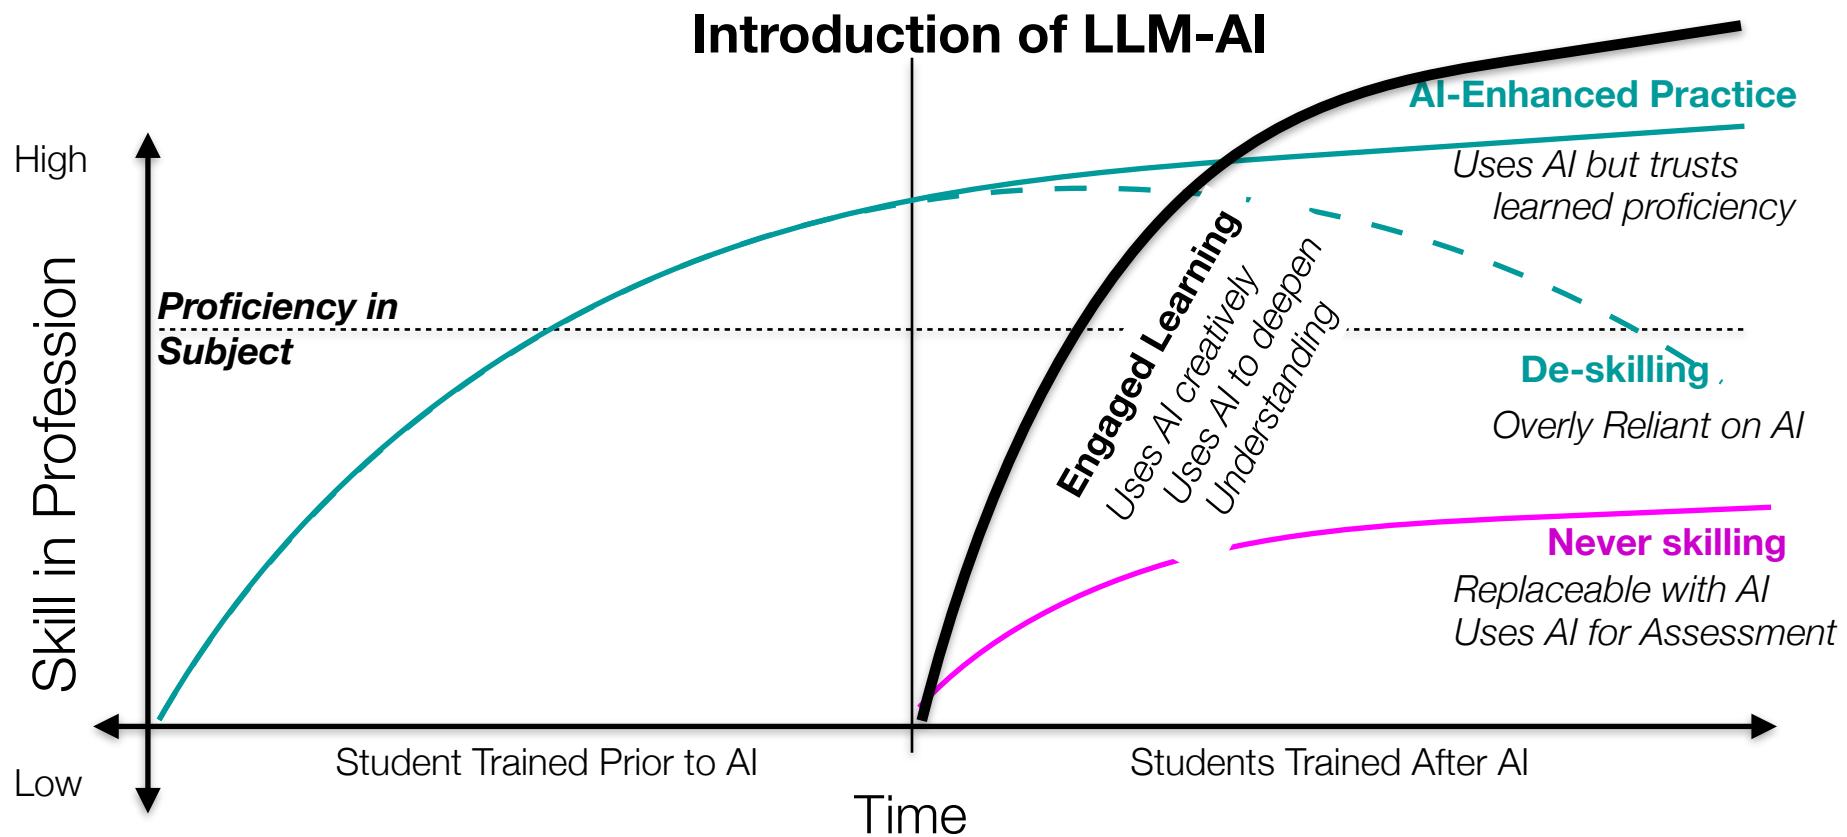
Introductions

- Name, Majors, and Designation
- When you took CS5/7324
- The Topic in this course you are most excited about
- Something true or false about you

- Do NOT forget:
 - Pick out papers on Canvas



DL Review

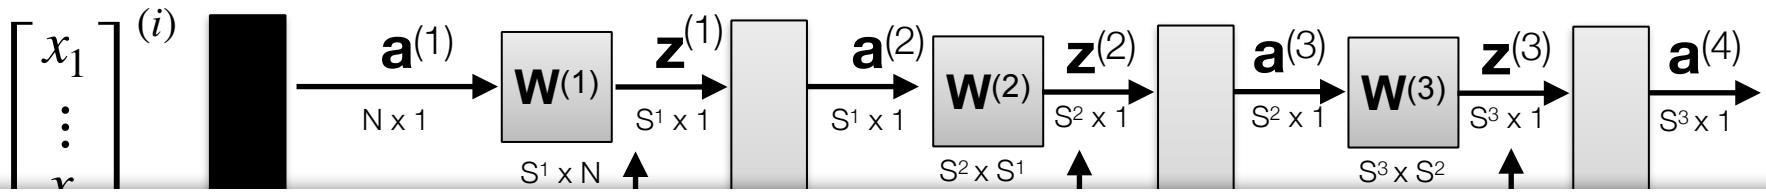


Prerequisite Topics, Review

- Feedforward networks and back propagation
- Optimization methods (SGD, AdaM)
- Computation graph and auto differentiation
- Embedding layers and multi-modal input
- Convolutional networks
- Sequential networks with emphasis on transformers
 - some transformer review will happen later in semester



Neural Network Structure



$$\mathbf{Z}^{(L)} = \mathbf{W}^{(L)} \cdot \mathbf{A}^{(L)} + \mathbf{b}^{(L)}$$



$$\mathbf{Z}^{(L)} = \mathbf{W}^{(L)} \cdot \phi(\mathbf{Z}^{(L-1)}) + \mathbf{b}^{(L)}$$

$$[\mathbf{z}^{(L)}]^{(i)} = \mathbf{W}^{(L)} \cdot [\mathbf{a}^{(L)}]^{(i)} + \mathbf{b}^{(L)} \quad \begin{bmatrix} z_1^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix} = \mathbf{W}^{(L)} \cdot \begin{bmatrix} a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix} + \mathbf{b}^{(L)}$$

b is broadcast added

$$\begin{bmatrix} \begin{bmatrix} z_1^{(L)} \\ z_2^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix}^{(1)}, \begin{bmatrix} z_1^{(L)} \\ z_2^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix}^{(2)}, \dots, \begin{bmatrix} z_1^{(L)} \\ z_2^{(L)} \\ \vdots \\ z_{S^L}^{(L)} \end{bmatrix}^{(M)} \end{bmatrix} = \mathbf{W}^{(L)} \cdot \begin{bmatrix} \begin{bmatrix} a_0^{(L)} \\ a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix}^{(1)}, \begin{bmatrix} a_0^{(L)} \\ a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix}^{(2)}, \dots, \begin{bmatrix} a_0^{(L)} \\ a_1^{(L)} \\ \vdots \\ a_{S^{L-1}}^{(L)} \end{bmatrix}^{(M)} \end{bmatrix} + \mathbf{b}^{(L)}$$



Adaptive Optimization

Adjust each element of gradient by the steepness (for each layer):

- AdaGrad

$$\rho_k = \frac{1}{\sqrt{\mathbf{G}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$

all operations are per element

where

$$\mathbf{G}_k = \gamma \cdot \mathbf{G}_{k-1} + \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

- RMSProp

$$\rho_k = \frac{1}{\sqrt{\mathbf{V}_k + \epsilon}} \odot \nabla J(\mathbf{W}_k)$$

all operations are per element

$$\mathbf{G}_k = \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k)$$

$$\mathbf{V}_k = \gamma \cdot \mathbf{V}_{k-1} + (1 - \gamma) \cdot \mathbf{G}_k$$

- AdaDelta

$$\rho_k = \frac{\mathbf{M}_k}{\sqrt{\mathbf{V}_k + \epsilon}}$$

all operations are per element

$$\mathbf{M}_{k+1} = \gamma \cdot \mathbf{M}_k + (1 - \gamma) \cdot \nabla J(\mathbf{W}_k)$$

- AdaM

$$\begin{array}{l} \text{update momentum} \\ \mathbf{M}_{k+1} \leftarrow \beta_1 \cdot \mathbf{M}_k + (1 - \beta_1) \cdot \nabla J(\mathbf{W}_k) \end{array}$$

$$\hat{\mathbf{M}}_k \leftarrow \frac{\mathbf{M}_k}{(1 - [\beta_1]^k)}$$

$$\begin{array}{l} \text{normalizer momentum} \\ \mathbf{V}_{k+1} \leftarrow \beta_2 \cdot \mathbf{V}_k + (1 - \beta_2) \cdot \nabla J(\mathbf{W}_k) \odot \nabla J(\mathbf{W}_k) \end{array}$$

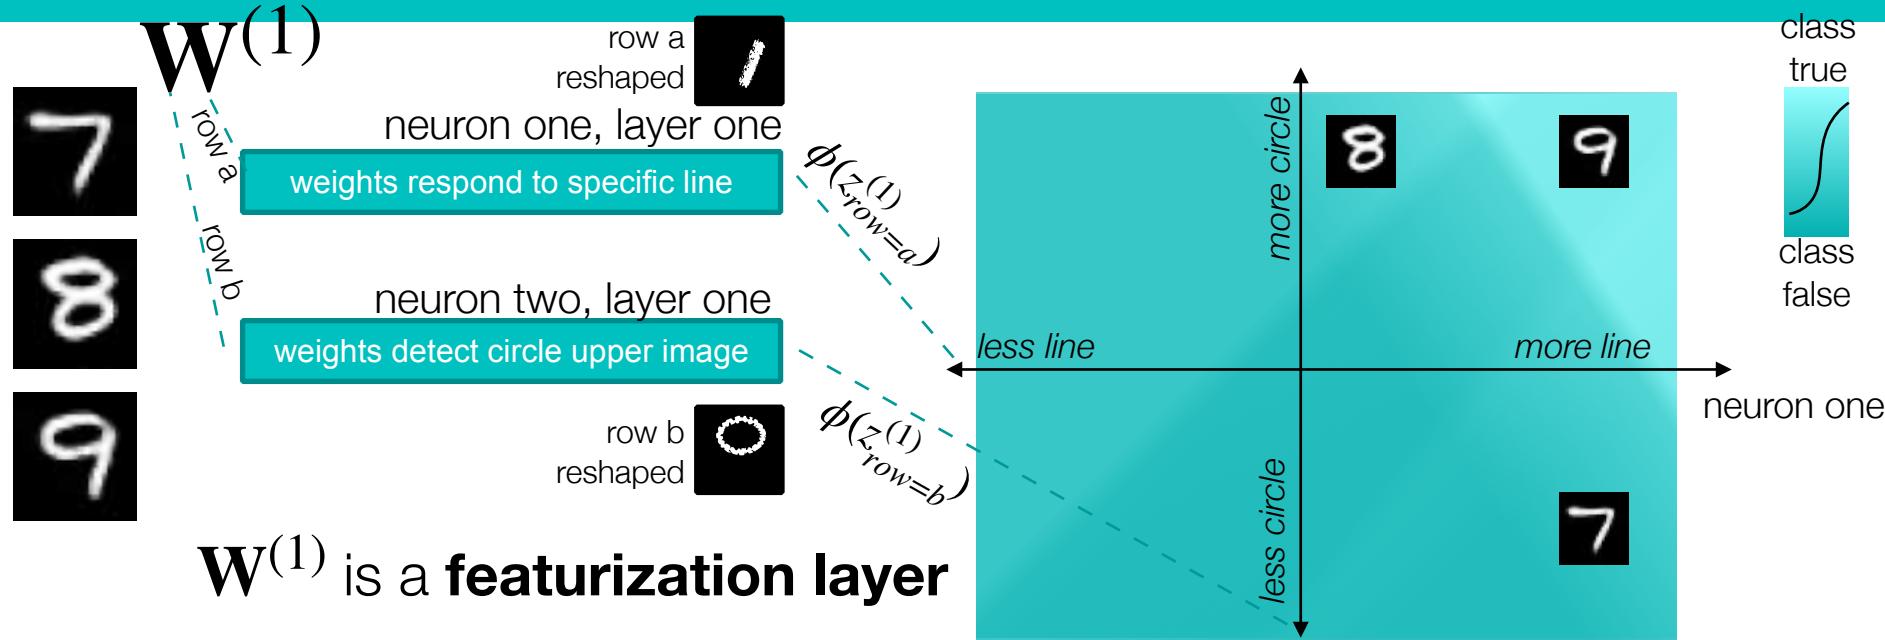
$$\hat{\mathbf{V}}_k \leftarrow \frac{\mathbf{V}_k}{(1 - [\beta_2]^k)}$$

$$\begin{array}{l} \text{full update} \\ \mathbf{W}_{k+1} \leftarrow \mathbf{W}_k - \eta \cdot \frac{\hat{\mathbf{M}}_k}{\sqrt{\hat{\mathbf{V}}_k + \epsilon}} \end{array}$$

exploitation, boosting



Universality



$W^{(1)}$ is a **featurization layer**

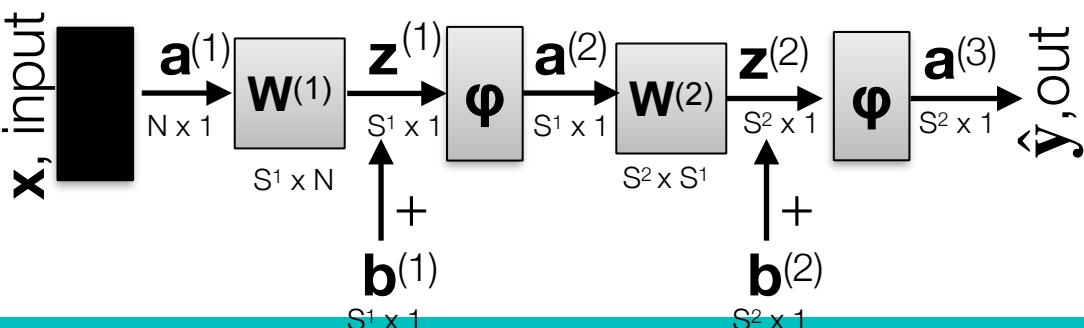
want



might get...



adding layers and convolution
help obtain desired behavior...

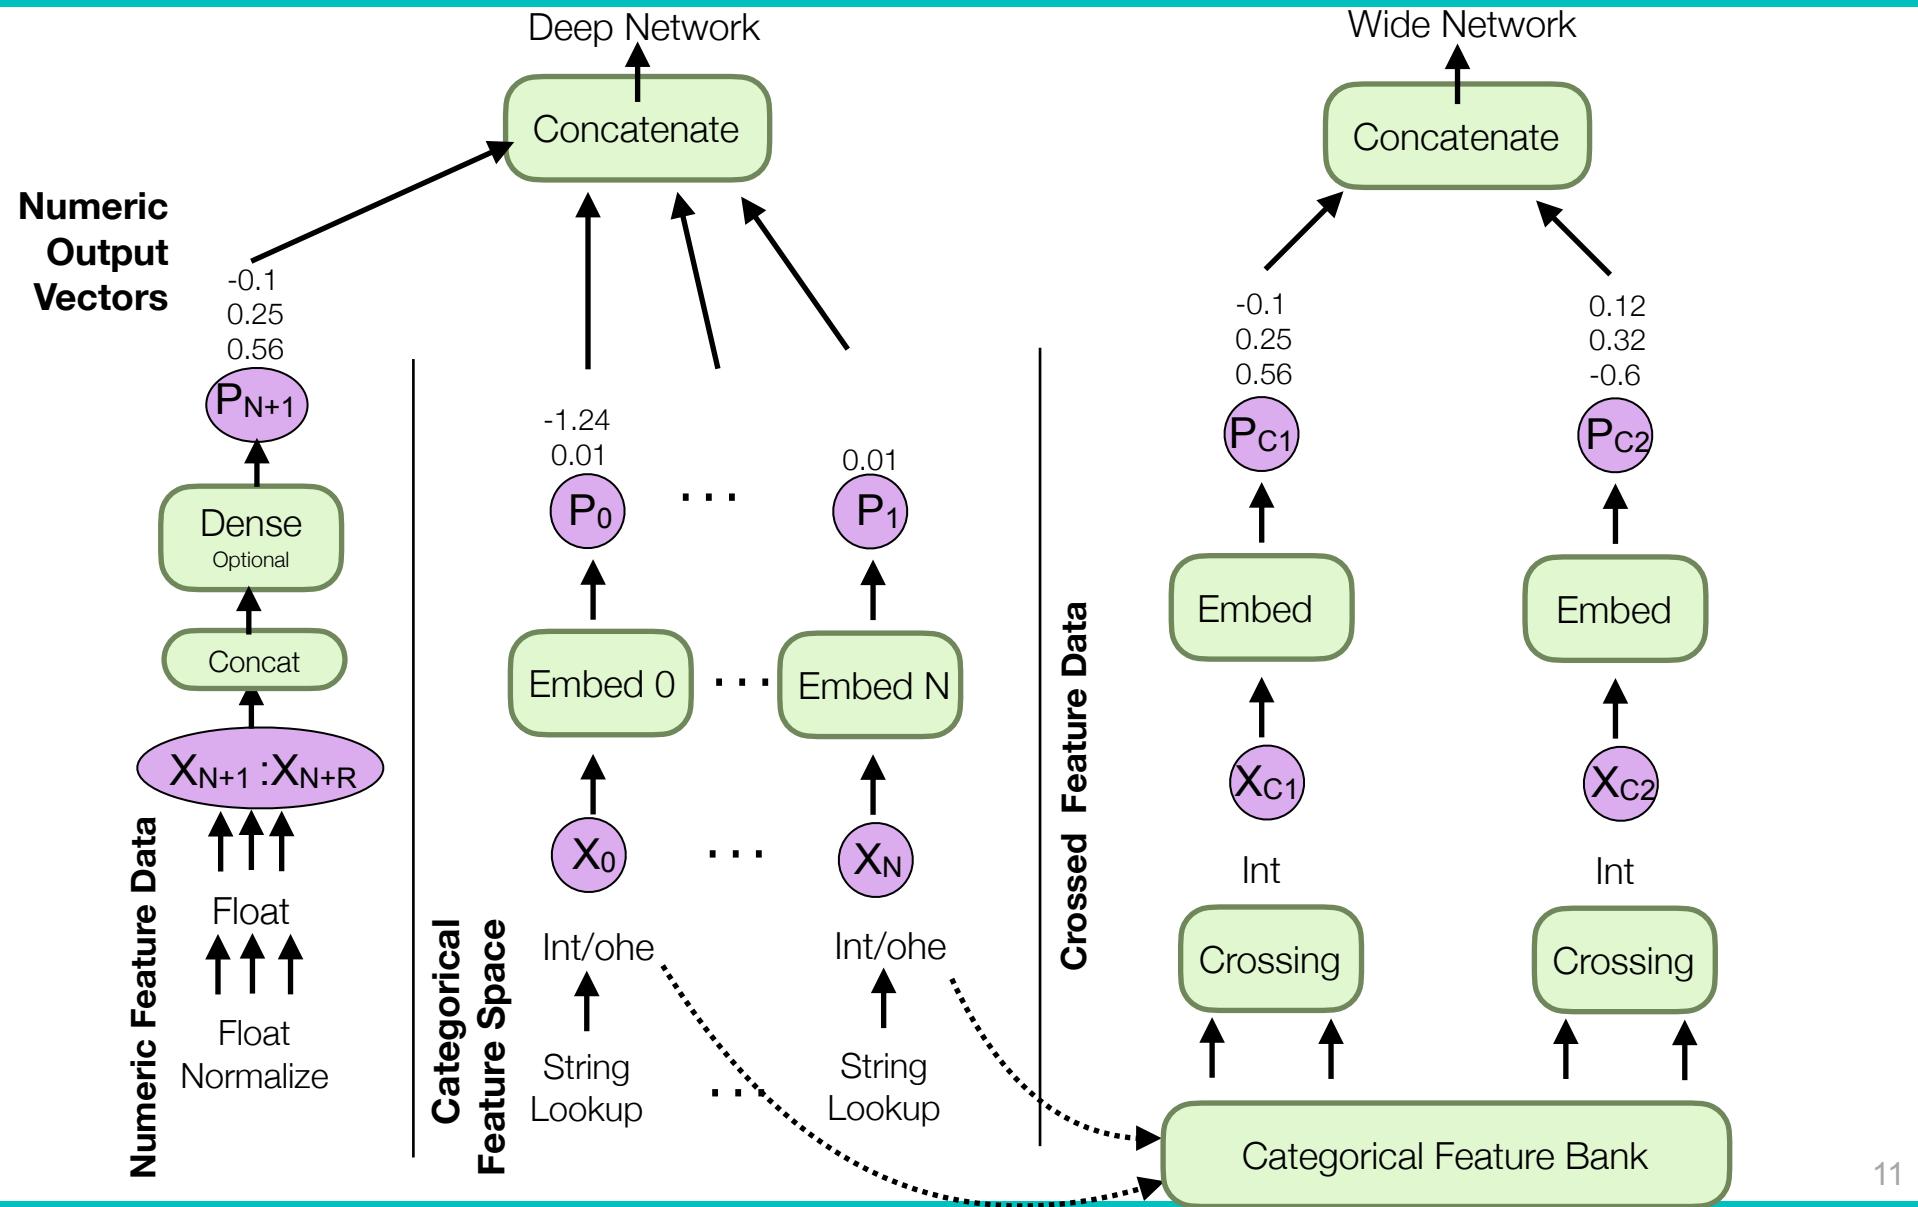


$W^{(2)}$ is **classification layer**, like
one versus all logistic regression

- One nonlinear hidden layer with an output layer can **perfectly train any problem with enough data**, but might be memorizing...
- ... could be better to have **even more layers** for more generalizing features

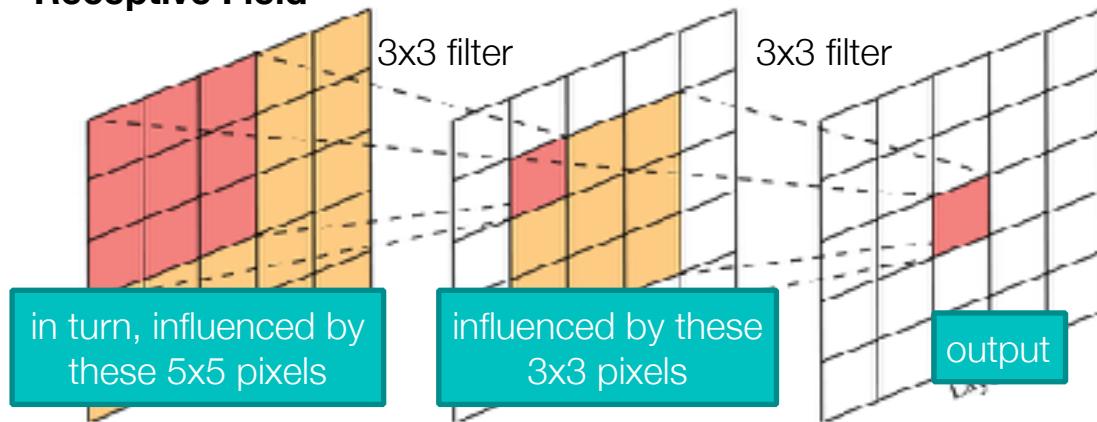


Computation Graph, Feature Spaces

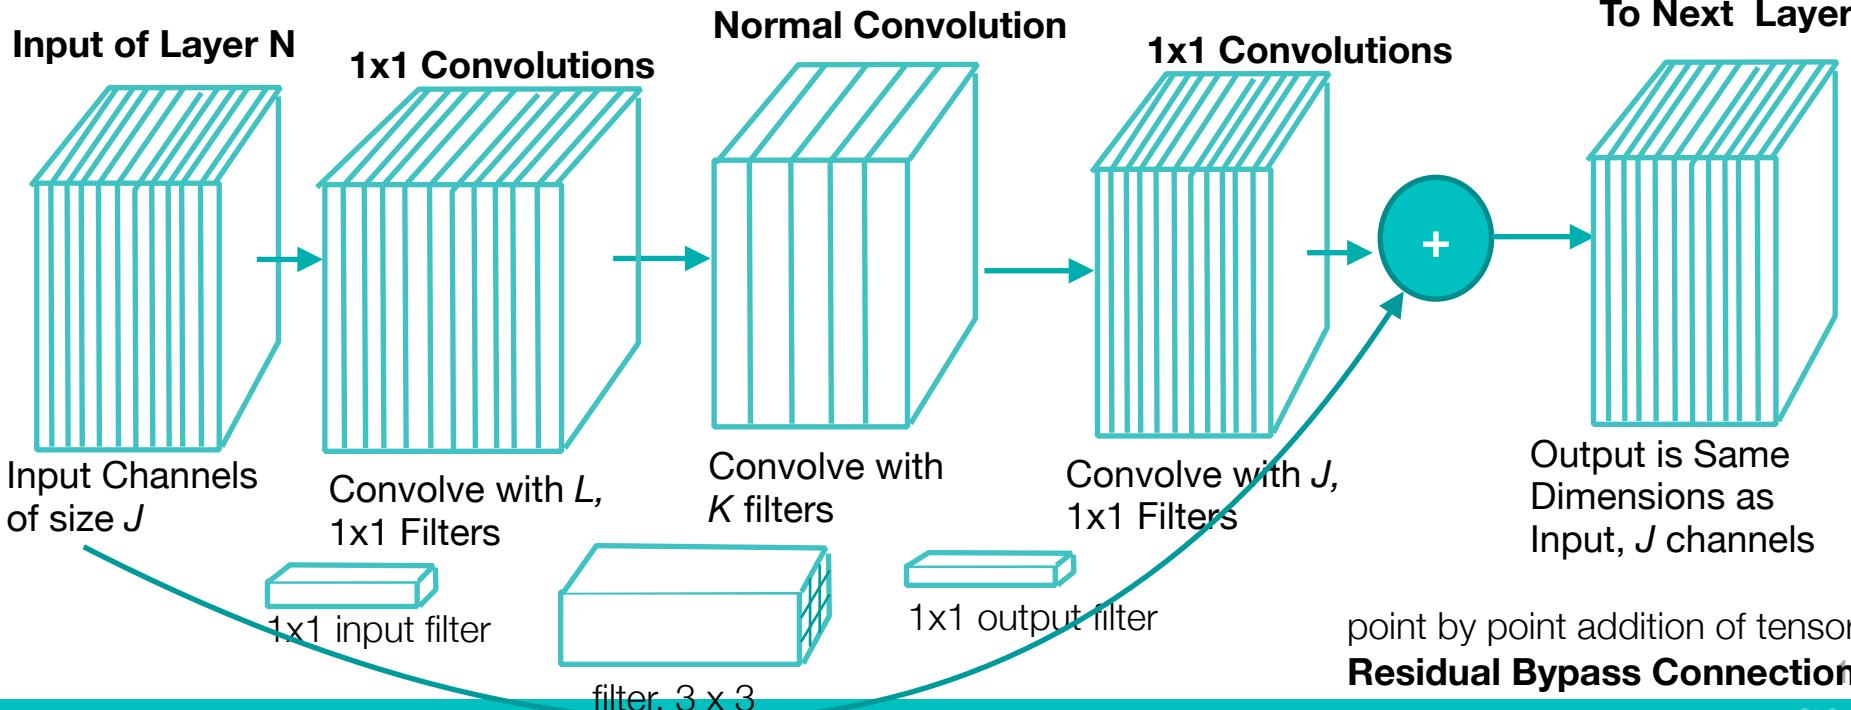


Convolutional Networks

Receptive Field

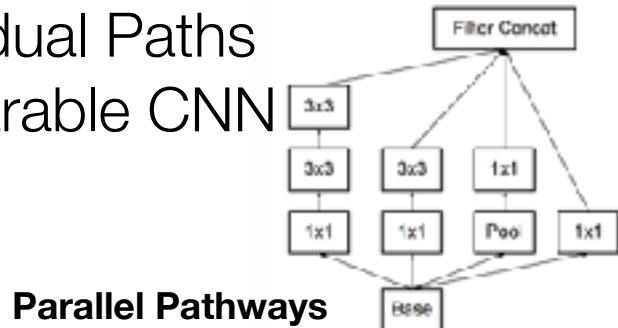


Input of Layer N

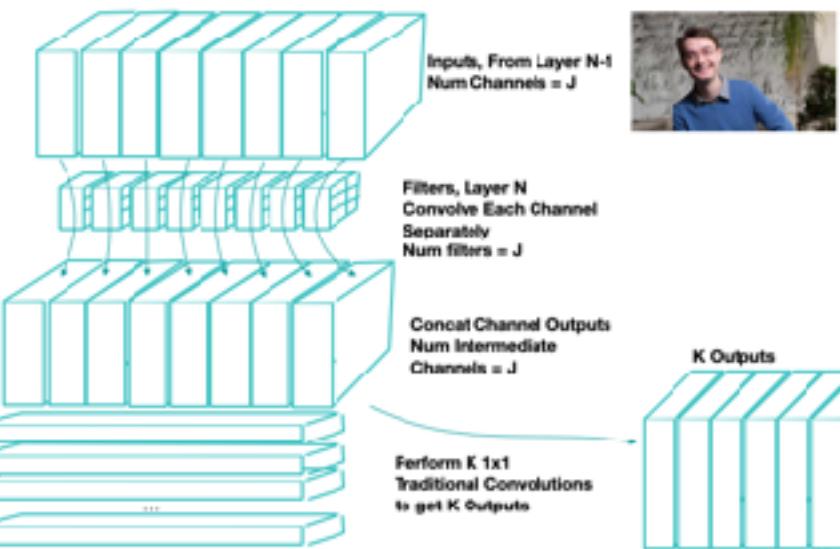


CNN Techniques

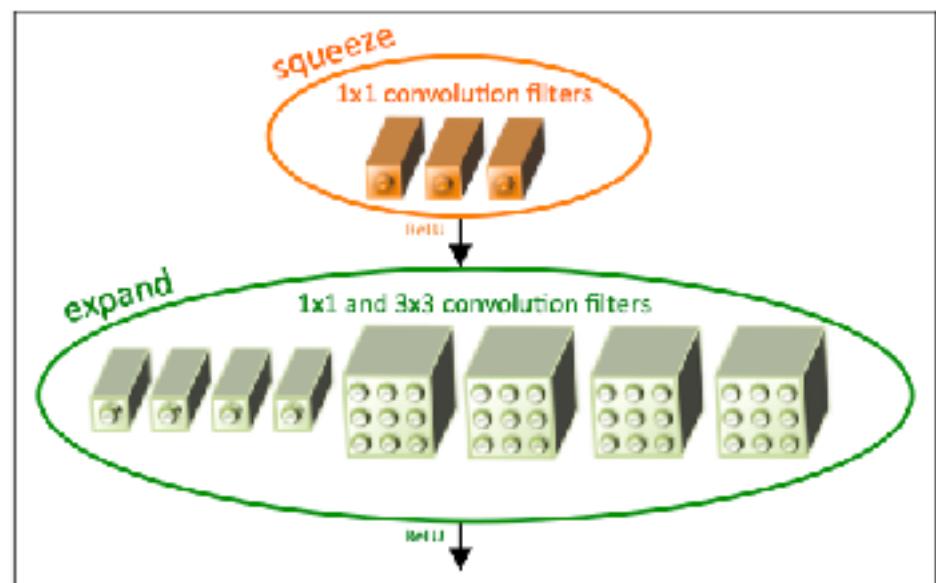
- Bottlenecks (1x1 filters)
- Parallel Paths, Concatenation
- Residual Paths
- Separable CNN



Parallel Pathways



Separable Convolution



Squeeze

$$SR = \frac{|F_{s1x1}|}{|F_{e1x1}| + |F_{e3x3}|}$$

Controls how much to bottleneck

$$PCT_{3x3} = \frac{|F_{e3x3}|}{|F_{e1x1}| + |F_{e3x3}|}$$

Controls num filter params
(how many 1x1 versus 3x3)



Lecture Notes for **Deep Learning II**



Course Introduction
Lecture: Deep Learning Review

