


# Lecture Notes for **Neural Networks and Machine Learning**



Vision Transformers  
Self-supervised Learning  
Consistency Loss



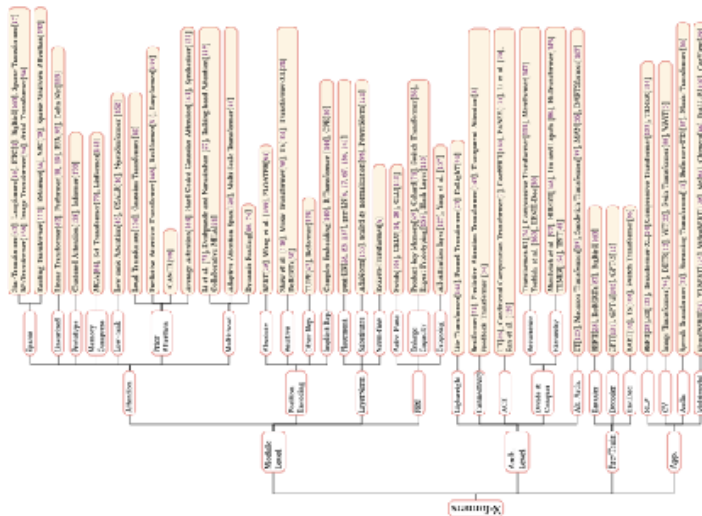
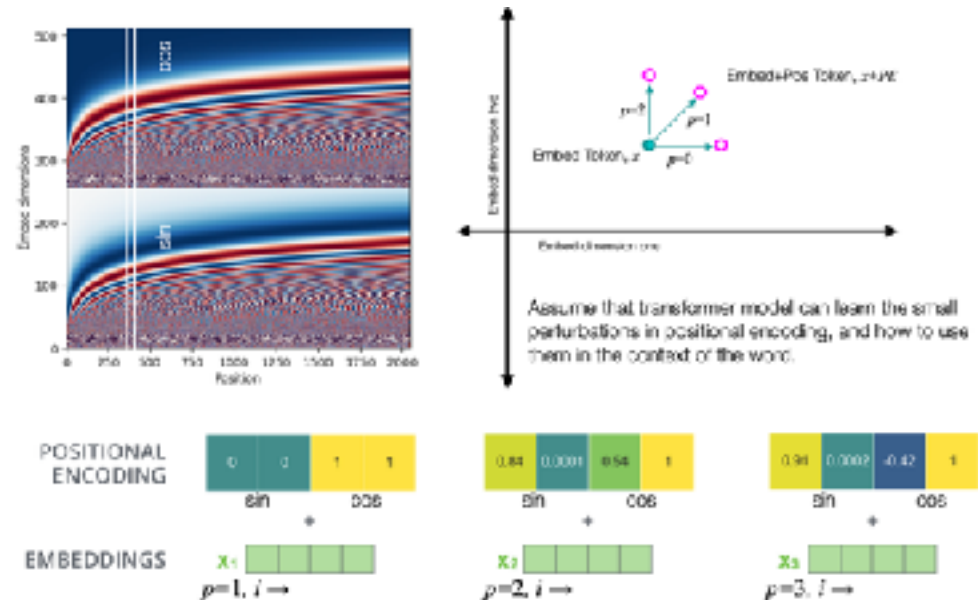

# Logistics and Agenda

- Logistics
  - None!
- Agenda
  - Vision Transformers
  - Self Supervised Learning and Consistency Loss
- Next Time:
  - Paper Presentation: *Language Models are Few Shot Learners*
  - Multi-modal Learning
    - ◆ Techniques
    - ◆ Applications and domains
  - Multi-Task and Demo



# Last Time: Transformers

## Transformer: Multi-headed Attention




# Vision Transformers



# Vision Transformers

- Divide image into patches
  - Treat each patch as something to encode separately
  - Flatten each patch
  - Put through dense layer
- Add positional encoding based on position of patch
  - for 7x7 patch, there are 49 positions
- Put into transformer. Same as text transformers ...
- **But you need a lot of data**
  - 14M or more images seems to be sweet spot



# Vision Transformers Video



<https://ai.googleblog.com/2020/12/transformers-for-image-recognition-at.html?m=1>



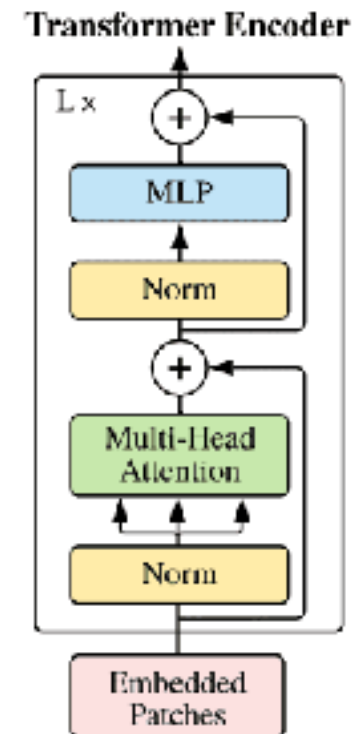
# ViT Architectures

- $D$  is size of patch embedding
- Uses skip connections (all size  $D$ )
- Multi-headed self attention (MSA) takes  $D$  input patch\_embed + pos\_embed
- Main difference in architectures
  - $L$  blocks used (*i.e.*, “layers”)
  - $H$  heads in each layer (*i.e.*, “heads”)
  - MLP head is final classifier


$$\begin{aligned} \mathbf{z}_0 &= [\mathbf{x}_{\text{class}}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{\text{pos}}, \\ \mathbf{z}'_\ell &= \text{MSA}(\text{LN}(\mathbf{z}_{\ell-1})) + \mathbf{z}_{\ell-1}, \\ \mathbf{z}_\ell &= \text{MLP}(\text{LN}(\mathbf{z}'_\ell)) + \mathbf{z}'_\ell, \\ \mathbf{y} &= \text{LN}(\mathbf{z}_L^0) \end{aligned}$$

Model	Layers	Hidden size $D$	MLP size	Heads	Params
ViT-Base	12	768	3072	12	86M
ViT-Large	24	1024	4096	16	307M
ViT-Huge	32	1280	5120	16	632M

ResNet50: 23M



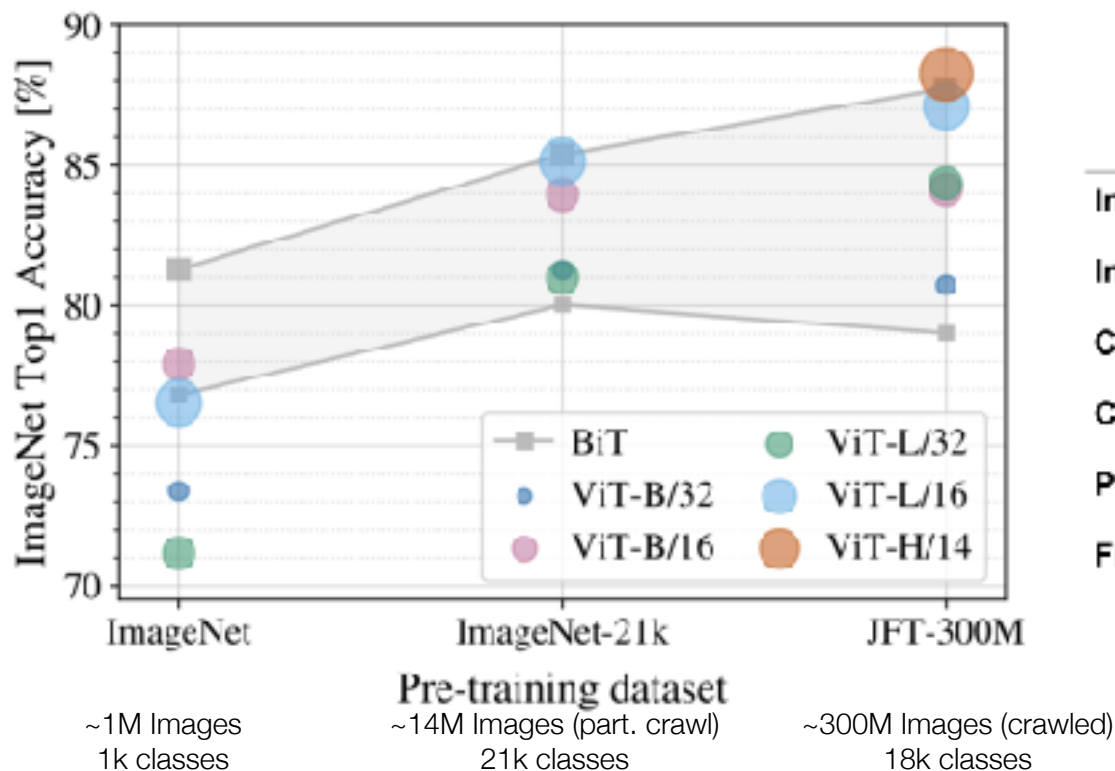
# What is the learnable class embedding?





# Do they work?

- Yes, but good luck getting weights for them or training them, even with the SuperPOD
- Less than 14M images for pre-training? Use ResNet.



Transfer Learning From Huge ViT

	ViT-H	Previous SOTA
ImageNet	88.55	88.5
ImageNet-Real	90.72	90.55
Cifar-10	99.50	99.37
Cifar-100	94.55	93.51
Pets	97.56	96.62
Flowers	99.68	99.63



# Self-Supervised Learning

From  
Yoshua Bengio

## Three challenges for Deep Learning

- ▶ Deep Supervised Learning works well for perception
  - ▶ When labeled data is abundant.
- ▶ Deep Reinforcement Learning works well for action generation
  - ▶ When trials are cheap, e.g. in simulation.

### Three problems the community is working on:

1. Learning with fewer labeled samples and/or fewer trials
  - ▶ Self-supervised learning / unsup learning / learning to fill in the blanks
    - ▶ learning to represent the world before learning tasks
2. Learning to reason, beyond "system 1" feed forward computation
  - ▶ Making reasoning compatible with gradient-based learning.
3. Learning to plan complex action sequences
  - ▶ Learning hierarchical representations of action plans




# Self-supervised Learning

- **Problem:** deep learning is not sample efficient
- **Idea:** learn about the world before learning the task
- **New Problem:** how do we learn about the world?
- **Solution:** transfer learning on toy problem
  - 1. train on auxiliary task that is easy to label
  - 2. throw away anything specific to auxiliary task
  - 3. train new network with task of interest, transferring knowledge (downstream task)
  - 4. profit



# Examples of Self Supervised Learning




[https://www.fast.ai/2020/01/13/self\\_supervised/](https://www.fast.ai/2020/01/13/self_supervised/)

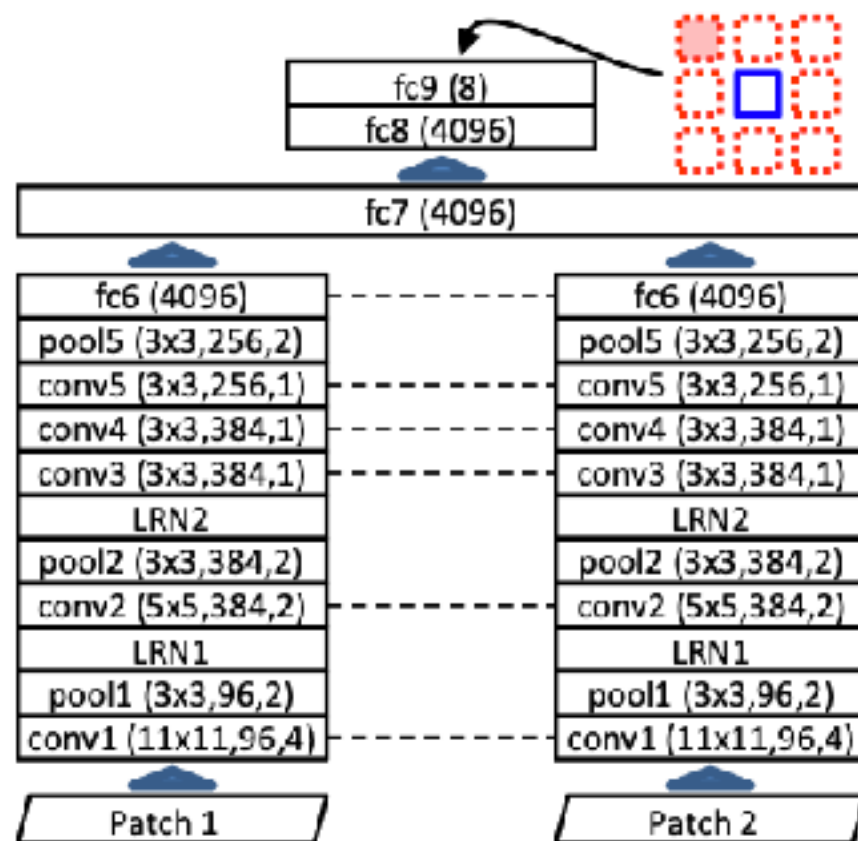
60



# Examples of Self Supervised Learning



$$X = \left( \begin{array}{c} \text{cat face} \\ \text{cat ear} \end{array} \right); Y = 3$$



Unsupervised Visual Representation Learning by Context Prediction

Carl Doersch<sup>1,2</sup>

Abhinav Gupta<sup>1</sup>

Alexei A. Efros<sup>2</sup>


<sup>1</sup> School of Computer Science  
Carnegie Mellon University

<sup>2</sup> Dept. of Electrical Engineering and Computer Science  
University of California, Berkeley






# Examples of Self Supervised Learning



Pretraining Method	Supervision	Pretraining time	Classification	Detection	Segmentation
ImageNet [26]	1000 class labels	3 days	78.2%	56.8%	48.0%
Random Gaussian	initialization	< 1 minute	53.3%	43.4%	19.8%
Autoencoder	-	14 hours	53.8%	41.9%	25.2%
Agrawal <i>et al.</i> [1]	egomotion	10 hours	52.9%	41.8%	-
Wang <i>et al.</i> [39]	motion	1 week	58.7%	47.4%	-
Doersch <i>et al.</i> [7]	relative context	4 weeks	55.3%	46.6%	-
Ours	context	14 hours	56.5%	44.5%	30.0%



Doesn't always work to increase performance...

Context Encoders: Feature Learning by Inpainting



# Consistency Loss

I'm from Canada, but live in the States now.

It took me a while to get used to writing boolean variables with an "Is" prefix, instead of the "Eh" suffix that Canadians use when programming.

For example:

```
MyObj.IsVisible
```

```
MyObj.VisibleEh
```



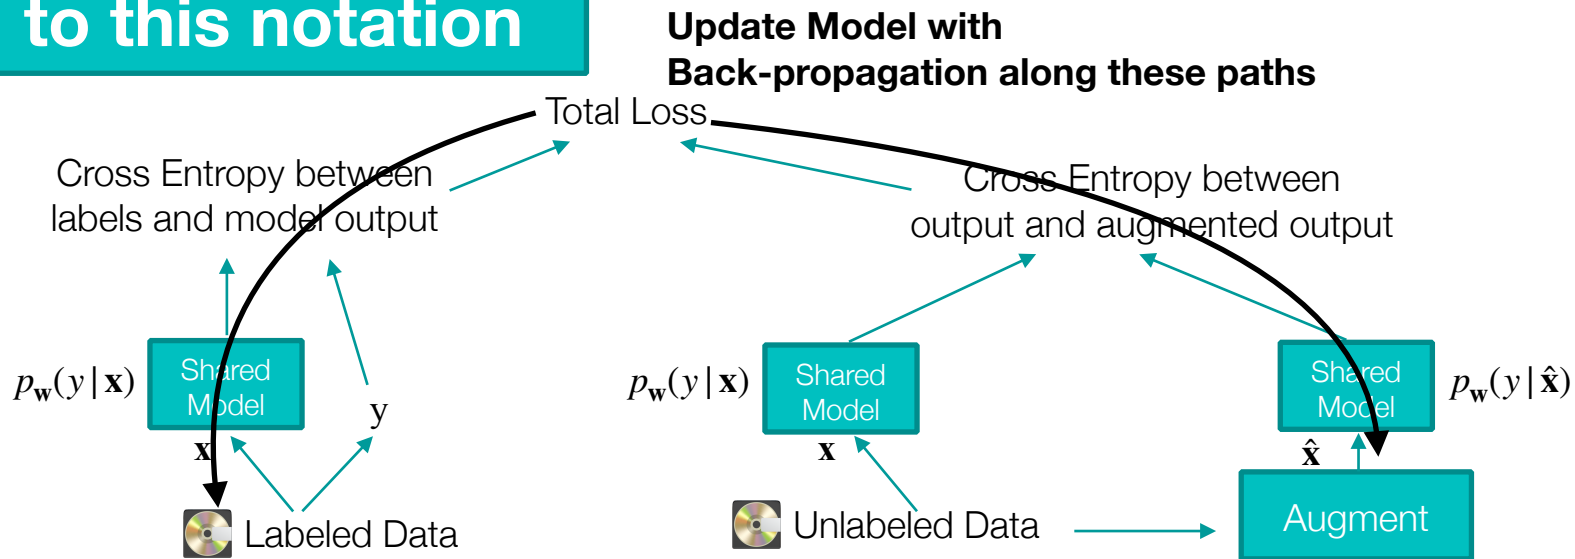
# Unsupervised Consistency Loss

$$\min_{\mathbf{w}} \underbrace{\mathbf{E}_{\mathbf{x}, y \in L} [-\log p_{\mathbf{w}}(y | \mathbf{x})]}_{\text{cross entropy}} + \lambda \underbrace{\mathcal{D}_{KL}(p_{\mathbf{w}}(y | \mathbf{x}) || p_{\mathbf{w}}(y | \hat{\mathbf{x}}))}_{\substack{\text{consistency in augmentation} \\ \text{no back prop} \quad \text{yes back prop}}}$$

Neural Network approximates  $p(y|\mathbf{x})$  by  $\mathbf{w}$   
Use labeled data to minimize network

Sample new  $\mathbf{x}$  from unlabeled pool with function  $q$   
function  $q$  is augmentation procedure  
Minimize cross entropy of two models

**Get accustomed  
to this notation**





# Unsupervised Consistency Loss

$$\min_{\mathbf{w}} \underbrace{\mathbb{E}_{\mathbf{x}, y \in L} [-\log p_{\mathbf{w}}(y | \mathbf{x})]}_{\text{cross entropy}} + \lambda \underbrace{\mathcal{D}_{KL}(p_{\mathbf{w}}(y | \mathbf{x}) || p_{\mathbf{w}}(y | \hat{\mathbf{x}}))}_{\text{consistency in augmentation}}$$

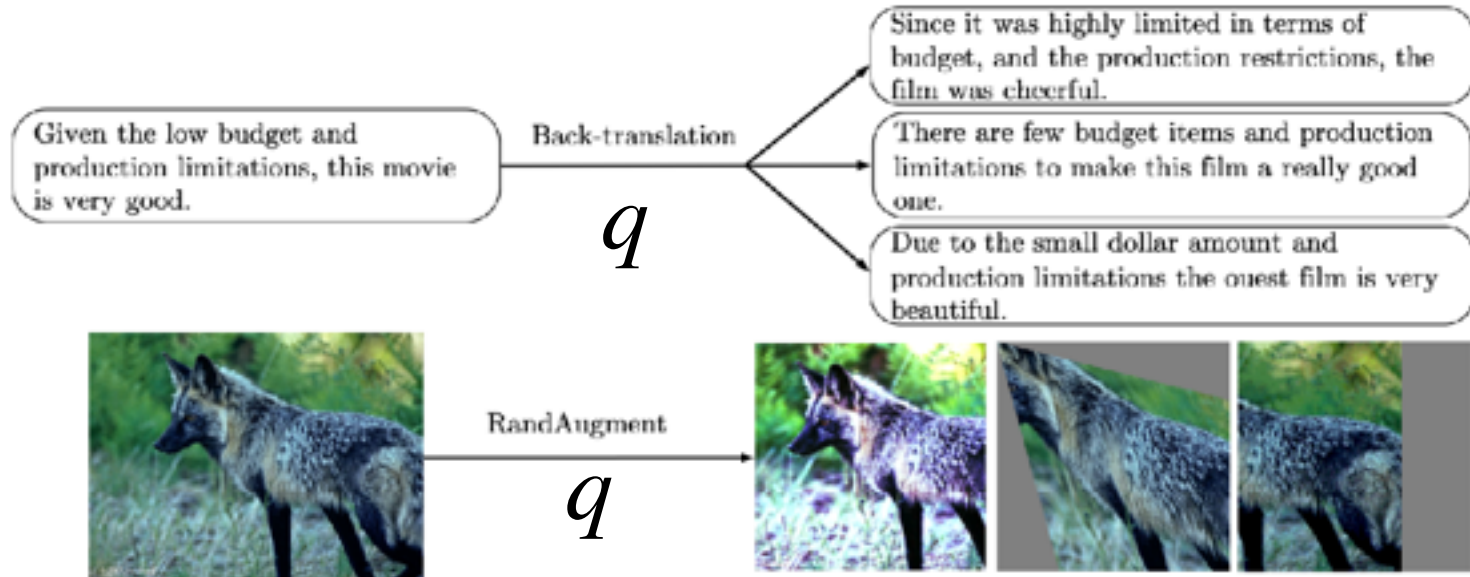


Figure 2: Augmented examples using back-translation and RandAugment.



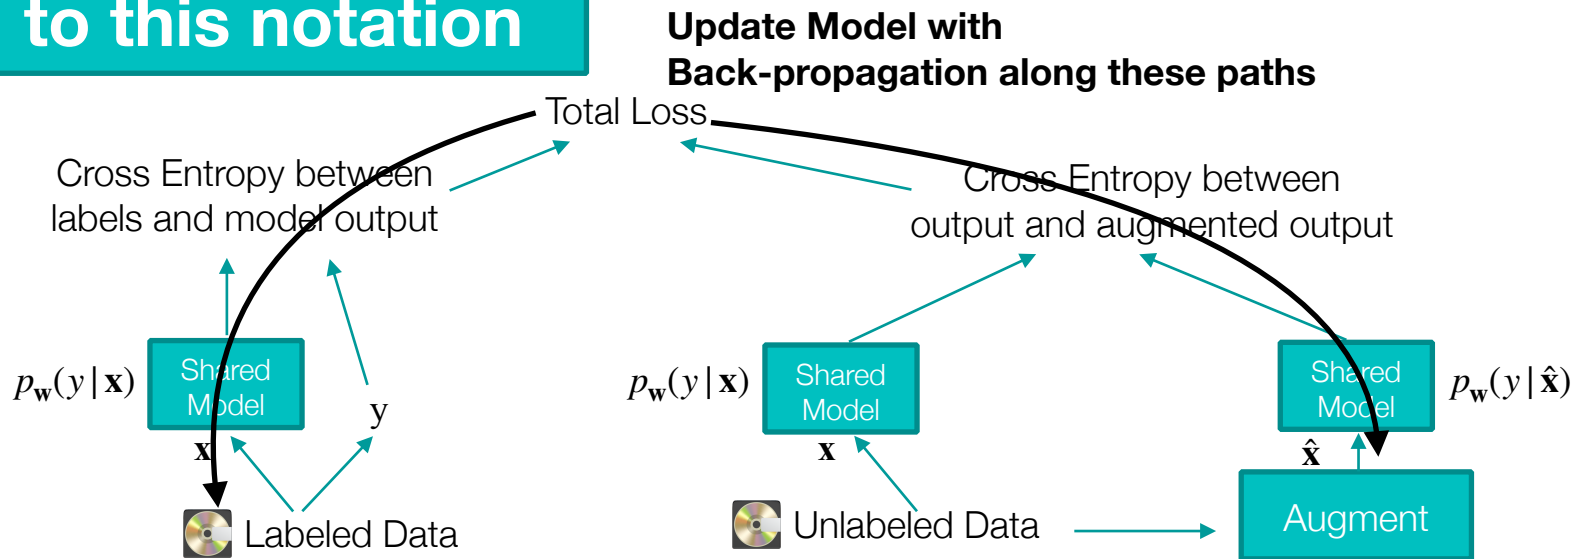
# Unsupervised Consistency Loss (review)

$$\min_{\mathbf{w}} \underbrace{\mathbf{E}_{\mathbf{x}, y \in L} [-\log p_{\mathbf{w}}(y | \mathbf{x})]}_{\text{cross entropy}} + \lambda \underbrace{\mathcal{D}_{KL}(p_{\mathbf{w}}(y | \mathbf{x}) || p_{\mathbf{w}}(y | \hat{\mathbf{x}}))}_{\substack{\text{consistency in augmentation} \\ \text{no back prop} \quad \text{yes back prop}}}$$

Neural Network approximates  $p(y|\mathbf{x})$  by  $\mathbf{w}$   
Use labeled data to minimize network

Sample new  $\mathbf{x}$  from unlabeled pool with function  $q$   
function  $q$  is augmentation procedure  
Minimize cross entropy of two models

**Get accustomed  
to this notation**



$$\min_{\mathbf{w}} \underbrace{\mathbf{E}_{\mathbf{x}, y \in L} [-\log p_{\mathbf{w}}(y | \mathbf{x})]}_{\text{cross entropy}} + \lambda \underbrace{\mathcal{D}_{KL}(p_{\mathbf{w}}(y | \mathbf{x}) || p_{\mathbf{w}}(y | \hat{\mathbf{x}}))}_{\text{consistency in augmentation}}$$


---

$$E[g] = \sum p(g) \cdot g \quad \text{definition of expected value}$$

$$E[-\log p_{\mathbf{w}}(y | \mathbf{x})] = - \sum p(y) \cdot \log p_{\mathbf{w}}(y | \mathbf{x}) \quad \text{insert -log probability, log likelihood}$$

$$NLL(y, p_{\mathbf{w}}(y | \mathbf{x})) = - \sum_c p(y = c) \cdot \log p_{\mathbf{w}}(y = c | \mathbf{x}) \quad \text{negative log likelihood}$$

$$CE(f, g) = - \sum f(x) \cdot \log g(x) \quad \text{cross entropy of two functions}$$

$$CE(y, p_{\mathbf{w}}(y | \mathbf{x})) = - \sum_c (y = c) \cdot \log p_{\mathbf{w}}(y = c | \mathbf{x}) \quad \text{if } y=c \text{ is a probability, these are same equation}$$

```
cce = tf.keras.losses.CategoricalCrossentropy()
cce(y_true, y_pred)
```

