

Lecture Notes for **Neural Networks and Machine Learning**



CNN Visualization



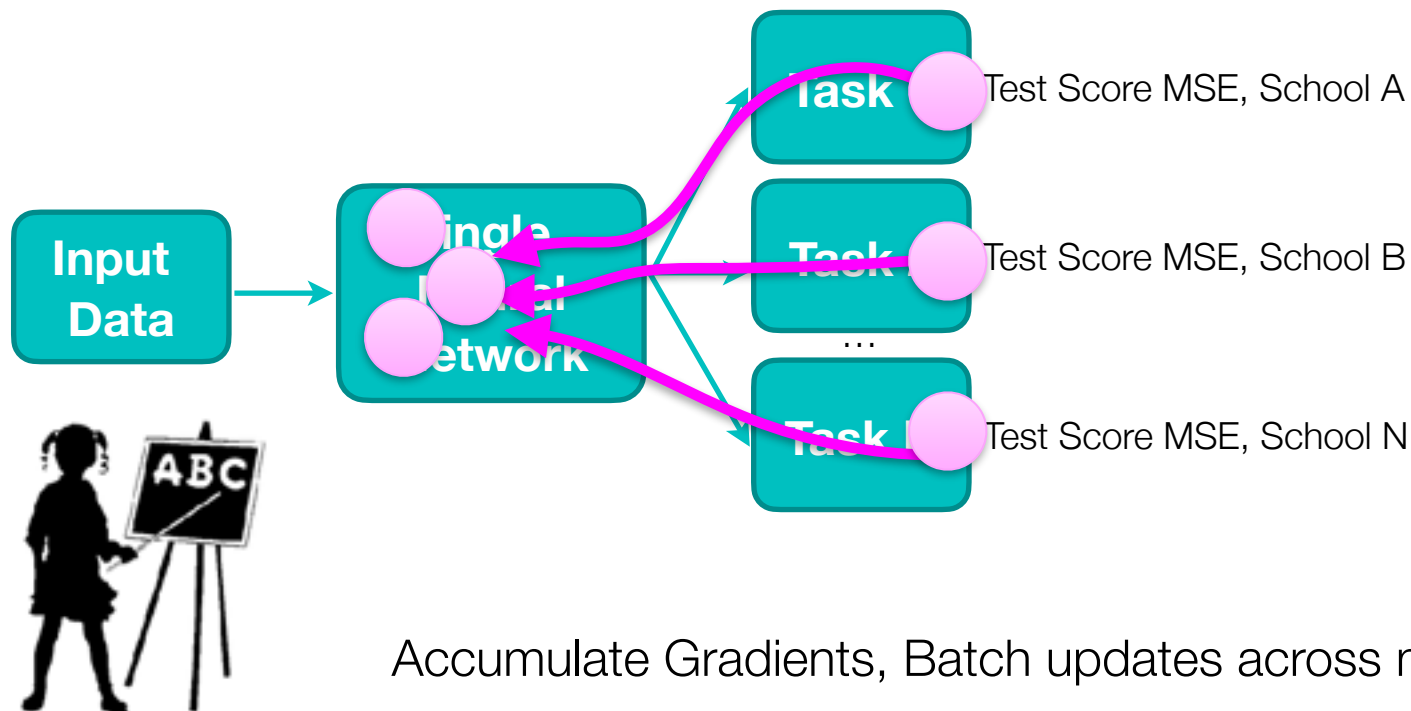
Logistics and Agenda

- Logistics
 - None
- Agenda
 - Multi-Task Demo, Revisit
 - Visualizing Convolutional Architectures and Demo
- Next Week:
 - Circuits in CNNs

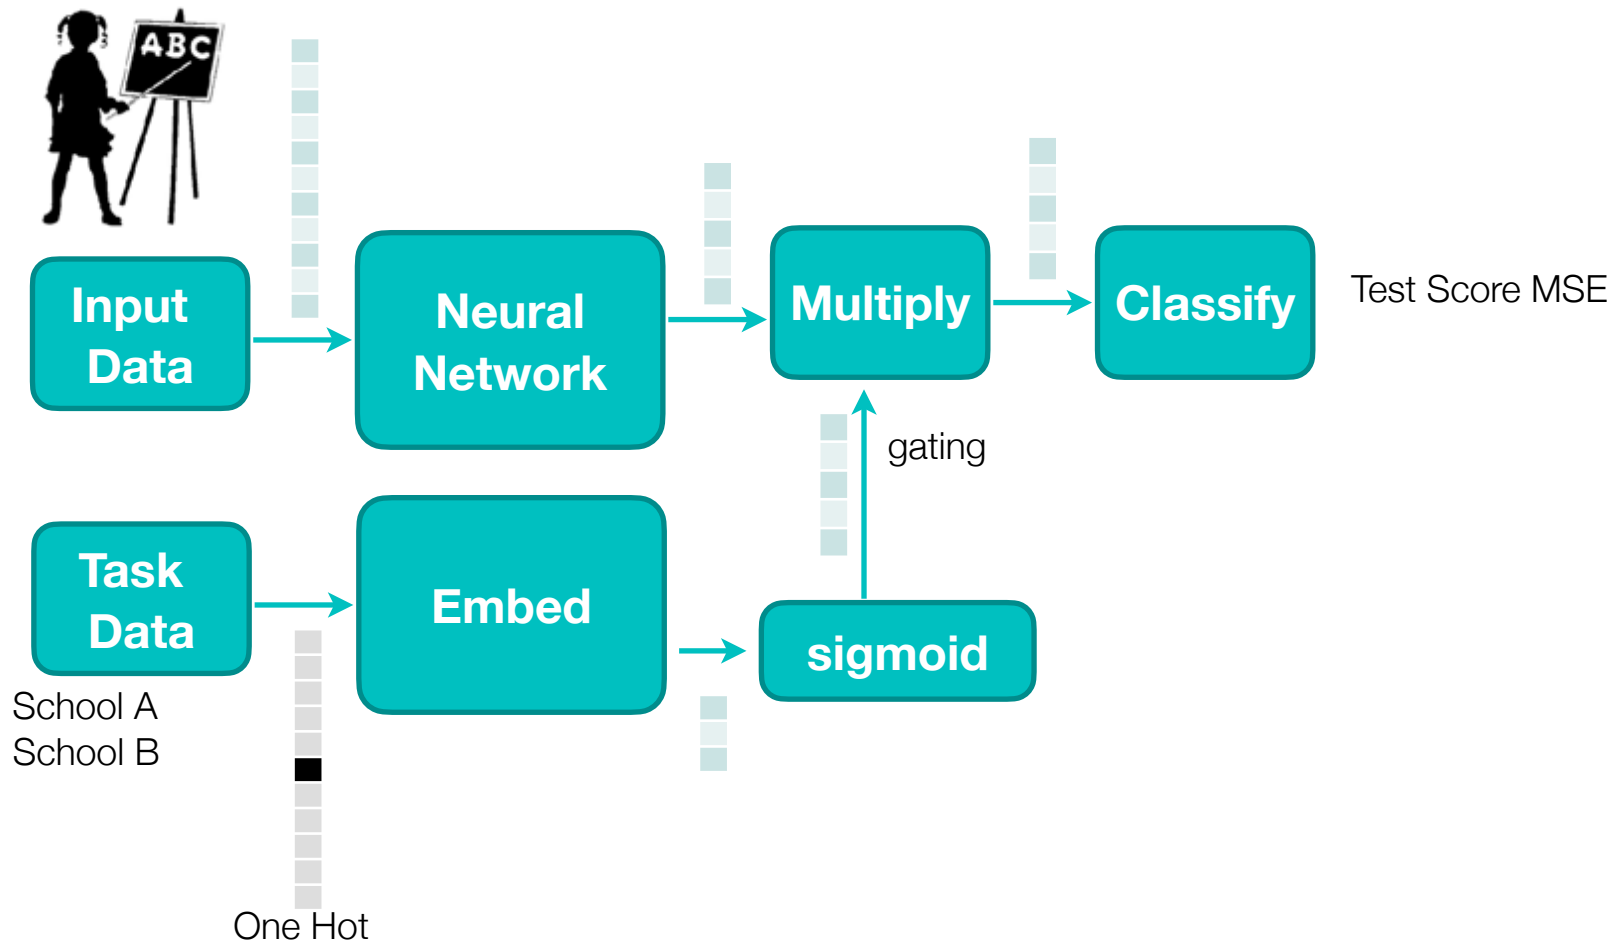


Multi-task Optimization, Review

Single Task Label per Input



An alternative: Task-Gating, Review





Multi-Task Learning

School Data, Computer Surveys



Traian-Pop Traian Pop



LukeWood Luke Wood

KerasCV Author, Full Time Keras team member & Machine Learning researcher @ Google, Part Time UCSD Ph.D student



♡ Sponsor

Follow

Testing out a new idea: structured demonstration with mix of code and pseudocode...

[LectureNotesMaster/03](#) [LectureMultiTask.ipynb](#)

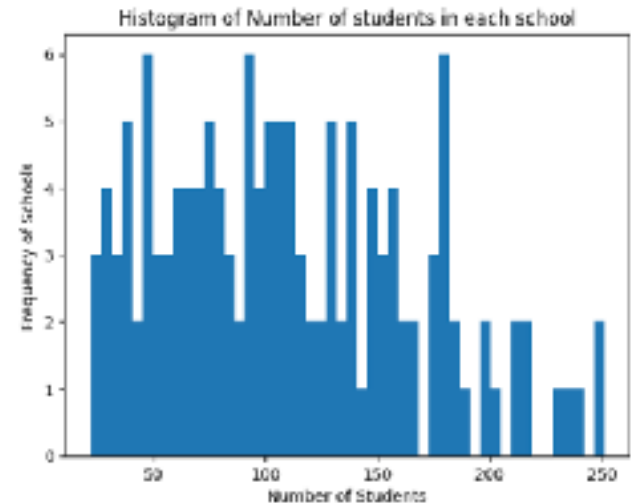


Demonstration Overview

- Scale features and outputs
- Organize schools into dictionaries of Bunch data:
- Baseline: one model per school (limited training)
- Baseline: one model overall (nothing per school)
- Multi-task: shared model with school specific layer
- Gated: one model with task gating

```
feature_scaler = StandardScaler()  
output_scaler = StandardScaler()  
  
X = feature_scaler.fit_transform(X)  
y = output_scaler.fit_transform(y)
```

```
sid = "School {}".format(i + 1)  
tasks[sid] = Bunch(data=X[start:end],  
                   target=y[start:end],  
                   DESCR=descr)
```



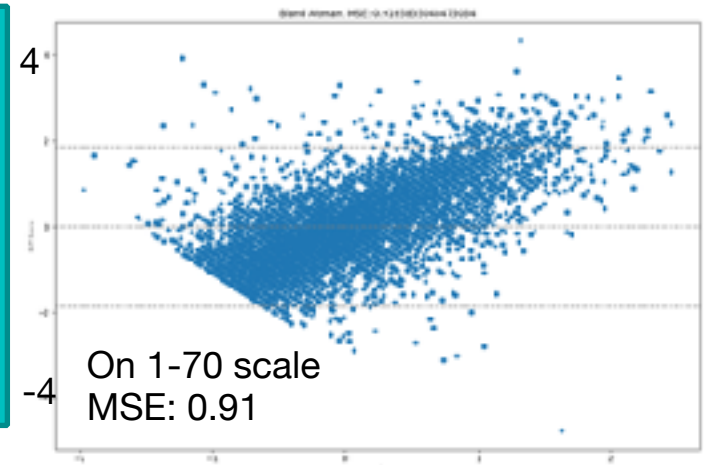
Baseline Modeling

```
for sid in tasks.keys():  
  
    mlp = Sequential()  
    mlp.add( ... layers ...)   
    mlp.compile( ...loss and optimizer... )  
    mlp.fit(X_train[sid], y_train[sid], ... )  
  
    # save the output results  
    yhat_mlp[sid] = mlp.predict(X_test[sid])
```


X_train_all, X_test_all = ... gather all data ...

```
mlp = Sequential()  
mlp.add( ... layers ...)   
mlp.compile( ...loss and optimizer... )  
mlp.fit(X_train_all, y_train_all, ... )  
  
# save the output results  
yhat_mlp_all = mlp.predict(X_test_all)
```

Per school model



Single Model



Multi-task Modeling, setup

```
w1, w2, w_output = ... get general model weights ...

inputs = Input(... single input to CG ...)
shared_input = Dense(... trainable=False)(inputs)
shared_mlp = Dense(..., trainable=False)(shared_input)
models = dict() # models for each task

for sid in tasks.keys():
    output_layer = Dense(1, ...)(shared_mlp) # new layer

    models[sid] = Model(inputs=inputs, outputs=output_layer)


    shared1 = models[sid].get_layer('shared_input')
    shared2 = models[sid].get_layer('shared_middle')
    personal = models[sid].layers[-1]

    # set weights from the general model, as starting point
    shared1.set_weights(w1)
    shared2.set_weights(w2)
    personal.set_weights(w_output)

    personal.trainable = True
```


Shared Layers

Multi-task Layers



Multi-task Modeling, setup

```
def step(keys, opt):
    # each key will be a separate school
    loss = {}, tapes = {}
    for sid in keys:
        with tf.GradientTape() as tape:
            # accumulate all the gradient updates for each school
            # make a prediction and calculate loss for each task
            tapes[sid] = tape # need to track
            preds = models[sid]( X_train[sid] )
            loss[sid] = mean_squared_error( y_train[sid], preds )
            Could weight the loss here to account for schools with low enrollment
    # now batch update all the models with the gradients
    for sid in keys:
        grads = tapes[sid].gradient(loss[sid],
                                    models[sid].trainable_variables)
        opt[sid].apply_gradients(zip(grads,
                                    models[sid].trainable_variables))
```



Accumulate MSE Loss

Gradient Accumulation

```
opt = {} # separate optimizers per task
for sid in all_keys:
    opt[sid] = Adam()

for i in range(EPOCHS):
    shuffle(all_keys) # shuffle
    step(all_keys, opt) # optimize
```

Apply Fitting

```
layer1, layer2 = ... shared layers ...
layer1.trainable = False
layer2.trainable = True

for i in range(EPOCHS):
    shuffle(all_keys) # shuffle
    step(all_keys, opt) # optimize
```

Fine tune



Task Gating Model

```
in_feature = Input(... feature data ..., name = 'student_features')
in_school = Input(... school id ..., name = 'school_id')

num_schools = 139, embed_sz = 32
x_school = Embedding(input_dim=num_schools, output_dim=embed_sz, ...)(in_school)
x_student = Dense(units=embed_size, ... )(in_feature)

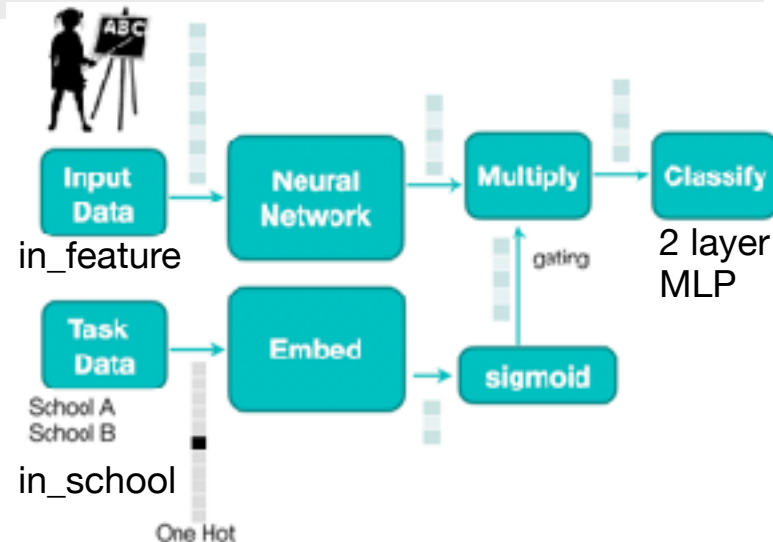
x = Multiply()( [ sigmoid(x_school), x_student ] ) # gating

x = Dense( ... )(x)
x_out = Dense(1, activation='linear', ... )(x)

gated_mlp = Model(inputs=[in_feature, in_school], outputs=x_out)
```

```
gated_mlp.compile( ... )
gated_mlp.fit({'student_features':X_train_all,
              'school_id':s_train_all },
              y_train_all, ... )

y_hat_gated = gated_mlp.predict(
    {'student_features':X_test_all,
     'school_id':s_test_all })
```

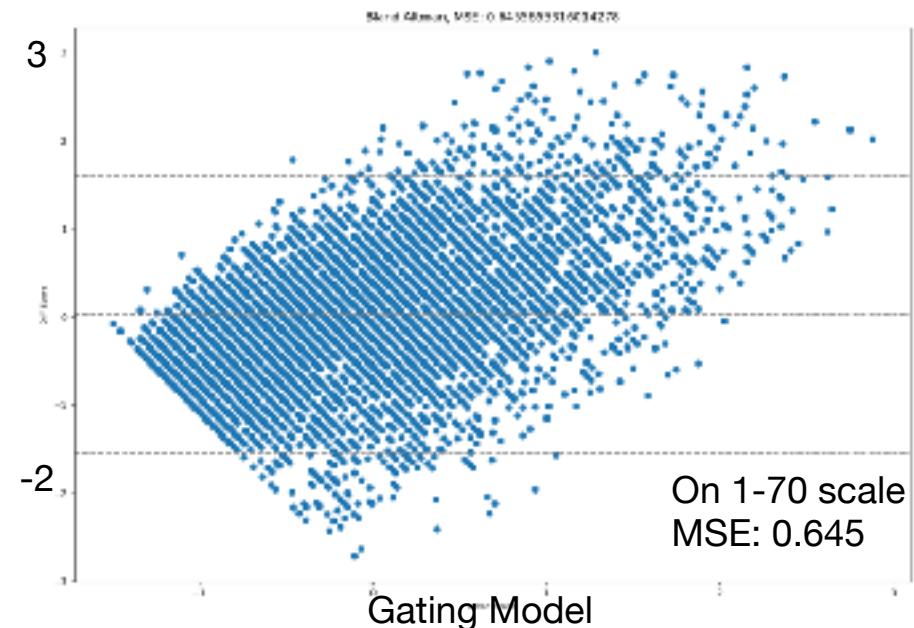



Results (schools and survey data)

<i>School Model</i>	<i>MSE</i>	<i>Training Time</i>
<i>Per School</i>	0.916	3 min
<i>Single Model</i>	0.660	1.5 min
<i>Multi-task</i>	0.649	4 min
<i>Gated</i>	0.645	1 min

Not always as Clear Cut

<i>Survey</i>	<i>MSE</i>	<i>Train Time</i>
<i>Per School</i>	21.70	3 hours
<i>Single Model</i>	9.03	15 min
<i>Multi-task</i>	6.22	50 min




Demonstration Comments?

- Pros/Cons



Basics of Convolutional Neural Network Visualization



Tools to Visualize Neurons and Filters


- Visualize **Filter Activation**
 - What parts of the inputs activate each filter?
- Visualize **Filters**
 - What does each filter look like? Is it similar to other filters?
 - Can we excite a certain filter by updating the input image?
- **Heatmaps** of Class Activation
 - What part of an input image most influences each final output?



Visualizing Intermediate Activations

Method
One


- Look layer by layer
- **Assume:** each filter learns something useful




Visualizing Intermediate Activations

Method
One


- **Recall:** general structure of most CNNs
 - Small kernels throughout (3x3)
 - Filtering followed by Pooling (spatial downsampling)
 - More filters in later layers



Early Activations
are larger but not as
numerous




Later Activations are
smaller and more
numerous



Visualizing Intermediate Activations


Method
One

- **Result:** Information Distillation Pipeline
 - Deeper layers have more abstract triggers
 - Deeper activations are increasingly sparse
 - Early layers are texture and edge detectors
 - Notion of “High Level Abstraction,” has biological motivation



Early Activations
are larger but not
as numerous



Later Activations are
smaller and more
numerous



Visualizing Filters: Class Neuron

Method
Two

- **Idea:** What Maximally Activates a Class Output?
 - Gradient Ascent in the Input Space



Flamingo

where c is a specific neuron in output layer
 f is the neural network function
 I is the input image, init to zeros (or random)
 ∇ is the gradient of f_c w.r.t I
CNN weights stay unchanged

<http://cs231n.github.io/understanding-cnn/> 18



Visualizing Filters: Maximal Activations


Method
Two

- **Idea:** What Maximally Activates a **Filter**?
 - **Again:** Gradient Ascent in the Input Space

$$I \leftarrow I + \eta \sum_{i,j} \nabla f_n(I)_{i,j}$$

“trick” use norm of gradient

where n is a specific **filter** in a layer
 f is the function to n^{th} filter in layer






Visualizing ConvNets


Part One: Filter Activations

Part Two: Image Gradients



Ian Stenbit

Google AI



Follow Along: `04 LectureVisualizingConvnets.ipynb`
`activation-demo`



Class Activation Mapping (CAM)

- **Idea:** What areas of the image contributed most to the classification result?
- Also, for each class, what areas of the image exhibit features of that class?
- Use change in output, w.r.t. final conv layer

normalize by $h \times w$ of A

final layer output in response to image I
 c is class of interest

$$\alpha_k^c = \frac{1}{|A_k^{(L)}|} \sum_{i,j} \frac{\partial f_c(I)}{\partial A_{i,j,k}^{(L)}}$$

final convolutional layer, L , activations for row, column, channel

gradient weight for channel k and class c in layer L
 k in $1 \dots K$ activations in final layer




Class Activation Mapping (CAM)

$$\alpha_k^c = \frac{1}{|A_k^{(L)}|} \sum_{i,j} \frac{\partial f_c(I)}{\partial A_{i,j,k}^{(L)}}$$

α_k^c : gradient weight for channel k and class c in layer L
 k in $1 \dots K$ activations in final layer

$\frac{\partial f_c(I)}{\partial A_{i,j,k}^{(L)}}$: final layer output in response to image I
 c is class of interest
 $A_{i,j,k}^{(L)}$: final convolutional layer, L , activations for row, column, channel



Sensitivity of Class to Activations



Class Activation Mapping (CAM)

$$\alpha_k^c = \frac{1}{|I \times J|} \sum_{i,j} \frac{\partial f_c(I)}{\partial A_{i,j,k}^{(L)}}$$

α_k^c : gradient weight for channel k and class c in layer L
 k in $1 \dots K$ activations in final layer


$\frac{\partial f_c(I)}{\partial A_{i,j,k}^{(L)}}$: final layer output in response to image I
 c is class of interest

$A_{i,j,k}^{(L)}$: final convolutional layer, L , activations for row, column, channel

Heatmap, S , is the weighted sum of final layer activations:

$$S_{i,j} = \frac{1}{S_{max}} \sum_k \phi(\alpha_k^c A_{i,j,k}^{(L)})$$

ϕ : relu activation






Visualizing ConvNets

Part Three: Grad-CAM



Ian Johnson



Follow Along: 04 LectureVisualizingConvnets.ipynb
activation-demo



Lecture Notes for Neural Networks and Machine Learning

CNN Visualization



Next Time:
CNN Circuits

Reading: OpenAI Circuits

