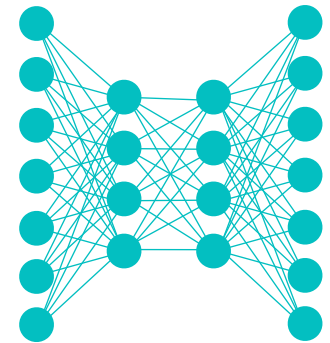


# Lecture Notes for **Neural Networks and Machine Learning**



## Fully Convolutional Learning II: Object Detection



# Logistics and Agenda

- Logistics
  - Lab grading update
- Agenda
  - Lab Town Hall Review and Final Project (if needed)
  - Upsampling
  - Full Convolutional Architectures
    - ◆ Semantic Segmentation Basics (last time)
    - ◆ Object Detection (this time):
      - RCNN, YOLO
    - ◆ Instance Segmentation (next time, probably):
      - Mask-RCNN, YOLACT



# Town Hall Revisited



# Final Project

Me: Predicts the next word correctly



- I have biometric data for pilots...
- Perhaps something to use for the final project...



# Basics: Upsampling Layers



**Shit Academics Say** @Academi... · 22h ...  
not wrong




**monstera adansonii** @yourn... · 2d


everything is peer reviewed if your  
friends are judgmental enough



# Last Time



# Decoder Network



Some researcher started calling this **deconvolution**.


If you use that term in this class, **you fail**.

This is upsampling and then convolution, but **now the interpolation filters are learned!!**





# Integer Upsampling via Interpolation

**Nearest Neighbor**





All are equivalent to inserting zeros and applying convolutional filter

**Linear**



**Cubic**





# Image Upsampling, Integer Factor

- Insert Zeros
- Convolve


1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16



1		2		3		4	
5		6		7		8	
9		10		11		12	
13		14		15		16	

0.25	0.5	0.25
0.5	1	0.5
0.25	0.5	0.25

Bilinear Filtering

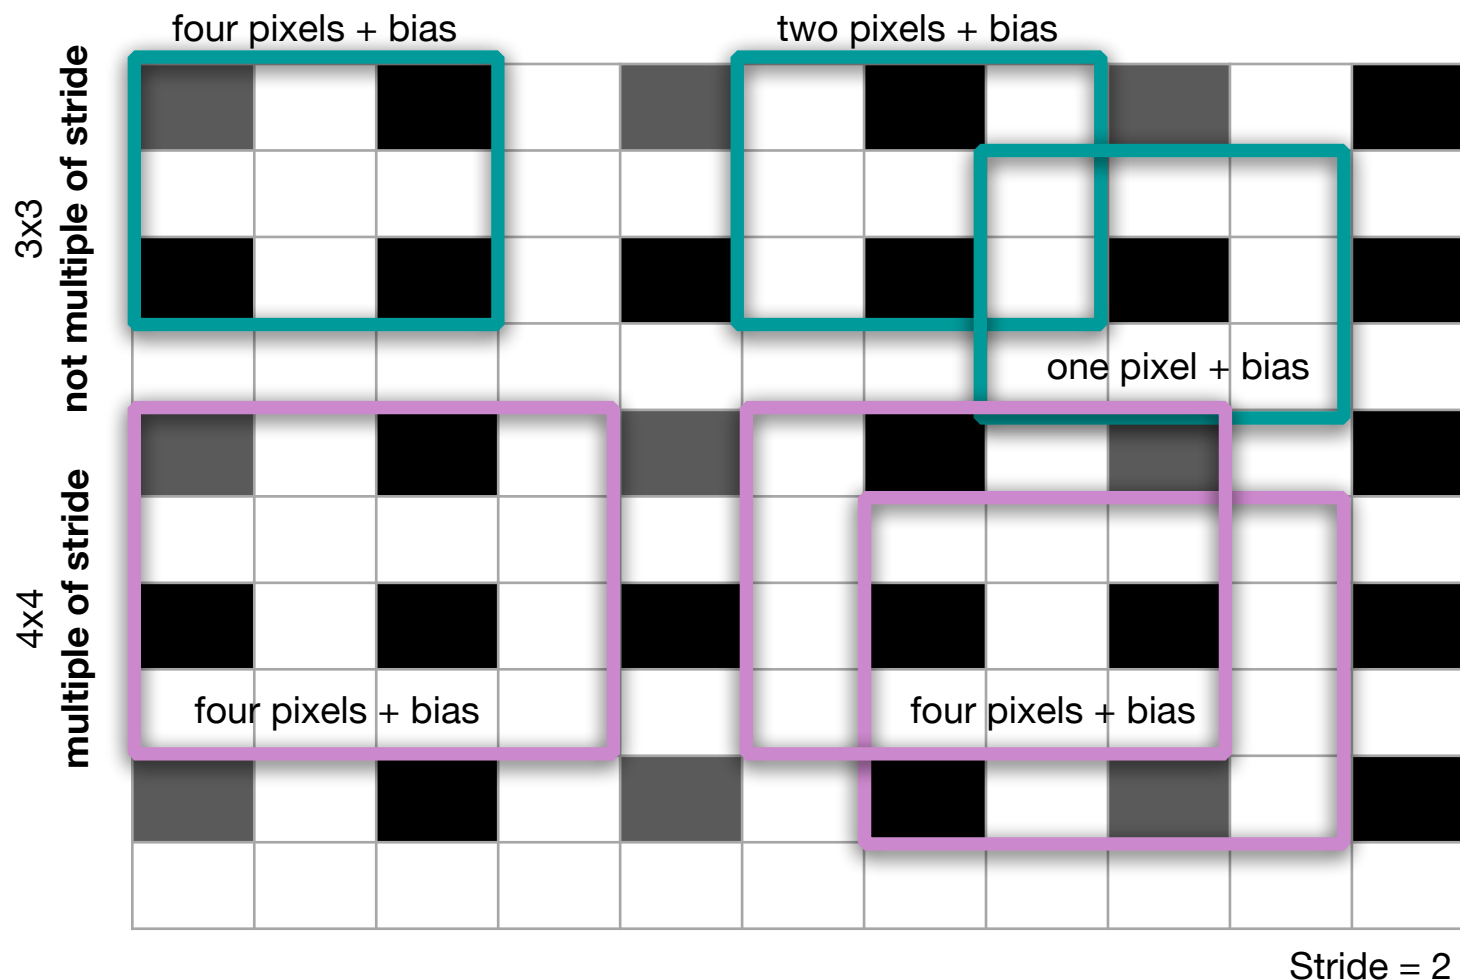


Bicubic Filter



# Learned Upsampling after Zero Insertion

- Learning the interpolation filter has some caveats:



Bias needs to account for both when different numbers of pixels overlap with the kernel

Multiple of stride ensures that same number of active pixels overlap the kernel.



# Image Upsampling, Integer Factor



**Nearest Neighbor**

`UpSampling2D()`




**Bilinear**

`UpSampling2D(interpolation='bilinear')`



**Bicubic**



**Many Types of Upsampling,  
with varying computational  
cost:**

area, bicubic, gaussian,  
lanczos3, lanczos5,  
mitchellcubic



# What about transpose convolution?

Convolution as Matrix Multiplication

$y$	$x$	$0$	$0$	$0$
$z$	$y$	$x$	$0$	$0$
$0$	$z$	$y$	$x$	$0$
$0$	$0$	$z$	$y$	$x$
$0$	$0$	$0$	$z$	$y$

 $\times$ 

$0$
$a$
$b$
$c$
$0$

 $=$ 

$ax$
$ay+bx$
$az+by+cx$
$bz+cy$
$cz$

Transpose

$y$	$z$	$0$	$0$	$0$
$x$	$y$	$z$	$0$	$0$
$0$	$x$	$y$	$z$	$0$
$0$	$0$	$x$	$y$	$z$
$0$	$0$	$0$	$x$	$y$

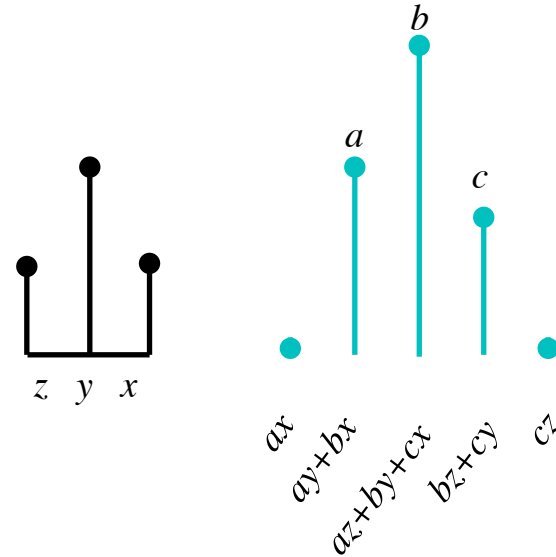
 $\times$ 

$0$
$a$
$b$
$c$
$0$


 $=$ 

$az$
$ay+bz$
$ax+by+cz$
$bx+cy$
$cx$

like convolving with “reversed coefficients”



Regular Convolution



Transpose Convolution



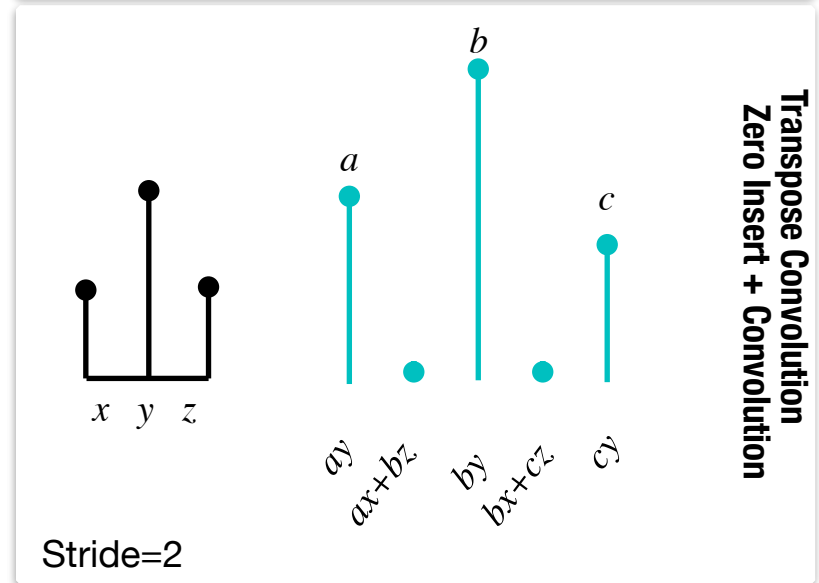

# Transpose Convolution: Strides

Strided Convolution as Matrix Multiplication

$$\begin{bmatrix} y & x & 0 & 0 & 0 \\ 0 & z & y & x & 0 \\ 0 & 0 & 0 & z & y \end{bmatrix} \times \begin{bmatrix} 0 \\ a \\ b \\ c \\ 0 \end{bmatrix} = \begin{bmatrix} ax \\ az+by+cx \\ cz \end{bmatrix}$$



Transpose

$$\begin{bmatrix} y & 0 & 0 \\ x & z & 0 \\ 0 & y & 0 \\ 0 & x & z \\ 0 & 0 & y \end{bmatrix} \times \begin{bmatrix} a \\ b \\ c \end{bmatrix} = \begin{bmatrix} ay \\ ax+bz \\ by \\ bx+cz \\ cy \end{bmatrix}$$



# This time... Object Detection Methods

- Semantic segmentation has good mIoU values (up to 90%) but this is exaggerated by background recognition, many classes are  $<40\%$
- How to adapt these techniques to get bounding boxes, not semantic segmentations?
  - Could this be easier? More stable?
  - More consistent labeling?
  - Suitable for “higher risk” tracking applications?



# Object Detection with RCNN



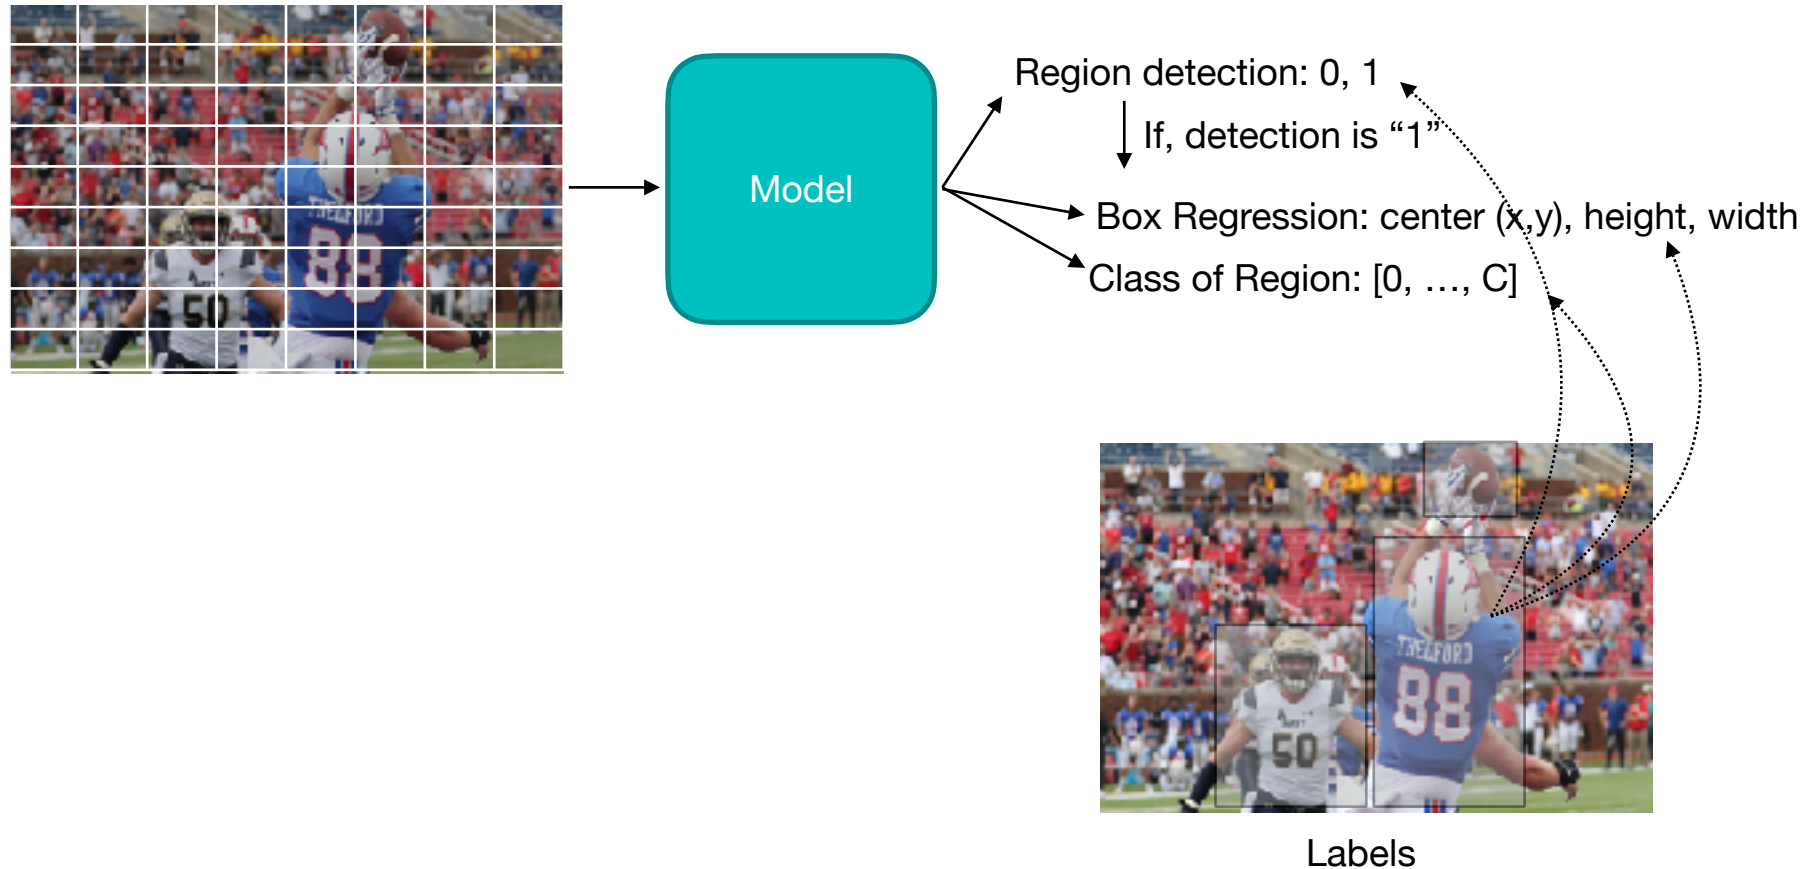
**A history in naming one network five different times  
with five different papers  
each time changing one thing about the architecture**

## Research!






# General Structure





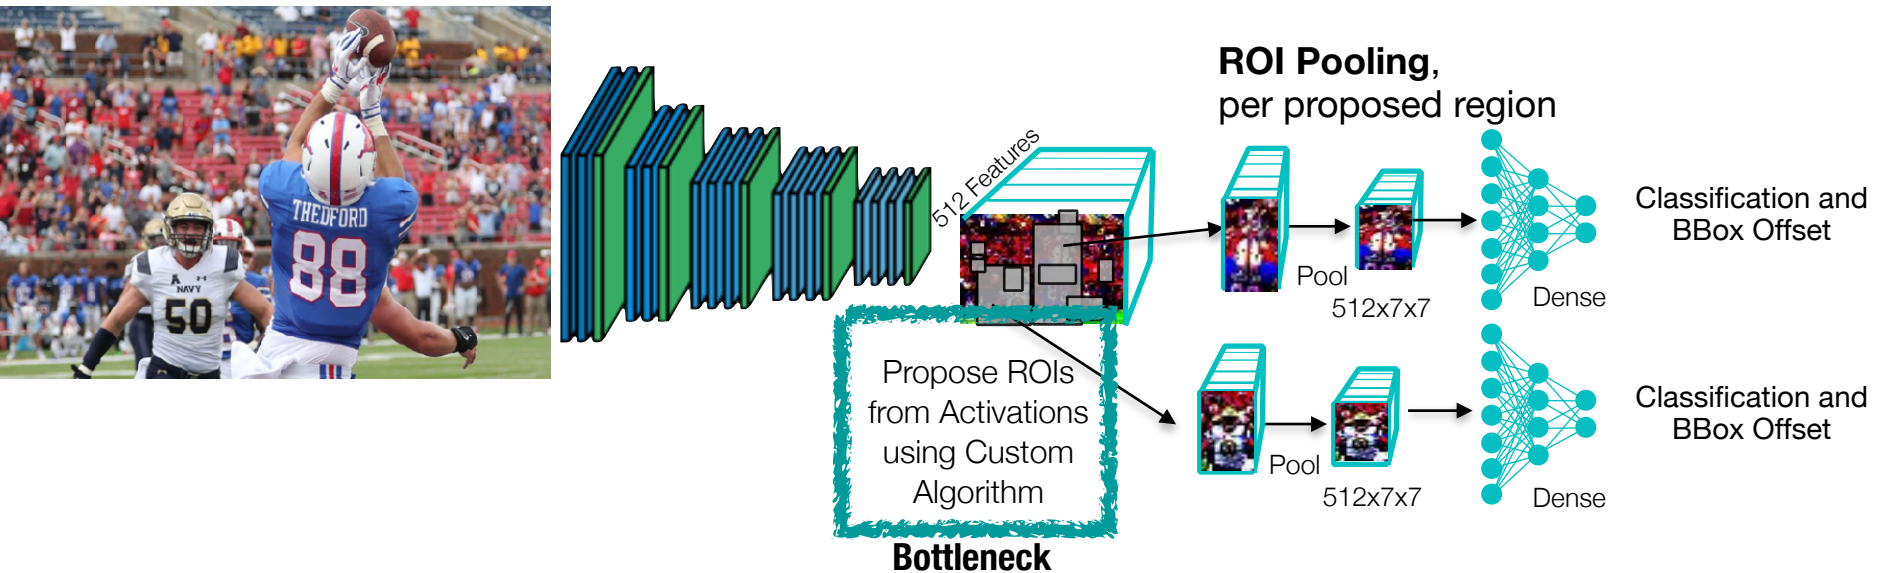
# 2014: R-CNN



- Too Slow to Be Useful
- SVM and BBox Regression Trained Separately
- Fine Tuned Existing ConvNet (for Warped Images)
- ~50 Seconds per Image when Deployed




# 2015: Fast R-CNN




- Fast! 2.3 seconds per image (not ~50)
- But still not real time...



# 2015: Faster R-CNN

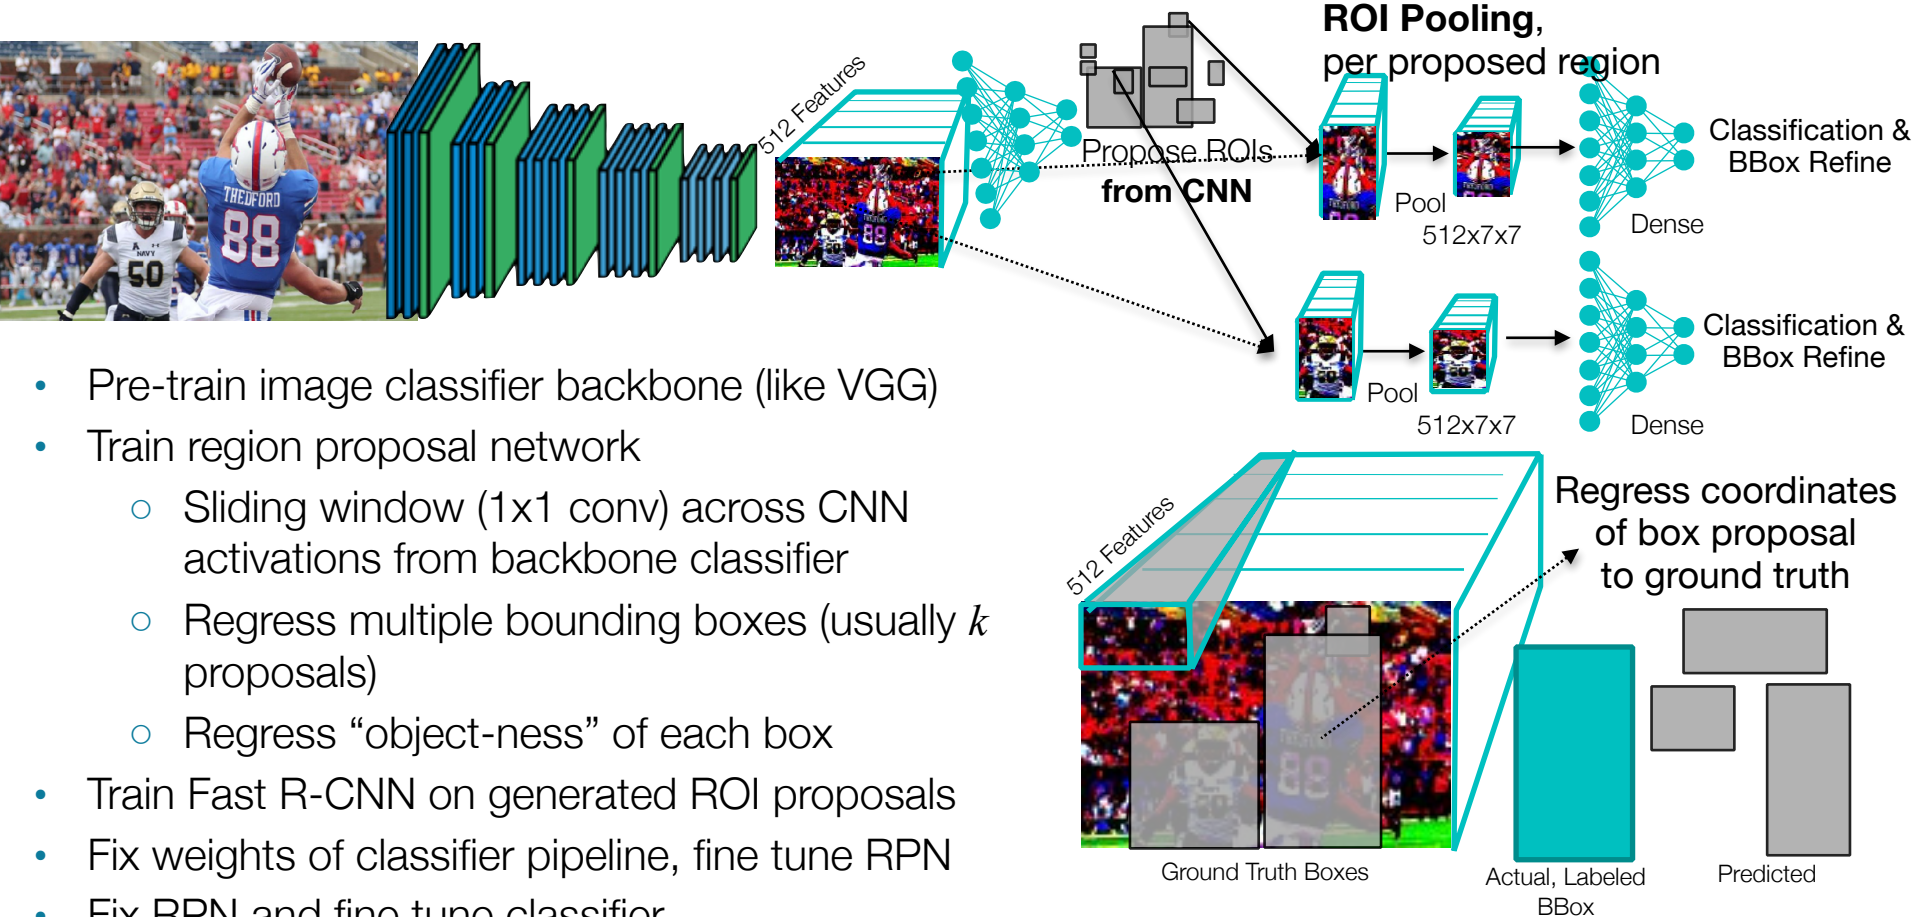


- **Faster!** 140ms per image (7 FPS)
- **Highly Accurate**





# 2015: Faster R-CNN, Training



- Pre-train image classifier backbone (like VGG)
- Train region proposal network
  - Sliding window (1x1 conv) across CNN activations from backbone classifier
  - Regress multiple bounding boxes (usually  $k$  proposals)
  - Regress “object-ness” of each box
- Train Fast R-CNN on generated ROI proposals
- Fix weights of classifier pipeline, fine tune RPN
- Fix RPN and fine tune classifier
- Rinse, repeat fine tuning

$$l_{box} = \sum_i \hat{p}_i \left[ (x - \hat{x}_i)^2 + (y - \hat{y}_i)^2 + (\log w - \log \hat{w}_i)^2 + (\log h - \log \hat{h}_i)^2 \right]$$

$$l_{class} = \sum_c CE(c, \hat{c})$$

$$l_{obj} = \sum_i CE(p_i, \hat{p}_i)$$

