

Lecture Notes for **Neural Networks and Machine Learning**



Introduction to
Reinforcement Learning




Logistics and Agenda

- Logistics
 - Student Presentation Next Time
- Agenda
 - Final Projects (if needed)
 - Basics of Reinforcement Learning
 - Markov Processes
 - Reinforcement Learning Categorization
 - OpenAI Gym
 - The Cross Entropy Method

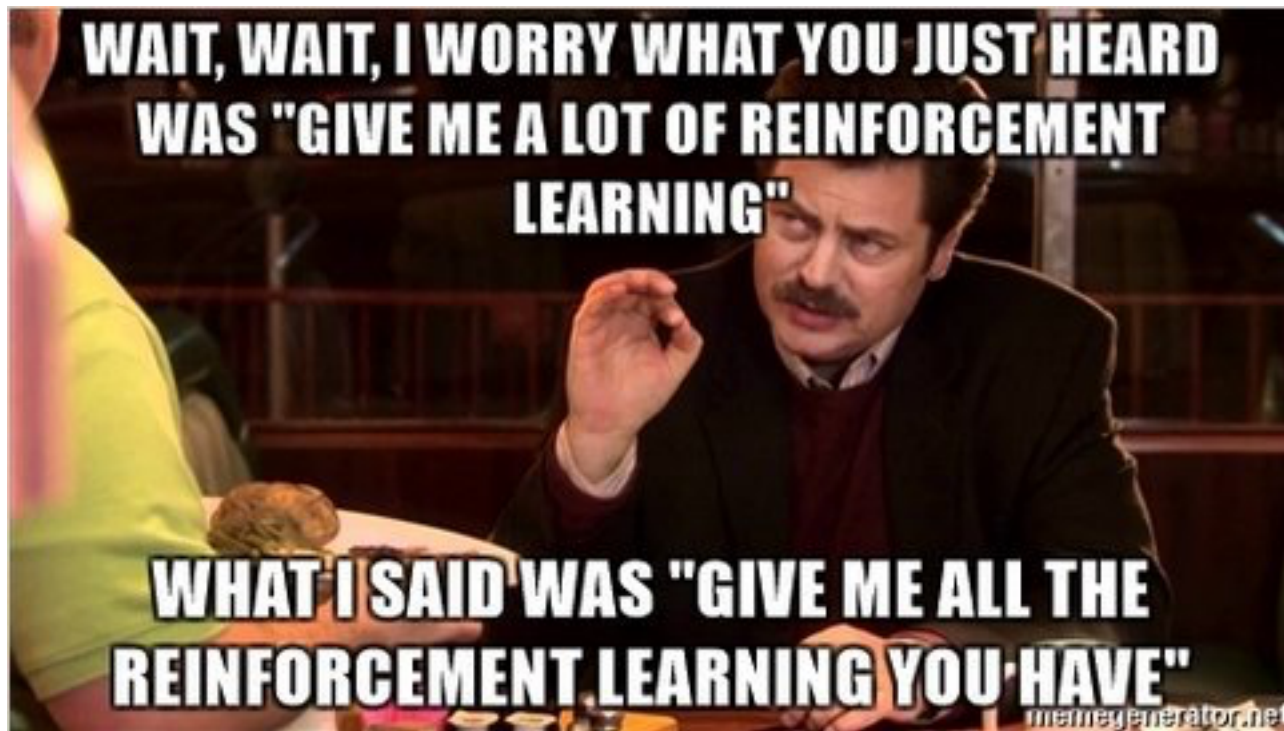


Final Project

One Idea from Professor Stephanie Langin-Hooper SMU Meadows



Reinforcement Learning Basics



History of RL from Two Paths

- **Optimal Control**

- Model processes via Markov property
- Optimal paths through states calculated through dynamic programming

- **Animal Behavioral Learning (psychology)**

- Animals learn by trial and error
- Formalized by Thorndike, 1911. Strengthen through pleasure and weaken through pain
- Pavlov and B.F. Skinner would conduct experiments proving that behavior could be influenced with RL

- **Motivation for many pioneering Researchers:**

Claude Shannon, J. Deutsch, Marvin Minsky, F. Rosenblatt, Widrow, Hoff



Edward Thorndike



B.F. Skinner



Ivan Pavlov



Bernard Widrow



Marvin Minsky



Ted Hoff



Claude Shannon



Conditioning, Skinner and Pavlov

Continuous Reinforcement



Desired behavior is reinforced every time it occurs



Most effective when teaching a new behavior

"SHAKE!"



Creates a strong association between behavior and response

Partial Reinforcement



Most effective once a behavior has been established



New behavior is less likely to disappear



Various partial reinforcement schedules available to suit individual needs



How to condition a machine learning model?

- Hybrid of **Supervised** and **Unsupervised** Learning
- **Reinforcement** Learning
 - Possibly specific labels given, but not necessarily with supervision for how labels are achieved
 - ◆ labels are typically stochastic
 - Uses many techniques from supervised learning, but applied towards a slightly different objective function
 - Rewards (positive and negative) are possible to assess behavior in an environment (just like with animals RL)
 - Not specific to Machine Learning community




Generic RL Landscape

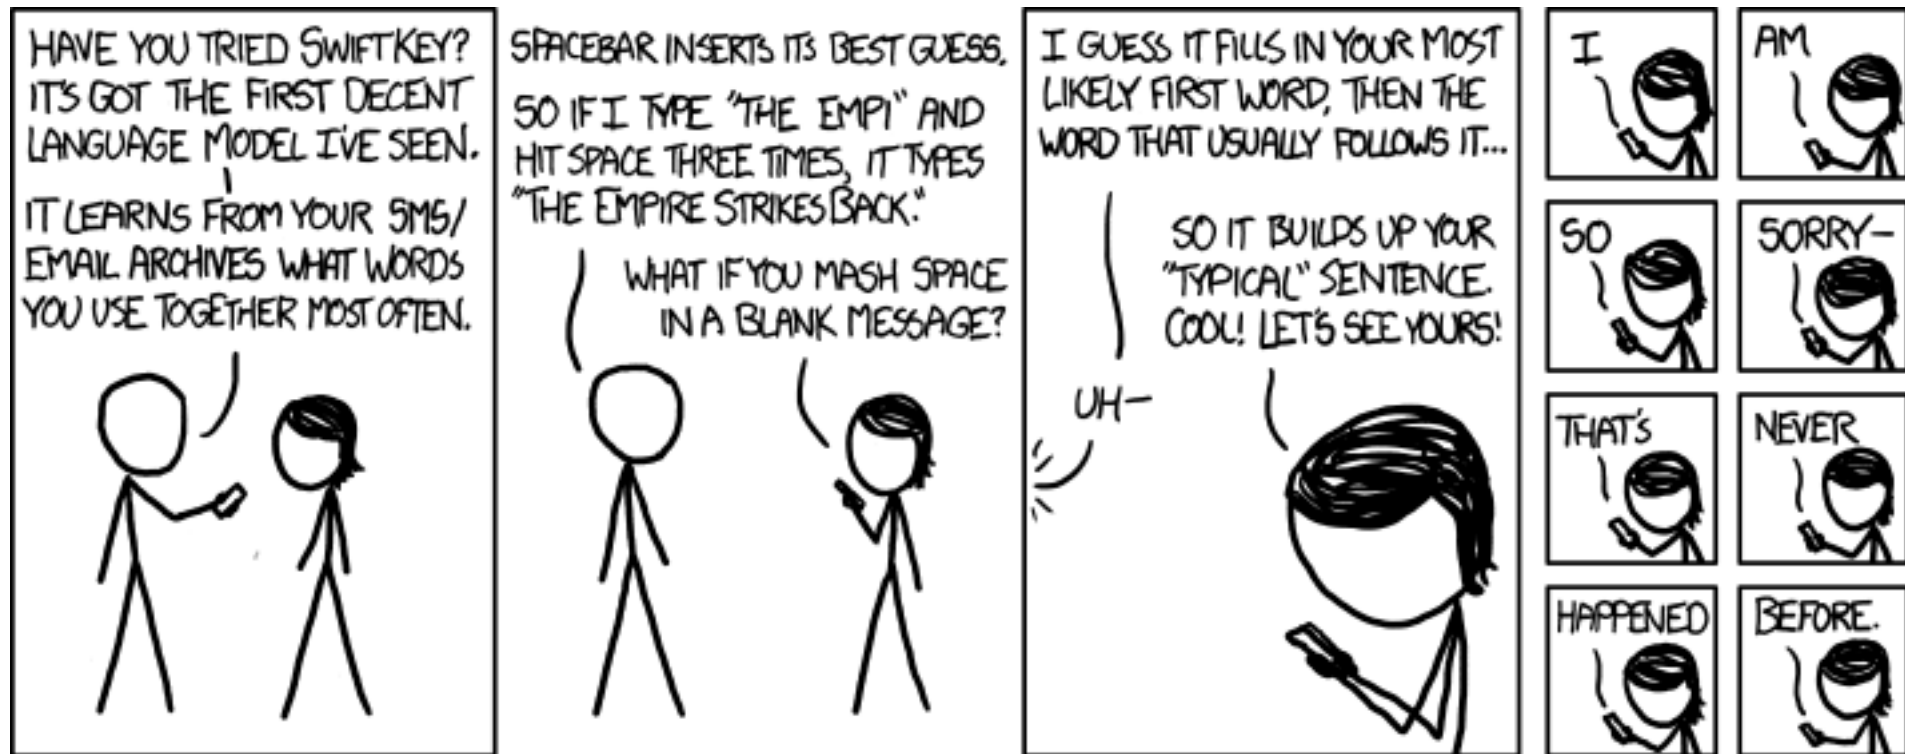
- **Agent**
 - Interacts with the environment. Your model guides the Agent's decisions
- **Environment**
 - Anything that is not the agent
- **Observations**
 - What the agent knows about the environment (usually state)
- **Actions**
 - What an agent can perform with the given environment (possibly stochastic)
- **Rewards**
 - Local measure of success
 - Can compound local rewards over time

State, Action, Reward, Next State

SARS 🤨



Markov Building Blocks



Markov Processes


- **Definition:** Any process that can be explained (or simplified) through a sequential set of states that depend only on the previous state
- **Practical Meaning:** For N states, there will be the probability of transition to any other state, encoded through an $N \times N$ transition matrix of discrete probabilities
- State sequences are not deterministic, they are sampled from these distributions
- Despite **simplicity**, they can model a number of real processes with **good enough** precision

	Next State, s_{t+1}				
	0.1	0.2	0.1	0.6	0.0
Current State, s_t	0.9	0.0	0.1	0.0	0.0
	0.0	0.4	0.0	0.4	0.2
	0.0	0.4	0.2	0.0	0.4
	0.0	0.0	0.6	0.0	0.4



MP Example from Maxim Lapan

	Sunny'	Rainy'
Sunny	0.8	0.2
Rainy	0.6	0.4




				...	
Sun+Summer					
Rainy+Summer					
Sun+Fall					
Rainy+Fall					
Sun+Else					
Rainy+Else					

**Adding One Variable Can Have
Drastic Effect on State Space Size**



Markov Decision Processes (MDP)

- **New Definition:** any state to state transition can be altered by an action that is given by a Markov Process
- **Definition:** An MDP consists of:
 - Env. States, s_t
 - Actions set for each time a_t
 - Reward function for each state, $r(s_t)$
 - A transition model, $P(s_{t+1}, s_t \mid a)$
a matrix of probabilities
 - ♦ Not **guaranteed** next state by given action



Markov Reward Process (MRP)

- **Total reward** is given by sum of all rewards in sequence

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots = \sum_k \gamma^k R_{t+k+1}$$

- Gamma defines future reward far- and short-sightedness
 - Common values are **0** (short), **0.9**, **0.99**, and **1** (far)
- This reward calculation can be used to estimate the “**Value**” of each state based upon the average total reward a state *should* give, $V(s) = \mathbf{E}[G \mid s_t=s]$
- Typically, this value must be estimated from the model over fixed sequences, otherwise some reward values can become arbitrarily large by looping actions




MDPs and MRPs

- The million dollar question:
How do we select a good action given a current state?
- If γ is not 0, this can get really complicated as we need to look at all possible future actions to measure value
- Instead of defining what is optimal, let's instead setup a comparison of different actions we might take (**policy**)
- A **policy** is defined as $\pi(a, s) = P(a_t = a \mid s_t = s)$
 - Given the current state, we have a certain probability of selecting each action
 - Action selection is **probabilistic**, but easy to discover **deterministic** actions (*set one action to 1.0, all others to 0.0*)
- Try different policies, select one with best average reward




An Illustrative Example: Grid World




- **State:** Every square in grid
- **Action:** Move to make (l,r,u,d), *with probability*
- **Reward:** Goal, Death
- **Policy:** Given state, where should we move?
- **Optimal Policy:**


$$\pi^* = \arg \max_{\pi} \mathbf{E} \left[\sum_k \gamma^k R_{t+k+1} \mid \pi \right]$$



Random Policy




Another Policy



Another Policy



RL Categorization



RL Categorizations

- Model-based versus Model-free
- Policy-based versus Value-based
- On-Policy, Off-Policy

- On-policy
 - We must interact with environment to learn a policy
- Off-policy
 - Can learn also from historical data or humans



Model-based versus Model-free

- Model Based
 - Predict the next observation and reward based on an understanding (model) of the rules in environment
 - Often look a number of moves ahead (like in chess or similar game)
 - Hard to construct in complex environments
 - NOT what we will be studying... needs domain expertise
- Model Free
 - Don't care what the environment is
 - Directly try to connect observations to actions (or values from which an action can be inferred)
 - Just use a neural network! That is our style!
- Mixed: Sure, like Alpha-Go




Policy Based versus Value Based

- Policy Based Learning
 - Directly approximate the policy of the agent
 - Policy is typically a probability distribution of actions that we sample from for next action
 - Could also be a “see this, do that” configuration
- Value Based
 - Calculate an intermediate value function for all possible actions
 - Policy becomes choosing the best action based on value function



OpenAI Gym



Object Oriented RL

- Basics:
 - Define object instance for **Agent ()** and the **Env ()**
 - Define what observations will return
 - Run **env.step(action)**
 - Get new observations and reward from env
- **action_space** and **observation_space**
 - Possible actions to execute, Observations to get
 - Discrete or continuous?
 - Can actions be given simultaneously?





Basics of Cartpole

```
import gym

if __name__ == "__main__":
    env = gym.make("CartPole-v0")

    total_reward = 0.0
    total_steps = 0
    obs = env.reset()

    while True:
        action = env.action_space.sample()
        obs, reward, done, _ = env.step(action)
        total_reward += reward
        total_steps += 1
        if done:
            break
```



Action Space: One input, $[0, 1]$ pull left or pull right

Obs Space: Dynamic state variables (continuous and four dimensional)

End: When more than 15 degrees off or too far from center

Reward: +1 for each time step



Wrapping the Environment

- When you want some extra action, observation, reward processing
- Expose function with **ActionWrapper**, **RewardWrapper**, **ObservationWrapper**

```
class RandomActionWrapper(gym.ActionWrapper):
    def __init__(self, env, epsilon=0.1):
        super(RandomActionWrapper, self).__init__(env)
        self.epsilon = epsilon

    def action(self, action):
        if random.random() < self.epsilon:
            print("Random!")
            return self.env.action_space.sample()
        return action
```

```
if __name__ == "__main__":
    env = RandomActionWrapper(gym.make("CartPole-v0"))

    obs = env.reset()
    total_reward = 0.0

    while True:
        obs, reward, done, _ = env.step(0)
        total_reward += reward
        if done:
            break
```

Might return different action than user supplied
with small probability



OpenAI Gym

<https://gym.openai.com>



We provide the environment; you provide the algorithm.
You can write your agent using your existing numerical computation library, such as TensorFlow or Theano.





Cross Entropy Method



Optimize Best Random Models

- Create a random neural network, with output $p(a|s)$
- Let it interact with the environment (randomly)
 - For some set of episodes (e.g., 20)
 - ◆ Use network output to sample from possible actions
 - ◆ Run episode to completion
 - ◆ Repeat
- Calculate reward for each episode
- Keep best episodes (some percentile, e.g., best five)
- For the given best episodes, develop loss function incentivizing the actions taken based upon the input observations



Repeat until desired performance!



Cross Entropy Method

- Model based or Model Free?
 - Model Free (no assumptions of problem)
- Value or Policy Based?
 - Policy Based (randomly sample actions based on policy)
- On-policy or Off-Policy?
 - On-Policy (need to interact with environment to get better)
- Has some similarity to **Simulated Annealing** Optimization



How to Make this More Mathy?

- If we have the optimal policy $p(x)$ and a reward function $H(x)$, then maximize

$$\mathbf{E}_{x \leftarrow p(x)}[H(x)] = \mathbf{E}_{x \leftarrow q(x)}\left[\frac{p(x)}{q(x)}H(x)\right]$$

- We can approximate the distribution by: $\frac{1}{N} \sum_i \frac{p(x_i)}{q(x_i)} H(x_i)$
- Proven that this is optimized when $\text{KL}(q(x) \parallel p(x)H(x))$ is minimized. But its intractable, so we drop terms ... and end up just minimizing (neg) cross entropy of samples

Performance

Measure

$$\pi_{k+1}(a | s) = \arg \max_{\pi_k} \mathbf{E}_{z \leftarrow \pi_k} [\mathbf{1}_{R(z) > \psi} \log \pi_k(a | s)]$$

min CrossEntropy(*net_actions*, *best_actions*)





Review: Basics of Cartpole

```
import gym

if __name__ == "__main__":
    env = gym.make("CartPole-v0")

    total_reward = 0.0
    total_steps = 0
    obs = env.reset()

    while True:
        action = env.action_space.sample()
        obs, reward, done, _ = env.step(action)
        total_reward += reward
        total_steps += 1
        if done:
            break
```



Action Space: One input, $[0, 1]$ pull left or pull right

Obs Space: Dynamic state variables (continuous and four dimensional)

End: When more than 15 degrees off or too far from center

Reward: +1 for each time step





Cross Entropy Reinforcement Learning

M. Lapan Implementation for CartPole
and Frozen Lake

Follow Along:

`08a_Basics_Of_Reinforcement_Learning.ipynb`



Lecture Notes for **Neural Networks and Machine Learning**

Intro to Reinforcement Learning



Next Time:
Q-Learning

Reading: Lapan CH4-CH6

