

# Lecture Notes for **Neural Networks and Machine Learning**



Cross Entropy and  
Value Iteration




# Logistics and Agenda

- Logistics
  - Atari paper next time!
  - Then, AlphaFold and SAC next week
- Agenda
  - OpenAI Gym
  - The Cross Entropy Method
  - Value Iteration
  - Q-Learning




# Final Project

## One Idea from Professor Stephanie Langin-Hooper SMU Meadows



# Last Time



- **State:** Every square in grid
- **Action:** Move to make (l,r,u,d), with probability
- **Reward:** Goal, Death
- **Policy:** Given state, where should we move?
- **Optimal Policy:**

$$\pi^* = \arg \max_{\pi} \mathbb{E} \left[ \sum_k \gamma^k R_{t+k+1} \mid \pi \right]$$



Random Policy



Another Policy



Another Policy



```

import gym

if __name__ == "__main__":
    env = gym.make("CartPole-v0")

    total_reward = 0.0
    total_steps = 0
    obs = env.reset()


    while True:
        action = env.action_space.sample()
        obs, reward, done, _ = env.step(action)
        total_reward += reward
        total_steps += 1
        if done:
            break
  
```



- Action Space:** One input, [0, 1] pull left or pull right
- Obs Space:** Dynamic state variables (continuous and four dimensional)
- End:** When more than 15 degrees off or too far from center
- Reward:** +1 for each time step



# OpenAI Gym



# Object Oriented RL

- Basics:
  - Define object instance for Agent() and the Env()
  - Define what observations will return
  - Run env.step(action)
  - Get new observations and reward from env
- **action\_space** and **observation\_space**
  - Possible actions to execute, Observations to get
  - Discrete or continuous?
  - Can actions be given simultaneously?





# Basics of Cartpole

```
import gym

if __name__ == "__main__":
    env = gym.make("CartPole-v0")

    total_reward = 0.0
    total_steps = 0
    obs = env.reset()

    while True:
        action = env.action_space.sample()
        obs, reward, done, _ = env.step(action)
        total_reward += reward
        total_steps += 1
        if done:
            break
```



**Action Space:** One input,  $[0, 1]$  pull left or pull right

**Obs Space:** Dynamic state variables (continuous and four dimensional)

**End:** When more than 15 degrees off or too far from center

**Reward:** +1 for each time step



# Wrapping the Environment

- When you want some extra action, observation, reward processing
- Expose function with **ActionWrapper**, **RewardWrapper**, **ObservationWrapper**

```
class RandomActionWrapper(gym.ActionWrapper):
    def __init__(self, env, epsilon=0.1):
        super(RandomActionWrapper, self).__init__(env)
        self.epsilon = epsilon

    def action(self, action):
        if random.random() < self.epsilon:
            print("Random!")
            return self.env.action_space.sample()
        return action
```

```
if __name__ == "__main__":
    env = RandomActionWrapper(gym.make("CartPole-v0"))

    obs = env.reset()
    total_reward = 0.0

    while True:
        obs, reward, done, _ = env.step(0)
        total_reward += reward
        if done:
            break
```

Might return different action than user supplied  
with small probability





# OpenAI Gym

<https://gym.openai.com>



We provide the environment; you provide the algorithm.  
You can write your agent using your existing numerical computation library, such as TensorFlow or Theano.


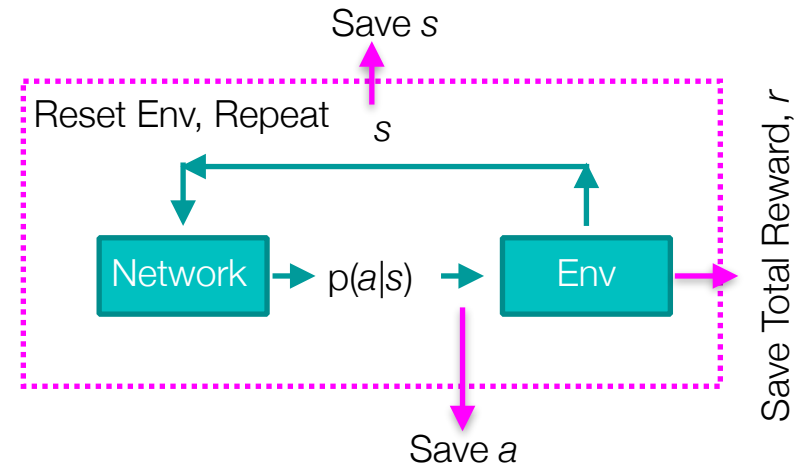


# Cross Entropy Method



# Optimize Best Random Models

- Create a random neural network
- Let it interact with the environment (randomly)
  - For some set of episodes (e.g., 20)
    - ◆ Use network output to sample from possible actions
    - ◆ Run episode to completion
    - ◆ Repeat
- Calculate reward for each episode
- Keep best episodes (some percentile, e.g., best five)
- For the given best episodes, develop loss function incentivizing the actions taken based upon the input observations



# Cross Entropy Method

- Model based or Model Free?
  - Model Free (no assumptions of problem)
- Value or Policy Based?
  - Policy Based (randomly sample actions based on policy)
- On-policy or Off-Policy?
  - On-Policy (need to interact with environment to get better)
- Has some similarity to **Simulated Annealing** Optimization



# How to Make this More Mathy?

- If we have the optimal policy  $p(x)$  and a reward function  $H(x)$ , then maximize

$$\mathbf{E}_{x \leftarrow p(x)}[H(x)] = \mathbf{E}_{x \leftarrow q(x)}\left[\frac{p(x)}{q(x)}H(x)\right]$$

- We can approximate the distribution by:  $\frac{1}{N} \sum_i \frac{p(x_i)}{q(x_i)} H(x_i)$
- Proven that this is optimized when  $\text{KL}(q(x) \parallel p(x)H(x))$  is minimized. But its intractable, so we drop terms ... and end up just minimizing (neg) cross entropy of samples

$$\pi_{k+1}(a \mid s) = \arg \max_{\pi_k} \mathbf{E}_{z \leftarrow \pi_k} \left[ \overset{\text{Performance Measure}}{\mathbf{1}_{R(z) > \psi}} \log \pi_k(a \mid s) \right]$$

`min CrossEntropy( net_actions, best_actions)`




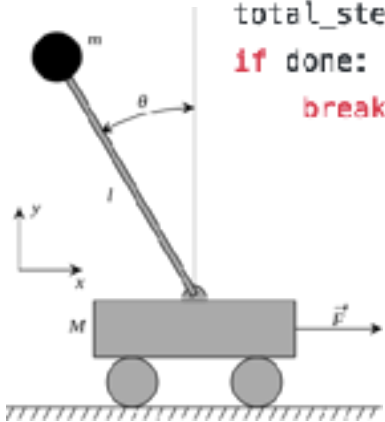
# Review: Basics of Cartpole

```
import gym

if __name__ == "__main__":
    env = gym.make("CartPole-v0")

    total_reward = 0.0
    total_steps = 0
    obs = env.reset()

    while True:
        action = env.action_space.sample()
        obs, reward, done, _ = env.step(action)
        total_reward += reward
        total_steps += 1
        if done:
            break
```



**Action Space:** One input,  $[0, 1]$  pull left or pull right

**Obs Space:** Dynamic state variables (continuous and four dimensional)

**End:** When more than 15 degrees off or too far from center

**Reward:** +1 for each time step





# Cross Entropy Reinforcement Learning

M. Lapan Implementation for CartPole  
and Frozen Lake

Follow Along:

`08a_Basics_Of_Reinforcement_Learning.ipynb`



# Value Iteration





# State Value Review

- Given: 
$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots = \sum_k \gamma^k R_{t+k+1}$$
- $V(s) = \mathbf{E}[G \mid s_t=s]$ , expected Value of a given state over all future iterations
- **Important:** we can only calculate this exactly if we know:
  - all the rewards for all the states
  - the probabilities of transitioning to a given state from selecting an action
  - likelihood of successful action
  - Most of the time **we know none of this** when we approach the problem, because it assumes a model of the system



# The Bellman Equation

- For the case when each action is successful and state is discrete, current  $V$  is easy to calculate:

$$V_0 = \max_{a \in 1 \dots A} (r_a + \gamma V_a)$$

current value is immediate reward plus value of next state with highest value

- Which feels like cheating because we assume we know  $V_a \dots$  just go with it for now
- General extension for when actions are probabilistic:

$$V_0 = \max_{a \in A} \mathbf{E}[r_{s,a} + \gamma V_s] = \max_{a \in A} \sum_{s \in S} p_{a,0 \rightarrow s} \cdot (r_{s,a} + \gamma V_s)$$

-probabilities of getting to next state  $s$  (current value is immediate reward plus value of next state)

- $p_{a,0 \rightarrow s}$  probability of getting to state  $s$  from state  $0$ , given that you perform action  $a$

- To select action with best value we need reward matrix,  $r_{s,a}$  and action transition matrix  $p_{a,0 \rightarrow s}$



# Defining the Q-Function

$$V_0 = \max_{a \in A} \mathbf{E}[r_{s,a} + \gamma V_s] = \max_{a \in A} \sum_{s \in S} p_{a,0 \rightarrow s} \cdot (r_{s,a} + \gamma V_s)$$

- Define intermediate function Q

$$Q(s, a) = \sum_{s' \in S} p_{a,s \rightarrow s'} \cdot (r_{s,a} + \gamma V_{s'})$$

- With some nice properties/relations:

$$V_s = \max_{a \in A} Q(s, a)$$

$$Q(s, a) = r_{s,a} + \gamma \max_{a' \in A} Q(s', a')$$



# Value Iteration (Value Based)

- **Direct:**

- Initialize  $V(s)$  to all zeros
- Take a series of random steps
- Perform for each state: 
$$V(s) \leftarrow \max_{a \in A} \sum_{s' \in S} p_{a,s \rightarrow s'} \cdot (r_{s,a} + \gamma V(s'))$$
- Repeat until  $V(s)$  stops changing

- **Q-Function Variant:**

- Initialize  $Q(s,a)$  to all zeros
- Take a series of random steps
- For each state and action: 
$$Q(s,a) \leftarrow \sum_{s' \in S} p_{a,s \rightarrow s'} \cdot (r_{s,a} + \gamma \max_{a'} Q(s', a'))$$
- Repeat until  $Q$  is not changing

Need to estimate  $p_{a,s \rightarrow s'}$   
Via observed **Transitions**

This Update Will **Converge to Optimal Policy**





# Value Iteration Reinforcement Learning

M. Lapan Implementation for  
and Frozen Lake

Follow Along:  
`08a_Basics_Of_Reinforcement_Learning.ipynb`



# Some Limitations

- Q function can get really big for **large states and action spaces**
- Infinite when the spaces are **continuous**
  - We will solve this by using a neural network to **approximate** the Q function
- Transition matrix, similarly, can get gigantic for large state and action spaces
  - We will solve this by dropping the transition probabilities in Q function update
- This Variant is known as Q-Learning



# Lecture Notes for **Neural Networks and Machine Learning**

CE and Value Iteration



**Next Time:**

Deep Q-Learning

**Reading:** Lapan CH6, CH7

